

prep code test code real code

```

private void checkUserGuess(String userGuess) {
    numOfGuesses++; (11)
    String result = "miss"; (12)
    for (DotCom dotComToTest : dotComsList) { (13)
        result = dotComToTest.checkYourself(userGuess); (14)
        if (result.equals("hit")) {
            break; (15)
        }
        if (result.equals("kill")) {
            dotComsList.remove(dotComToTest); (16)
            break;
        }
    } // close for
    System.out.println(result); (17)
} // close method

private void finishGame() {
    System.out.println("All Dot Coms are dead! Your stock is now worthless.");
    if (numOfGuesses <= 18) {
        System.out.println("It only took you " + numOfGuesses + " guesses.");
        System.out.println("You got out before your options sank.");
    } else {
        System.out.println("Took you long enough. " + numOfGuesses + " guesses.");
        System.out.println("Fish are dancing with your options.");
    }
} // close method

public static void main (String[] args) {
    DotComBust game = new DotComBust(); (19)
    game.setUpGame(); (20)
    game.startPlaying(); (21)
} // close method
}

```

**Whatever you do,
DON'T turn the
page!**

**Not until you've
finished this
exercise.**

**Our version is on
the next page.**



(18)

— repeat with all DotComs in the list
 — this guy's dead, so take him out of the DotComs list then get out of the loop
 — increment the number of guesses the user has made
 — get out of the loop early, no point in testing the others

— print a message telling the user how he did in the game
 — assume it's a 'miss', unless told otherwise
 — tell the game object to start the main game play loop (keeps asking for user input and checking the guess)

— print the result for the user
 — ask the DotCom to check the user guess, looking for a hit (or kill)
 — create the game object

— tell the game object to set up the game

the DotComBust code (the game)

prep code test code real code

```
import java.util.*;
public class DotComBust {
    private GameHelper helper = new GameHelper();
    private ArrayList<DotCom> dotComsList = new ArrayList<DotCom>();
    private int numOfGuesses = 0;

    private void setUpGame() {
        // first make some dot coms and give them locations
        DotCom one = new DotCom();
        one.setName("Pets.com");
        DotCom two = new DotCom();
        two.setName("eToys.com");
        DotCom three = new DotCom();
        three.setName("Go2.com");
        dotComsList.add(one);
        dotComsList.add(two);
        dotComsList.add(three);

        System.out.println("Your goal is to sink three dot coms.");
        System.out.println("Pets.com, eToys.com, Go2.com");
        System.out.println("Try to sink them all in the fewest number of guesses");

        for (DotCom dotComToSet : dotComsList) { ← Repeat with each DotCom in the list.
            ArrayList<String> newLocation = helper.placeDotCom(3); ← Ask the helper for a
            dotComToSet.setLocationCells(newLocation); ← Call the setter method on this
        } // close for loop
    } // close setUpGame method

    private void startPlaying() {
        while(!dotComsList.isEmpty()) { ← As long as the DotCom list is NOT empty (the ! means NOT, it's
            String userGuess = helper.getUserInput("Enter a guess"); ← Get user input.
            checkUserGuess(userGuess); ← Call our own checkUserGuess method.
        } // close while
        finishGame(); ← Call our own finishGame method.
    } // close startPlaying method
```

Declare and initialize the variables we'll need.

Make three DotCom objects, give 'em names, and stick 'em in the ArrayList.

Print brief instructions for user.

← Repeat with each DotCom in the list.

← Ask the helper for a DotCom location (an ArrayList of Strings).

← Call the setter method on this DotCom to give it the location you just got from the helper.

← As long as the DotCom list is NOT empty (the ! means NOT, it's the same as (dotComsList.isEmpty() == false)).

← Get user input.

← Call our own checkUserGuess method.

← Call our own finishGame method.

prep code **test code** **real code**

```

private void checkUserGuess (String userGuess) {
    numOfGuesses++;           ← increment the number of guesses the user has made
    String result = "miss";   ← assume it's a 'miss', unless told otherwise
    for(int x = 0; x < dotComsList.size(); x++) { ← repeat with all DotComs in the list
        result = dotComsList.get(x).checkYourself(userGuess); ← ask the DotCom to check the user
        if (result.equals("hit")) {                            guess, looking for a hit (or kill)
            break;          ← get out of the loop early, no point
        }
        if (result.equals("kill")) {
            dotComsList.remove(x);                         ← this guy's dead, so take him out of the
            break;                                      DotComs list then get out of the loop
        }
    } // close for
    System.out.println(result); ← print the result for the user
} // close method
                                         print a message telling the
                                         user how he did in the game
                                         ↓
private void finishGame () {
    System.out.println("All Dot Coms are dead! Your stock is now worthless.");
    if (numOfGuesses <= 18) {
        System.out.println("It only took you " + numOfGuesses + " guesses.");
        System.out.println("You got out before your options sank.");
    } else {
        System.out.println("Took you long enough. " + numOfGuesses + " guesses.");
        System.out.println("Fish are dancing with your options");
    }
} // close method
                                         ↓
public static void main (String[] args) {
    DotComBust game = new DotComBust(); ← create the game object
    game.setUpGame();                  ← tell the game object to set up the game
    game.startPlaying();              ← tell the game object to start the game
} // close method
                                         play loop (keeps asking for user
                                         input and checking the guess)
                                         ↓

```

the DotCom code

prep code test code real code

The final version of the DotCom class

```
import java.util.*;  
  
public class DotCom {  
    private ArrayList<String> locationCells;  
    private String name;  
  
    public void setLocationCells(ArrayList<String> loc) { ← A setter method that updates  
        locationCells = loc;  
    } ← the DotCom's location.  
    (Random location provided by  
    the GameHelper placeDotCom()  
    method.)  
  
    public void setName(String n) { ← Your basic setter method  
        name = n;  
    } ←  
  
    public String checkYourself(String userInput) { ← The ArrayList indexOf() method in  
        String result = "miss"; ← action! If the user guess is one of the  
        int index = locationCells.indexOf(userInput); ← entries in the ArrayList, indexOf()  
        if (index >= 0) { ← will return its ArrayList location. If  
            locationCells.remove(index); ← not, indexOf() will return -1.  
            ← Using ArrayList's remove() method to delete an entry.  
  
            if (locationCells.isEmpty()) { ← Using the isEmpty() method to see if all  
                result = "kill"; ← of the locations have been guessed  
                System.out.println("Ouch! You sunk " + name + " : ( ");  
            } else { ← Tell the user when a DotCom has been sunk.  
                result = "hit";  
            } // close if  
        } // close if  
        return result; ← Return: 'miss' or 'hit' or 'kill'.  
    } // close method  
} // close class
```

Super Powerful Boolean Expressions

So far, when we've used boolean expressions for our loops or `if` tests, they've been pretty simple. We will be using more powerful boolean expressions in some of the Ready-Bake code you're about to see, and even though we know you wouldn't peek, we thought this would be a good time to discuss how to energize your expressions.

'And' and 'Or' Operators (`&&`, `||`)

Let's say you're writing a `chooseCamera()` method, with lots of rules about which camera to select. Maybe you can choose cameras ranging from \$50 to \$1000, but in some cases you want to limit the price range more precisely. You want to say something like:

'If the price range is between \$300 *and* \$400 then choose X.'

```
if (price >= 300 && price < 400) {
    camera = "X";
}
```

Let's say that of the ten camera brands available, you have some logic that applies to only a few of the list:

```
if (brand.equals("A") || brand.equals("B")) {
    // do stuff for only brand A or brand B
}
```

Boolean expressions can get really big and complicated:

```
if ((zoomType.equals("optical") &&
    (zoomDegree >= 3 && zoomDegree <= 8)) ||
    (zoomType.equals("digital") &&
    (zoomDegree >= 5 && zoomDegree <= 12))) {
    // do appropriate zoom stuff
}
```

If you want to get *really* technical, you might wonder about the *precedence* of these operators. Instead of becoming an expert in the arcane world of precedence, we recommend that you *use parentheses* to make your code clear.

Not equals (`!=` and `!`)

Let's say that you have a logic like, "of the ten available camera models, a certain thing is *true for all but one*."

```
if (model != 2000) {
    // do non-model 2000 stuff
}
or for comparing objects like strings...
if (!brand.equals("X")) {
    // do non-brand X stuff
}
```

Short Circuit Operators (`&&` , `||`)

The operators we've looked at so far, `&&` and `||`, are known as *short circuit* operators. In the case of `&&`, the expression will be true only if *both* sides of the `&&` are true. So if the JVM sees that the left side of a `&&` expression is false, it stops right there! Doesn't even bother to look at the right side.

Similarly, with `||`, the expression will be true if *either* side is true, so if the JVM sees that the left side is true, it declares the entire statement to be true and doesn't bother to check the right side.

Why is this great? Let's say that you have a reference variable and you're not sure whether it's been assigned to an object. If you try to call a method using this null reference variable (i.e. no object has been assigned), you'll get a `NullPointerException`. So, try this:

```
if (refVar != null &&
    refVar.isValidType()) {
    // do 'got a valid type' stuff
}
```

Non Short Circuit Operators (`&` , `|`)

When used in boolean expressions, the `&` and `|` operators act like their `&&` and `||` counterparts, except that they force the JVM to *always* check *both* sides of the expression. Typically, `&` and `|` are used in another context, for manipulating bits.

Ready-bake: GameHelper



```
import java.io.*;
import java.util.*;
```

```
public class GameHelper {  
  
    private static final String alphabet = "abcdefg";  
    private int gridLength = 7;  
    private int gridSize = 49;  
    private int [] grid = new int[gridSize];  
    private int comCount = 0;  
  
    public String getUserInput(String prompt) {  
        String inputLine = null;  
        System.out.print(prompt + " ");  
        try {  
            BufferedReader is = new BufferedReader(  
                new InputStreamReader(System.in));  
            inputLine = is.readLine();  
            if (inputLine.length() == 0) return null;  
        } catch (IOException e) {  
            System.out.println("IOException: " + e);  
        }  
        return inputLine.toLowerCase();  
    }  
  
    public ArrayList<String> placeDotCom(int comSize) {  
        ArrayList<String> alphaCells = new ArrayList<String>();  
  
        String temp = null;  
        int [] coords = new int[comSize];  
        int attempts = 0;  
        boolean success = false;  
        int location = 0;  
  
        comCount++;  
        int incr = 1;  
        if ((comCount % 2) == 1) {  
            incr = gridLength;  
        }  
  
        while ( !success & attempts++ < 200 ) {  
            location = (int) (Math.random() * gridSize);  
            //System.out.print(" try " + location);  
            int x = 0;  
            success = true;  
            while (success && x < comSize) {  
                if (grid[location] == 0) {  
                    alphaCells.add(location + "" + x);  
                    grid[location] = 1;  
                    x++;  
                } else {  
                    success = false;  
                }  
            }  
        }  
    }  
}
```

This is the helper class for the game. Besides the user input method (that prompts the user and reads input from the command-line), the helper's Big Service is to create the cell locations for the DotComs. If we were you, we'd just back away slowly from this code, except to type it in and compile it. We tried to keep it fairly small so you wouldn't have to type so much, but that means it isn't the most readable code. And remember, you won't be able to compile the DotComBust game class until you have this class.

Note: For extra credit, you might try 'un-commenting' the `System.out.print(``ln`)'s in the `placeDotCom()` method, just to watch it work! These print statements will let you "cheat" by giving you the location of the DotComs, but it will help you test it.



Ready-bake Code

GameHelper class code continued...

```

        coords[x++] = location;           // save location
        location += incr;               // try 'next' adjacent
        if (location >= gridSize){      // out of bounds - 'bottom'
            success = false;           // failure
        }
        if (x>0 && (location % gridLength == 0)) { // out of bounds - right edge
            success = false;           // failure
        }
    } else {                           // found already used location
        // System.out.print(" used " + location);
        success = false;               // failure
    }
}
}                                     // end while

int x = 0;                            // turn location into alpha coords
int row = 0;
int column = 0;
// System.out.println("\n");
while (x < comSize) {
    grid[coords[x]] = 1;              // mark master grid pts. as 'used'
    row = (int) (coords[x] / gridLength); // get row value
    column = coords[x] % gridLength;   // get numeric column value
    temp = String.valueOf(alphabet.charAt(column)); // convert to alpha

    alphaCells.add(temp.concat(Integer.toString(row)));
    x++;
    // System.out.print(" coord "+x+" = " + alphaCells.get(x-1));
}
// System.out.println("\n");

return alphaCells;
}
}

```

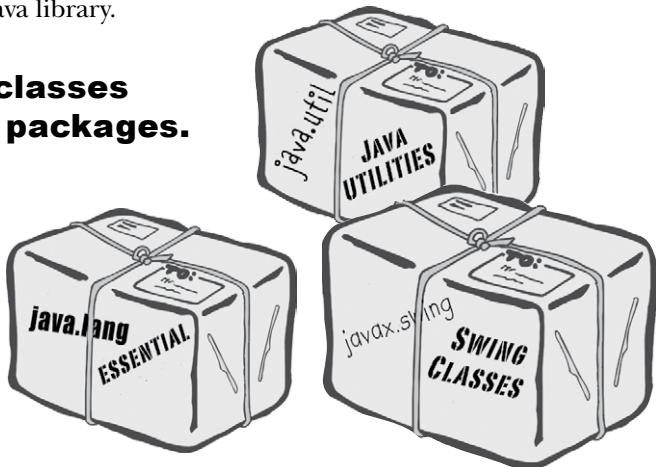
This is the statement that
tells you exactly where the
DotCom is located.

API packages

Using the Library (the Java API)

You made it all the way through the DotComBust game, thanks to the help of ArrayList. And now, as promised, it's time to learn how to fool around in the Java library.

In the Java API, classes are grouped into packages.



To use a class in the API, you have to know which package the class is in.

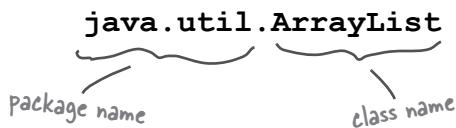
Every class in the Java library belongs to a package. The package has a name, like **javax.swing** (a package that holds some of the Swing GUI classes you'll learn about soon). ArrayList is in the package called **java.util**, which surprise surprise, holds a pile of *utility* classes. You'll learn a lot more about packages in chapter 17, including how to put your *own* classes into your *own* packages. For now though, we're just looking to *use* some of the classes that come with Java.

Using a class from the API, in your own code, is simple. You just treat the class as though you wrote it yourself... as though you compiled it, and there it sits, waiting for you to use it. With one big difference: somewhere in your code you have to indicate the *full* name of the library class you want to use, and that means package name + class name.

Even if you didn't know it, *you've already been using classes from a package*. System (System.out.println), String, and Math (Math.random()), all belong to the **java.lang** package.

You have to know the full name* of the class you want to use in your code.

ArrayList is not the *full name* of ArrayList, just as ‘Kathy’ isn’t a full name (unless it’s like Madonna or Cher, but we won’t go there). The full name of ArrayList is actually:



You have to tell Java which ArrayList you want to use. You have two options:

A IMPORT

Put an import statement at the top of your source code file:

```
import java.util.ArrayList;
public class MyClass { ... }
```

OR

B TYPE

Type the full name everywhere in your code. Each time you use it. *Anywhere* you use it.

When you declare and/or instantiate it:

```
java.util.ArrayList<Dog> list = new java.util.ArrayList<Dog>();
```

When you use it as an argument type:

```
public void go(java.util.ArrayList<Dog> list) { }
```

When you use it as a return type:

```
public java.util.ArrayList<Dog> foo() { ... }
```

*Unless the class is in the java.lang package.

there are no
Dumb Questions

Q: Why does there have to be a full name? Is that the only purpose of a package?

A: Packages are important for three main reasons. First, they help the overall organization of a project or library. Rather than just having one horrendously large pile of classes, they’re all grouped into packages for specific kinds of functionality (like GUI, or data structures, or database stuff, etc.)

Second, packages give you a name-scoping, to help prevent collisions if you and 12 other programmers in your company all decide to make a class with the same name. If you have a class named Set and someone else (including the Java API) has a class named Set, you need some way to tell the JVM *which* Set class you’re trying to use.

Third, packages provide a level of security, because you can restrict the code you write so that only other classes in the same package can access it. You’ll learn all about that in chapter 17.

Q: OK, back to the name collision thing. How does a full name really help? What’s to prevent two people from giving a class the same package name?

A: Java has a naming convention that usually prevents this from happening, as long as developers adhere to it. We’ll get into that in more detail in chapter 17.



BULLET POINTS

- **ArrayList** is a class in the Java API.
- To put something into an ArrayList, use **add()**.
- To remove something from an ArrayList use **remove()**.
- To find out where something is (and if it is) in an ArrayList, use **indexOf()**.
- To find out if an ArrayList is empty, use **isEmpty()**.
- To get the size (number of elements) in an ArrayList, use the **size() method**.
- To get the **length** (number of elements) in a regular old array, remember, you use the **length variable**.
- An ArrayList **resizes dynamically** to whatever size is needed. It grows when objects are added, and it **shrinks** when objects are removed.
- You declare the type of the array using a **type parameter**, which is a type name in angle brackets. Example: `ArrayList<Button>` means the ArrayList will be able to hold only objects of type Button (or subclasses of Button as you'll learn in the next couple of chapters).
- Although an ArrayList holds objects and not primitives, the compiler will automatically "wrap" (and "unwrap" when you take it out) a primitive into an Object, and place that object in the ArrayList instead of the primitive. (More on this feature later in the book.)
- Classes are grouped into packages.
- A class has a full name, which is a combination of the package name and the class name. Class `ArrayList` is really `java.util.ArrayList`.
- To use a class in a package other than `java.lang`, you must tell Java the full name of the class.
- You use either an import statement at the top of your source code, or you can type the full name every place you use the class in your code.

Where'd that 'x' come from? (or, what does it mean when a package starts with javax?)



In the first and second versions of Java (1.02 and 1.1), all classes that shipped with Java (in other words, the standard library) were in packages that began with **java**. There was always **java.lang**, of course — the one you don't have to import. And there was **java.net**, **java.io**, **java.util** (although there was no such thing as ArrayList way back then), and a few others, including the **java.awt** package that held GUI-related classes.

Looming on the horizon, though, were other packages not included in the standard library. These classes were known as **extensions**, and came in two main flavors: *standard*, and *not standard*. Standard extensions were those that Sun considered official, as opposed to experimental, early access, or beta packages that might or might not ever see the light of day.

Standard extensions, by convention, all began with an 'x' appended to the regular **java** package starter. The mother of all standard extensions was the Swing library. It included several packages, all of which began with **javax.swing**.

But standard extensions can get promoted to first-class, ships-with-Java, standard-out-of-the-box library packages. And that's what happened to Swing, beginning with version 1.2 (which eventually became the first version dubbed 'Java 2').

"Cool", everyone thought (including us). "Now everyone who has Java will have the Swing classes, and we won't have to figure out how to get those classes installed with our end-users."

Trouble was lurking beneath the surface, however, because when packages get promoted, well of COURSE they have to start with **java**, not **javax**. Everyone KNOWS that packages in the standard library don't have that "x", and that only extensions have the "x". So, just (and we mean just) before version 1.2 went final, Sun changed the package names and deleted the "x" (among other changes). Books were printed and in stores featuring Swing code with the new names. Naming conventions were intact. All was right with the Java world.

Except the 20,000 or so screaming developers who realized that with that simple name change came disaster! All of their Swing-using code had to be changed! The horror! Think of all those import statements that started with **javax**...

And in the final hour, desperate, as their hopes grew thin, the developers convinced Sun to "screw the convention, save our code". The rest is history. So when you see a package in the library that begins with **javax**, you know it started life as an extension, and then got a promotion.

there are no
Dumb Questions

Q: Does `import` make my class bigger? Does it actually compile the imported class or package into my code?

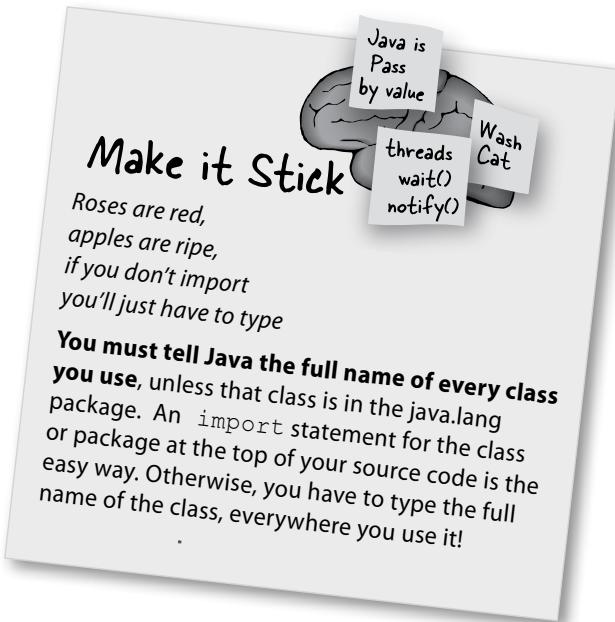
A: Perhaps you're a C programmer? An `import` is not the same as an `include`. So the answer is no and no. Repeat after me: "an `import` statement saves you from typing." That's really it. You don't have to worry about your code becoming bloated, or slower, from too many imports. An `import` is simply the way you give Java the *full name of a class*.

Q: OK, how come I never had to import the `String` class? Or `System`?

A: Remember, you get the `java.lang` package sort of "pre-imported" for free. Because the classes in `java.lang` are so fundamental, you don't have to use the full name. There is only one `java.lang.String` class, and one `java.lang.System` class, and Java darn well knows where to find them.

Q: Do I have to put my own classes into packages? How do I do that? Can I do that?

A: In the real world (which you should try to avoid), yes, you *will* want to put your classes into packages. We'll get into that in detail in chapter 17. For now, we won't put our code examples in a package.



One more time, in the unlikely event that you don't already have this down:



getting to know the API

"Good to know there's an ArrayList in the java.util package. But by myself, how would I have figured that out?"

- Julia, 31, hand model

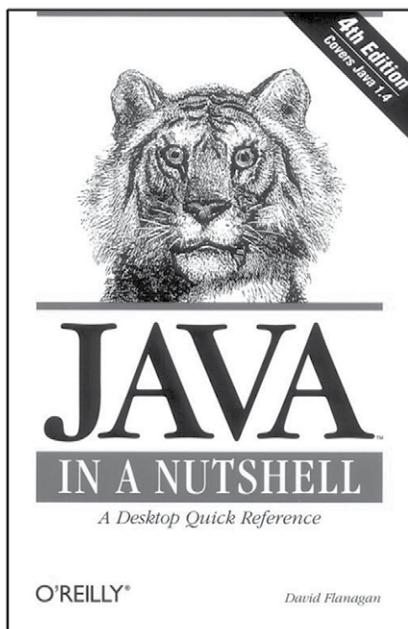
How to play with the API

Two things you want to know:

- 1 What classes are in the library?**
- 2 Once you find a class, how do you know what it can do?**



1 Browse a Book



2 Use the HTML API docs

Java™ 2 SDK, Standard Edition Documentation

Version 1.4.0

Search General Info API & Language Guide to Features Tool Docs Demos/Tutorials

Your feedback is important to us. Please send us comments: [Contacting Java™ Software.](#)

Search the Documentation Location

[Search the online documentation](#) [website](#)

1**Browse a Book**

Flipping through a reference book is the best way to find out what's in the Java library. You can easily stumble on a class that looks useful, just by browsing pages.

class name
package name
class description

methods (and other things we'll talk about later)

java.util.Currency

Returned By: java.text.DecimalFormat.getCurrency(), java.text.DecimalFormatSymbols.getCurrency(), java.text.NumberFormat.getCurrency(), Currency.getInstance()

Date java.util	Java 1.0 <i>cloneable serializable comparable</i>
---------------------------------	-------------------------------------------------------------

This class represents dates and times and lets you work with them in a system-independent way. You can create a Date by specifying the number of milliseconds from the epoch (midnight GMT, January 1st, 1970) or the year, month, date, and, optionally, the hour, minute, and second. Years are specified as the number of years since 1900. If you call the Date constructor with no arguments, the Date is initialized to the current time and date. The instance methods of the class allow you to get and set the various date and time fields, to compare dates and times, and to convert dates to and from string representations. As of Java 1.1, many of the date methods have been deprecated in favor of the methods of the Calendar class.

```

Object ---> Date
Cloneable Comparable Serializable

```

```

public class Date implements Cloneable, Comparable, Serializable {
// Public Constructors
    public Date();
    public Date(long date);
    # public Date(String s);
    # public Date(int year, int month, int date);
    # public Date(int year, int month, int date, int hrs, int min);
    # public Date(int year, int month, int date, int hrs, int min, int sec);
// Property Accessor Methods (by property name)
    public long getTime();
    public void setTime(long time);
// Public Instance Methods
    public boolean after(java.util.Date when);
    public boolean before(java.util.Date when);
    1.2 public int compareTo(java.util.Date anotherDate);
// Methods Implementing Comparable
    1.2 public int compareTo(Object o);
// Public Methods Overriding Object
    1.2 public Object clone();
    public boolean equals(Object obj);
    public int hashCode();
    public String toString();
// Deprecated Public Methods
    # public int getDate();
    # public int getDay();
    # public int getHours();
    # public int getMinutes();
    # public int getMonth();
    # public int getSeconds();
    # public int gettimezoneOffset();
    # public int getYear();
    # public static long parse(String s);
    # public void setDate(int date);
    # public void setHours(int hours);
    # public void setMinutes(int minutes);
    # public void setMonth(int month);

```

you are here ▶ **159**

using the Java API documentation



Use the HTML API docs

Java comes with a fabulous set of online docs called, strangely, the Java API. They're part of a larger set called the Java 5 Standard Edition Documentation (which, depending on what day of the week you look, Sun may be referring to as "Java 2 Standard Edition 5.0"), and you have to download the docs separately; they don't come shrink-wrapped with the Java 5 download. If you have a high-speed internet connection, or tons of patience, you can also browse them at java.sun.com. Trust us, you probably want these on your hard drive.

The API docs are the best reference for getting more details about a class and its methods. Let's say you were browsing through the reference book and found a class called `Calendar`, in `java.util`. The book tells you a little about it, enough to know that this is indeed what you want to use, but you still need to know more about the methods.

The reference book, for example, tells you what the methods take, as arguments, and what they return. Look at `ArrayList`, for example. In the reference book, you'll find the method `indexOf()`, that we used in the `DotCom` class. But if all you knew is that there is a method called `indexOf()` that takes an object and returns the index (an int) of that object, you still need to know one crucial thing: what happens if the object is not in the `ArrayList`? Looking at the method signature alone won't tell you how that works. But the API docs will (most of the time, anyway). The API docs tell you that the `indexOf()` method returns a -1 if the object parameter is not in the `ArrayList`. That's how we knew we could use it both as a way to check if an object is even *in* the `ArrayList`, and to get its index at the same time, if the object was there. But without the API docs, we might have thought that the `indexOf()` method would blow up if the object wasn't in the `ArrayList`.

① Scroll through the packages and select one (click it) to restrict the list in the lower frame to only classes from that package.

② Scroll through the classes and select one (click it) to choose the class that will fill the main browser frame.

This is where all the good stuff is. You can scroll through the methods for a brief summary, or click on a method to get full details.



Exercise

Code Magnets

Can you reconstruct the code snippets to make a working Java program that produces the output listed below? **NOTE:** To do this exercise, you need one NEW piece of info—if you look in the API for ArrayList, you'll find a *second* add method that takes two arguments:

`add(int index, Object o)`

It lets you specify to the ArrayList where to put the object you're adding.

`a.remove(2);`

`printAL(a);`

`a.add(0, "zero");
a.add(1, "one");`

`printAL(a);`

`if (a.contains("two")) {
 a.add("2.2");
}`

`a.add(2, "two");`

`public static void main (String[] args) {`

`System.out.print(element + " ");
}`

`System.out.println(" ");`

`if (a.contains("three")) {
 a.add("four");
}`

`public class ArrayListMagnet {`

`if (a.indexOf("four") != 4) {
 a.add(4, "4.2");
}`

`import java.util.*;`

`printAL(a);`

`ArrayList<String> a = new ArrayList<String>();`

`for (String element : al) {`

`a.add(3, "three");
printAL(a);`

File Edit Window Help Dance

```
% java ArrayListMagnet
zero  one  two  three
zero  one  three  four
zero  one  three  four  4.2
zero  one  three  four  4.2
```

puzzle: crossword



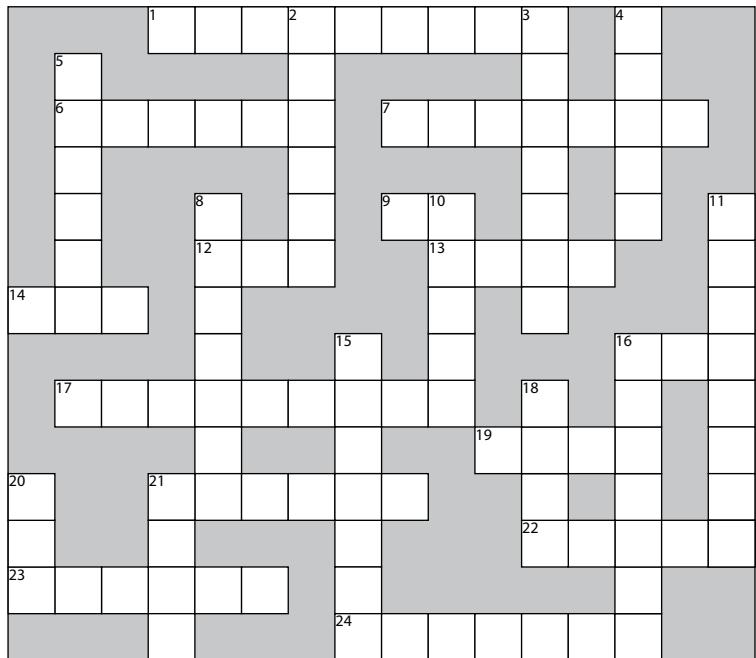
JavaCross 7.0

How does this crossword puzzle help you learn Java? Well, all of the words **are** Java related (except one red herring).

Hint: When in doubt, remember ArrayList.

Across

1. I can't behave
6. Or, in the courtroom
7. Where it's at baby
9. A fork's origin
12. Grow an ArrayList
13. Wholly massive
14. Value copy
16. Not an object
17. An array on steroids
19. Extent
21. 19's counterpart
22. Spanish geek snacks (Note: This has nothing to do with Java.)
23. For lazy fingers
24. Where packages roam



Down

2. Where the Java action is.
3. Addressable unit
4. 2nd smallest
5. Fractional default
8. Library's grandest
10. Must be low density
11. He's in there somewhere
15. As if
16. dearth method
18. What shopping and arrays have in common
20. Library acronym
21. What goes around

More Hints:

- | | | | | | | | | |
|--------|-----------------------|------------------------|--------------------|----------------------|----------------------|---------------------|-----------------------------------------|-----------------------------------------|
| Across | 1. 8 varieties | 2. What's overridable? | 3. Think ArrayList | 4. & 10. Primitive | 16. Common primitive | 18. Think ArrayList | 21. Array's extent | 22. Not about Java - Spanish appetizers |
| Down | 2. What's observable? | 3. Think ArrayList | 4. & 10. Primitive | 16. Common primitive | 18. Think ArrayList | 21. Array's extent | 22. Not about Java - Spanish appetizers | |



Exercise Solutions

File Edit Window Help Dance

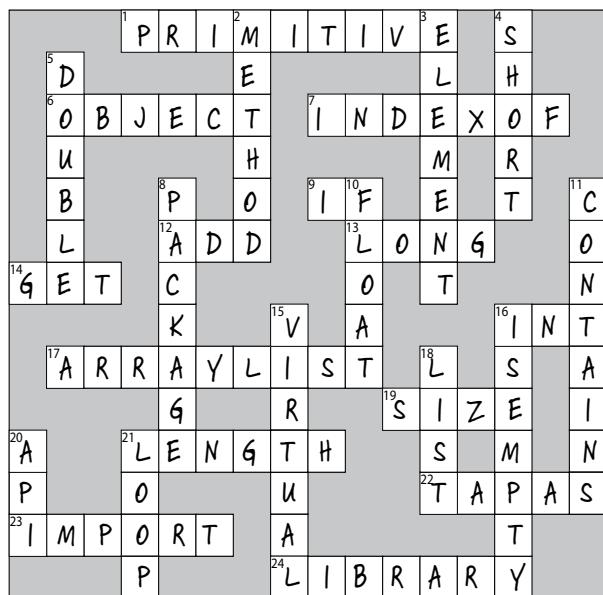
```
% java ArrayListMagnet
zero one two three
zero one three four
zero one three four 4.2
zero one three four 4.2
```

```
import java.util.*;
public class ArrayListMagnet {
    public static void main (String[] args) {
        ArrayList<String> a = new ArrayList<String>();
        a.add(0,"zero");
        a.add(1,"one");
        a.add(2,"two");
        a.add(3,"three");
        printAL(a);
        if (a.contains("three")) {
            a.add("four");
        }
        a.remove(2);
        printAL(a);
        if (a.indexOf("four") != 4) {
            a.add(4, "4.2");
        }
        printAL(a);
        if (a.contains("two")) {
            a.add("2.2");
        }
        printAL(a);
    }
    public static void printAL(ArrayList<String> al) {
        for (String element : al) {
            System.out.print(element + " ");
        }
        System.out.println(" ");
    }
}
```

puzzle answers



JavaCross answers



Sharpen your pencil

Write your OWN set of clues! Look at each word, and try to write your own clues. Try making them easier, or harder, or more technical than the ones we have.

Across

1. _____
6. _____
7. _____
9. _____
12. _____
13. _____
14. _____
16. _____
17. _____
19. _____
21. _____
22. _____
23. _____
24. _____

Down

2. _____
3. _____
4. _____
5. _____
8. _____
10. _____
11. _____
15. _____
16. _____
18. _____
20. _____
21. _____

7 inheritance and polymorphism

Better Living in Objectville



We were underpaid,
overworked coders 'till we
tried the Polymorphism Plan. But
thanks to the Plan, our future is
bright. Yours can be too!

Plan your programs with the future in mind. If there were a way to write Java code such that you could take more vacations, how much would it be worth to you? What if you could write code that someone *else* could extend, **easily**? And if you could write code that was flexible, for those pesky last-minute spec changes, would that be something you're interested in? Then this is your lucky day. For just three easy payments of 60 minutes time, you can have all this. When you get on the Polymorphism Plan, you'll learn the 5 steps to better class design, the 3 tricks to polymorphism, the 8 ways to make flexible code, and if you act now—a bonus lesson on the 4 tips for exploiting inheritance. Don't delay, an offer this good will give you the design freedom and programming flexibility you deserve. It's quick, it's easy, and it's available now. Start today, and we'll throw in an extra level of abstraction!

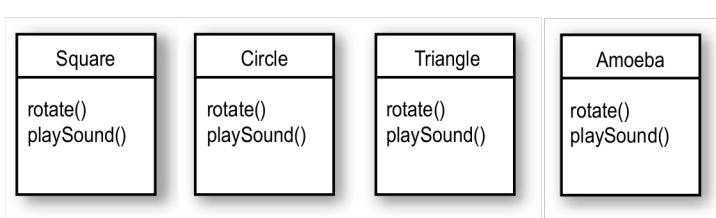
the power of inheritance

Chair Wars Revisited...

Remember way back in chapter 2, when Larry (procedural guy) and Brad (OO guy) were vying for the Aeron chair? Let's look at a few pieces of that story to review the basics of inheritance.

LARRY: You've got duplicated code! The rotate procedure is in all four Shape things. It's a stupid design. You have to maintain four different rotate "methods". How can that ever be good?

BRAD: Oh, I guess you didn't see the final design. Let me show you how OO **inheritance** works, Larry.

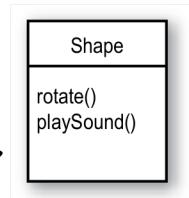


1

I looked at what all four classes have in common.

2

They're Shapes, and they all rotate and playSound. So I abstracted out the common features and put them into a new class called Shape.

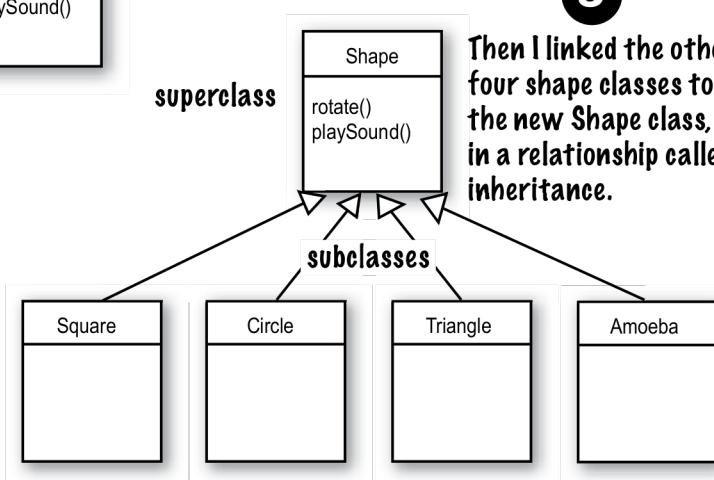


3

Then I linked the other four shape classes to the new Shape class, in a relationship called inheritance.

You can read this as, "Square inherits from Shape", "Circle inherits from Shape", and so on. I removed rotate() and playSound() from the other shapes, so now there's only one copy to maintain.

The Shape class is called the **superclass** of the other four classes. The other four are the **subclasses** of Shape. The subclasses inherit the methods of the superclass. In other words, if the Shape class has the functionality, then the subclasses automatically get that same functionality.



What about the Amoeba rotate()?

LARRY: Wasn't that the whole problem here — that the amoeba shape had a completely different rotate and playSound procedure?

How can amoeba do something different if it *inherits* its functionality from the Shape class?

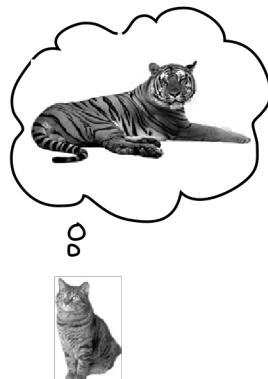
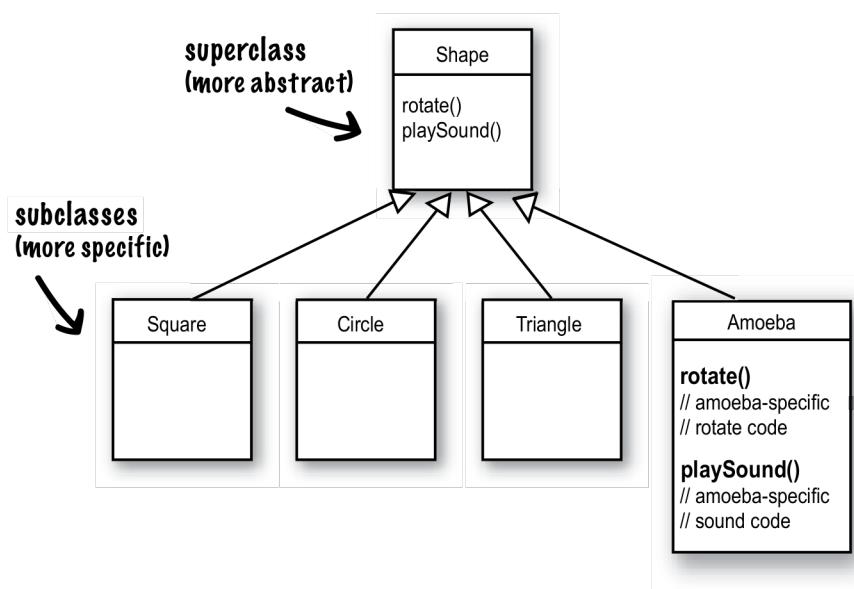
BRAD: That's the last step. The Amoeba class *overrides* the methods of the Shape class. Then at runtime, the JVM knows exactly which *rotate()* method to run when someone tells the Amoeba to rotate.



4

I made the Amoeba class override the *rotate()* and *playSound()* methods of the superclass Shape. Overriding just means that a subclass redefines one of its inherited methods when it needs to change or extend the behavior of that method.

Overriding methods



How would you represent a house cat and a tiger, in an inheritance structure. Is a domestic cat a specialized version of a tiger? Which would be the subclass and which would be the superclass? Or are they both subclasses to some other class?

How would you design an inheritance structure? What methods would be overridden?

Think about it. *Before* you turn the page.

the way inheritance works

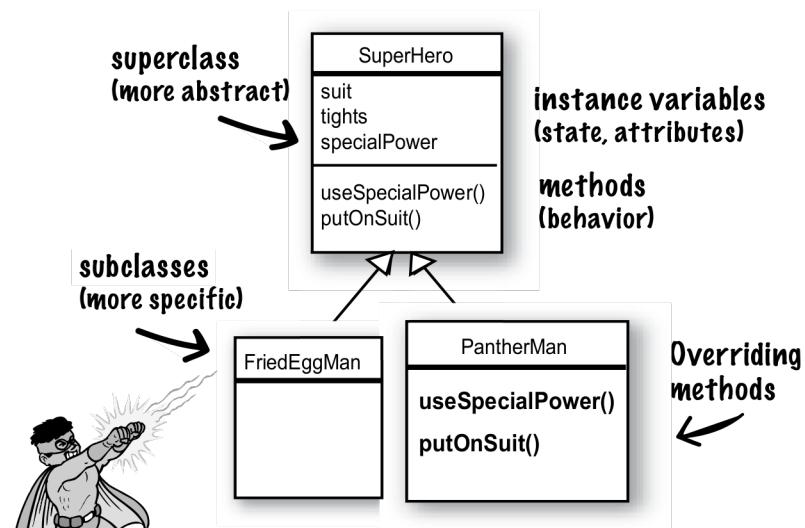
Understanding Inheritance

When you design with inheritance, you put common code in a class and then tell other more specific classes that the common (more abstract) class is their superclass. When one class inherits from another, **the subclass inherits from the superclass**.

In Java, we say that the **subclass extends the superclass**.

An inheritance relationship means that the subclass inherits the **members** of the superclass. When we say “members of a class” we mean the instance variables and methods.

For example, if PantherMan is a subclass of SuperHero, the PantherMan class automatically inherits the instance variables and methods common to all superheroes including `suit`, `tights`, `specialPower`, `useSpecialPower()` and so on. But the PantherMan **subclass can add new methods and instance variables** of its own, and it **can override the methods it inherits from the superclass** SuperHero.



FriedEggMan doesn't need any behavior that's unique, so he doesn't override any methods. The methods and instance variables in SuperHero are sufficient.

PantherMan, though, has specific requirements for his suit and special powers, so `useSpecialPower()` and `putOnSuit()` are both overridden in the PantherMan class.

Instance variables are not overridden because they don't need to be. They don't define any special behavior, so a subclass can give an inherited instance variable any value it chooses. PantherMan can set his inherited `tights` to purple, while FriedEggMan sets his to white.

inheritance and polymorphism

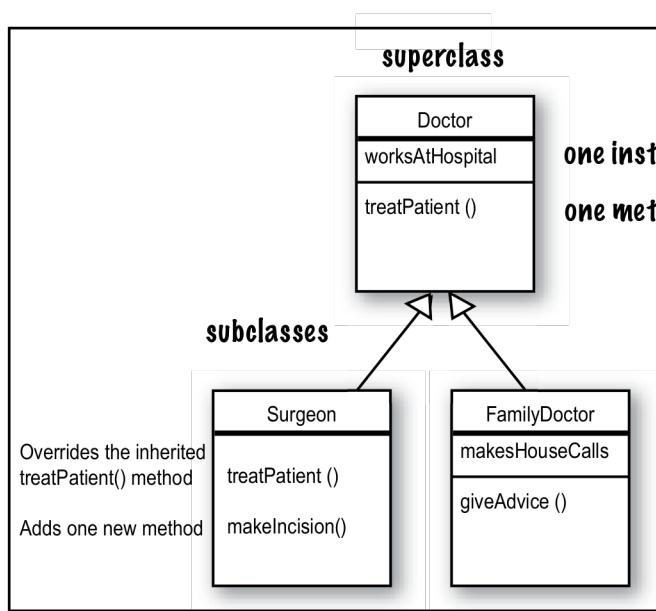
An inheritance example:

```
public class Doctor {
    boolean worksAtHospital;
    void treatPatient() {
        // perform a checkup
    }
}

public class FamilyDoctor extends Doctor {
    boolean makesHouseCalls;
    void giveAdvice() {
        // give homespun advice
    }
}

public class Surgeon extends Doctor{
    void treatPatient() {
        // perform surgery
    }

    void makeIncision() {
        // make incision (yikes!)
    }
}
```



How many instance variables does Surgeon have? _____

How many instance variables does FamilyDoctor have? _____

How many methods does Doctor have? _____

How many methods does Surgeon have? _____

How many methods does FamilyDoctor have? _____

Can a FamilyDoctor do treatPatient()? _____

Can a FamilyDoctor do makeIncision()? _____

designing for inheritance

Let's design the inheritance tree for an Animal simulation program

Imagine you're asked to design a simulation program that lets the user throw a bunch of different animals into an environment to see what happens. We don't have to code the thing now, we're mostly interested in the design.

We've been given a list of *some* of the animals that will be in the program, but not all. We know that each animal will be represented by an object, and that the objects will move around in the environment, doing whatever it is that each particular type is programmed to do.

And we want other programmers to be able to add new kinds of animals to the program at any time.

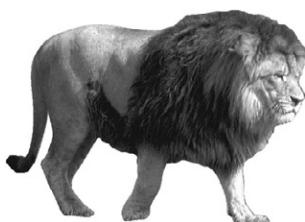
First we have to figure out the common, abstract characteristics that all animals have, and build those characteristics into a class that all animal classes can extend.

1

Look for objects that have common attributes and behaviors.

What do these six types have in common? This helps you to abstract out behaviors. (step 2)

How are these types related? This helps you to define the inheritance tree relationships (step 4-5)



Using inheritance to avoid duplicating code in subclasses

We have five *instance variables*:

picture – the file name representing the JPEG of this animal

food – the type of food this animal eats. Right now, there can be only two values: *meat* or *grass*.

hunger – an int representing the hunger level of the animal. It changes depending on when (and how much) the animal eats.

boundaries – values representing the height and width of the ‘space’ (for example, 640 x 480) that the animals will roam around in.

location – the X and Y coordinates for where the animal is in the space.

We have four *methods*:

makeNoise() – behavior for when the animal is supposed to make noise.

eat() – behavior for when the animal encounters its preferred food source, *meat* or *grass*.

sleep() – behavior for when the animal is considered asleep.

roam() – behavior for when the animal is not eating or sleeping (probably just wandering around waiting to bump into a food source or a boundary).

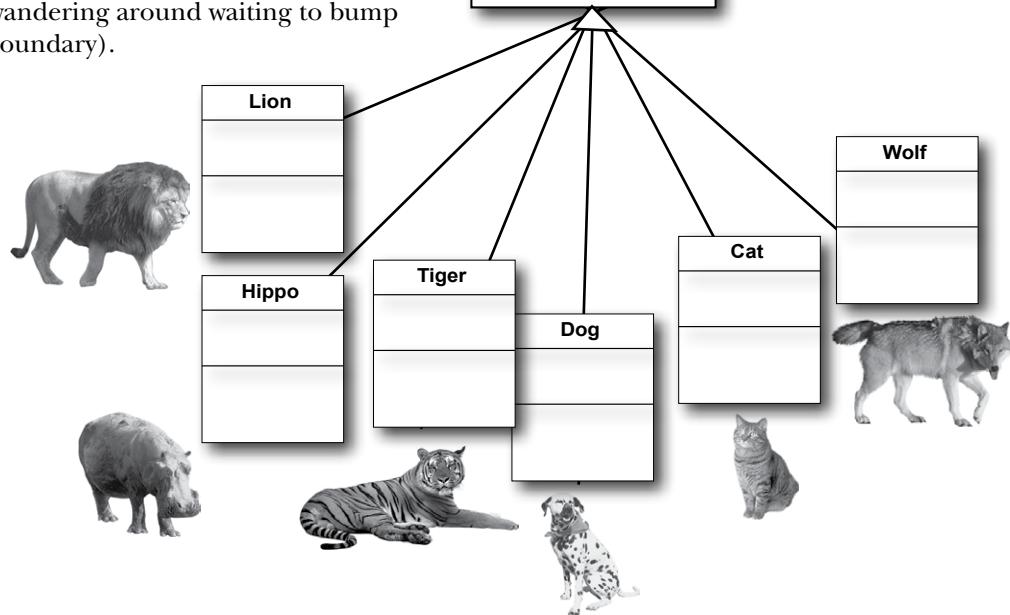
2

Design a class that represents the common state and behavior.

These objects are all animals, so we'll make a common superclass called *Animal*.

We'll put in methods and instance variables that all animals might need.

Animal
picture food hunger boundaries location
makeNoise() eat() sleep() roam()



designing for inheritance

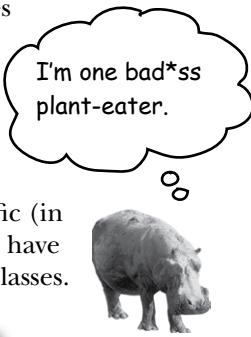
Do all animals eat the same way?

Assume that we all agree on one thing: the instance variables will work for *all* Animal types. A lion will have his own value for picture, food (we're thinking *meat*), hunger, boundaries, and location. A hippo will have different *values* for his instance variables, but he'll still have the same variables that the other Animal types have. Same with dog, tiger, and so on. But what about *behavior*?

Which methods should we override?

Does a lion make the same **noise** as a dog? Does a cat **eat** like a hippo? Maybe in *your* version, but in ours, eating and making noise are Animal-type-specific. We can't figure out how to code those methods in such a way that they'd work for any animal. OK, that's not true. We could write the `makeNoise()` method, for example, so that all it does is play a sound file defined in an instance variable for that type, but that's not very specialized. Some animals might make different noises for different situations (like one for eating, and another when bumping into an enemy, etc.)

So just as with the Amoeba overriding the Shape class `rotate()` method, to get more amoeba-specific (in other words, *unique*) behavior, we'll have to do the same for our Animal subclasses.



3

Decide if a subclass needs behaviors (method implementations) that are specific to that particular subclass type.

Looking at the Animal class, we decide that `eat()` and `makeNoise()` should be overridden by the individual subclasses.



Animal
picture
food
hunger
boundaries
location
makeNoise()
eat()
sleep()
roam()



We better override these two methods, `eat()` and `makeNoise()`, so that each animal type can define its own specific behavior for eating and making noise. For now, it looks like `sleep()` and `roam()` can stay generic.

Looking for more inheritance opportunities

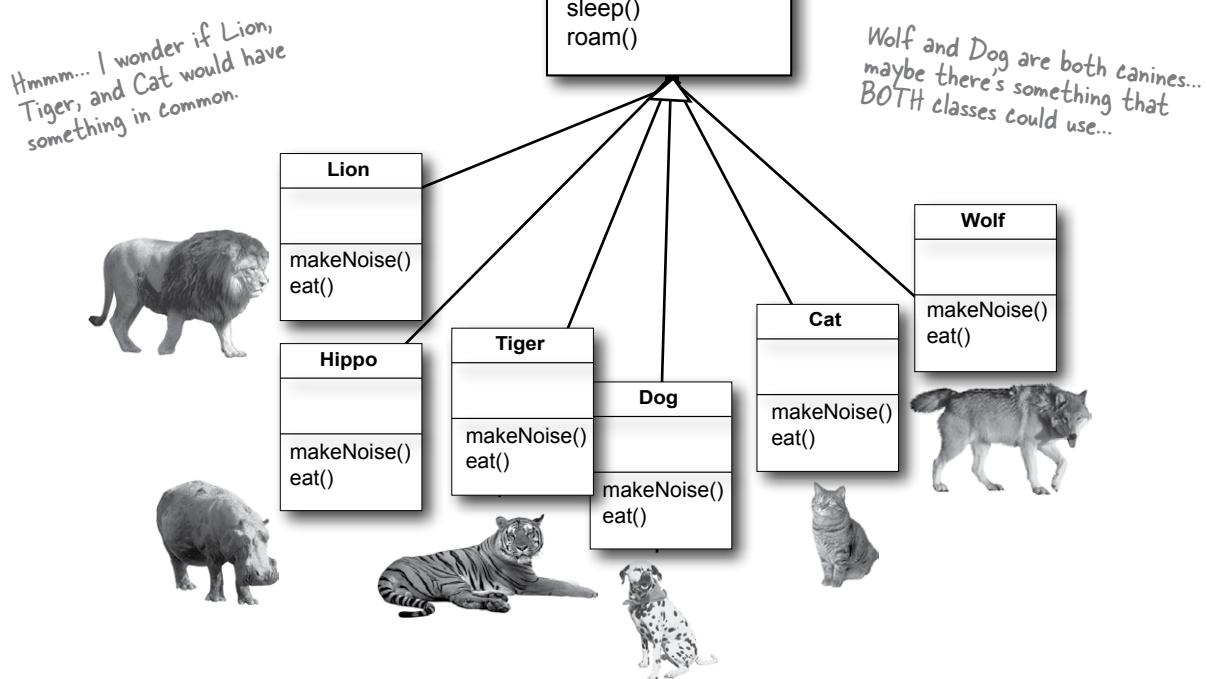
The class hierarchy is starting to shape up. We have each subclass override the `makeNoise()` and `eat()` methods, so that there's no mistaking a Dog bark from a Cat meow (quite insulting to both parties). And a Hippo won't eat like a Lion.

But perhaps there's more we can do. We have to look at the subclasses of `Animal`, and see if two or more can be grouped together in some way, and given code that's common to only *that* new group. Wolf and Dog have similarities. So do Lion, Tiger, and Cat.

4

Look for more opportunities to use abstraction, by finding two or more subclasses that might need common behavior.

We look at our classes and see that Wolf and Dog might have some behavior in common, and the same goes for Lion, Tiger, and Cat.



designing for inheritance

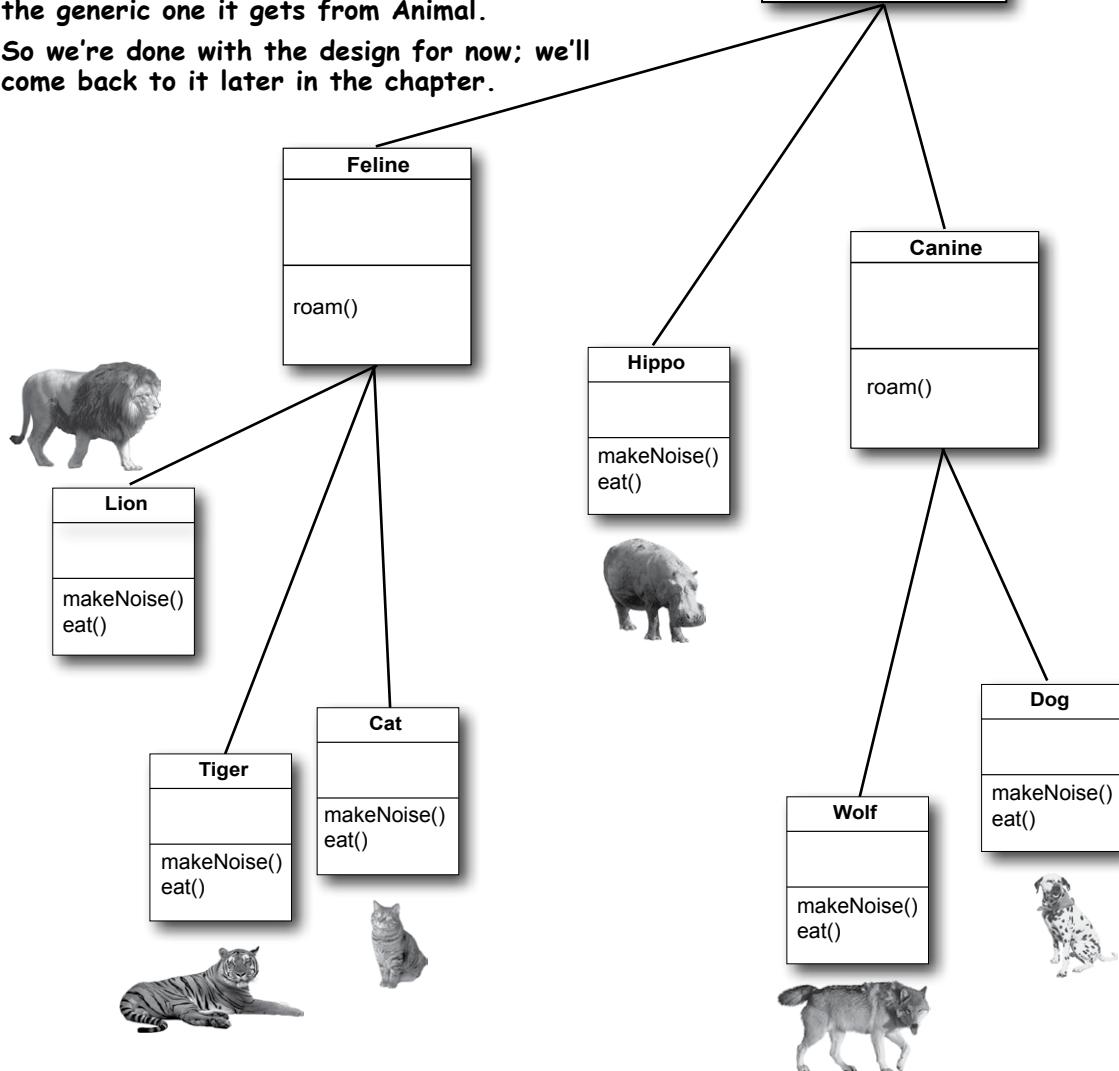
5 Finish the class hierarchy

Since animals already have an organizational hierarchy (the whole kingdom, genus, phylum thing), we can use the level that makes the most sense for class design. We'll use the biological "families" to organize the animals by making a Feline class and a Canine class.

We decide that Canines could use a common `roam()` method, because they tend to move in packs. We also see that Felines could use a common `roam()` method, because they tend to avoid others of their own kind. We'll let Hippo continue to use its inherited `roam()` method—the generic one it gets from `Animal`.

So we're done with the design for now; we'll come back to it later in the chapter.

Animal
picture food hunger boundaries location
makeNoise() eat() sleep() roam()



Which method is called?

The Wolf class has four methods. One inherited from Animal, one inherited from Canine (which is actually an overridden version of a method in class Animal), and two overridden in the Wolf class. When you create a Wolf object and assign it to a variable, you can use the dot operator on that reference variable to invoke all four methods. But which *version* of those methods gets called?

make a new Wolf object

```
Wolf w = new Wolf();
```

calls the version in Wolf

```
w.makeNoise();
```

calls the version in Canine

```
w.roam();
```

calls the version in Wolf

```
w.eat();
```

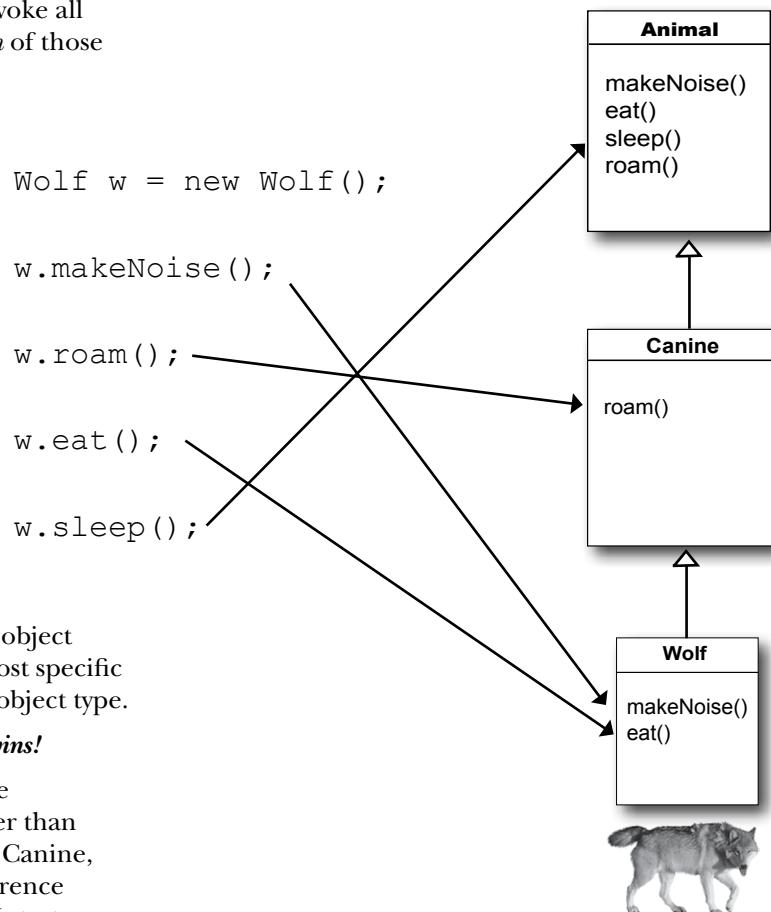
calls the version in Animal

```
w.sleep();
```

When you call a method on an object reference, you're calling the most specific version of the method for that object type.

In other words, ***the lowest one wins!***

“Lowest” meaning lowest on the inheritance tree. Canine is lower than Animal, and Wolf is lower than Canine, so invoking a method on a reference to a Wolf object means the JVM starts looking first in the Wolf class. If the JVM doesn't find a version of the method in the Wolf class, it starts walking back up the inheritance hierarchy until it finds a match.



practice designing an inheritance tree

Designing an Inheritance Tree

Class	Superclasses	Subclasses
Clothing	---	Boxers, Shirt
Boxers	Clothing	
Shirt	Clothing	

Inheritance Table



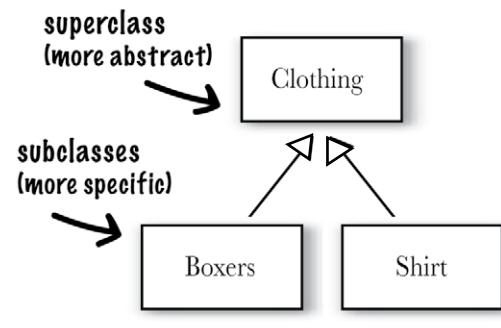
Sharpen your pencil

Find the relationships that make sense. Fill in the last two columns

Class	Superclasses	Subclasses
Musician		
Rock Star		
Fan		
Bass Player		
Concert Pianist		

Hint: not everything can be connected to something else.

Hint: you're allowed to add to or change the classes listed.



Draw an inheritance diagram here.

Dumb Questions

Q: You said that the JVM starts walking up the inheritance tree, starting at the class type you invoked the method on (like the Wolf example on the previous page). But what happens if the JVM doesn't ever find a match?

A: Good question! But you don't have to worry about that. The compiler guarantees that a particular method is callable for a specific reference type, but it doesn't say (or care) from which class that method actually comes from at runtime. With the Wolf example, the compiler checks for a sleep() method, but doesn't care that sleep() is actually defined in (and inherited from) class Animal. Remember that if a class inherits a method, it has the method.

Where the inherited method is defined (in other words, in which superclass it is defined) makes no difference to the compiler. But at runtime, **the JVM will always pick the right one**. And the right one means, **the most specific version for that particular object**.

Using IS-A and HAS-A

Remember that when one class inherits from another, we say that the subclass *extends* the superclass. When you want to know if one thing should extend another, apply the IS-A test.

Triangle IS-A Shape, yeah, that works.

Cat IS-A Feline, that works too.

Surgeon IS-A Doctor, still good.

Tub extends Bathroom, sounds reasonable.

Until you apply the IS-A test.

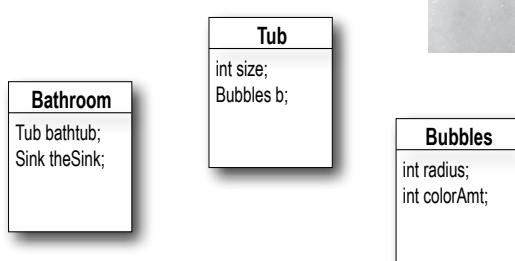
To know if you've designed your types correctly, ask, "Does it make sense to say type X IS-A type Y?" If it doesn't, you know there's something wrong with the design, so if we apply the IS-A test, Tub IS-A Bathroom is definitely false.

What if we reverse it to Bathroom extends Tub? That still doesn't work, Bathroom IS-A Tub doesn't work.

Tub and Bathroom *are* related, but not through inheritance. Tub and Bathroom are joined by a HAS-A relationship. Does it make sense to say "Bathroom HAS-A Tub"? If yes, then it means that Bathroom has a Tub instance variable. In other words, Bathroom has a *reference* to a Tub, but Bathroom does not *extend* Tub and vice-versa.



Does it make sense to say a Tub IS-A Bathroom? Or a Bathroom IS-A Tub? Well it doesn't to me. The relationship between my Tub and my Bathroom is HAS-A. Bathroom HAS-A Tub. That means Bathroom has a Tub instance variable.



Bathroom HAS-A Tub and Tub HAS-A Bubbles.
But nobody inherits from (extends) anybody else.

exploiting the power of objects

But wait! There's more!

The IS-A test works *anywhere* in the inheritance tree. If your inheritance tree is well-designed, the IS-A test should make sense when you ask *any* subclass if it IS-A *any* of its supertypes.

If class B extends class A, class B IS-A class A.

This is true anywhere in the inheritance tree. If class C extends class B, class C passes the IS-A test for both B and A.

Canine extends Animal

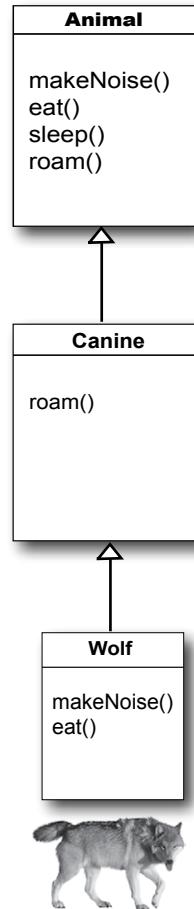
Wolf extends Canine

Wolf extends Animal

Canine IS-A Animal

Wolf IS-A Canine

Wolf IS-A Animal



With an inheritance tree like the one shown here, you're *always* allowed to say "**Wolf extends Animal**" or "**Wolf IS-A Animal**". It makes no difference if Animal is the superclass of the superclass of Wolf. In fact, **as long as Animal is somewhere in the inheritance hierarchy above Wolf, Wolf IS-A Animal will always be true**.

The structure of the Animal inheritance tree says to the world:

"Wolf IS-A Canine, so Wolf can do anything a Canine can do. And Wolf IS-A Animal, so Wolf can do anything an Animal can do."

It makes no difference if Wolf overrides some of the methods in Animal or Canine. As far as the world (of other code) is concerned, a Wolf can do those four methods. *How* he does them, or *in which class they're overridden* makes no difference. A Wolf can makeNoise(), eat(), sleep(), and roam() because a Wolf extends from class Animal.

How do you know if you've got your inheritance right?

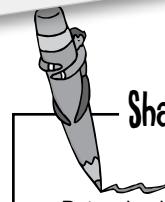
There's obviously more to it than what we've covered so far, but we'll look at a lot more OO issues in the next chapter (where we eventually refine and improve on some of the design work we did in *this* chapter).

For now, though, a good guideline is to use the IS-A test. If "X IS-A Y" makes sense, both classes (X and Y) should probably live in the same inheritance hierarchy. Chances are, they have the same or overlapping behaviors.

Keep in mind that the inheritance IS-A relationship works in only one direction!

Triangle IS-A Shape makes sense, so you can have Triangle extend Shape.

But the reverse—Shape IS-A Triangle—does *not* make sense, so Shape should not extend Triangle. Remember that the IS-A relationship implies that if X IS-A Y, then X can do anything a Y can do (and possibly more).



Sharpen your pencil

Put a check next to the relationships that make sense.

- Oven extends Kitchen
- Guitar extends Instrument
- Person extends Employee
- Ferrari extends Engine
- FriedEgg extends Food
- Beagle extends Pet
- Container extends Jar
- Metal extends Titanium
- GratefulDead extends Band
- Blonde extends Smart
- Beverage extends Martini

Hint: apply the IS-A test

who inherits what

there are no
Dumb Questions

Q: So we see how a subclass gets to inherit a superclass method, but what if the superclass wants to use the subclass version of the method?

A: A superclass won't necessarily know about any of its subclasses. You might write a class and much later someone else comes along and extends it. But even if the superclass creator does know about (and wants to use) a subclass version of a method, there's no sort of *reverse or backwards* inheritance. Think about it, children inherit from parents, not the other way around.

Q: In a subclass, what if I want to use BOTH the superclass version and my overriding subclass version of a method? In other words, I don't want to completely replace the superclass version, I just want to add more stuff to it.

A: You can do this! And it's an important design feature. Think of the word "extends" as meaning, "I want to extend the functionality of the superclass".

```
public void roam() {  
    super.roam();  
    // my own roam stuff  
}
```

You can design your superclass methods in such a way that they contain method implementations that will work for any subclass, even though the subclasses may still need to 'append' more code. In your subclass overriding method, you can call the superclass version using the keyword **super**. It's like saying, "first go run the superclass version, then come back and finish with my own code..."

this calls the inherited version of
roam(), then comes back to do
your own subclass-specific code

Who gets the Porsche, who gets the porcelain? (how to know what a subclass can inherit from its superclass)



A subclass inherits members of the superclass. Members include instance variables and methods, although later in this book we'll look at other inherited members. A superclass can choose whether or not it wants a subclass to inherit a particular member by the level of access the particular member is given.

There are four access levels that we'll cover in this book. Moving from most restrictive to least, the four access levels are:

private default protected public

Access levels control *who sees what*, and are crucial to having well-designed, robust Java code. For now we'll focus just on public and private. The rules are simple for those two:

public members are inherited
private members are not inherited

When a subclass inherits a member, it is *as if the subclass defined the member itself*. In the Shape example, Square inherited the `rotate()` and `playSound()` methods and to the outside world (other code) the Square class simply has a `rotate()` and `playSound()` method.

The members of a class include the variables and methods defined in the class plus anything inherited from a superclass.

Note: get more details about default and protected in chapter 17 (deployment) and appendix B.

When designing with inheritance, are you using or abusing?

Although some of the reasons behind these rules won't be revealed until later in this book, for now, simply *knowing* a few rules will help you build a better inheritance design.

DO use inheritance when one class is a more specific type of a superclass. Example: Willow *is a* more specific type of Tree, so Willow *extends* Tree makes sense.

DO consider inheritance when you have behavior (implemented code) that should be shared among multiple classes of the same general type. Example: Square, Circle, and Triangle all need to rotate and play sound, so putting that functionality in a superclass Shape might make sense, and makes for easier maintenance and extensibility. Be aware, however, that while inheritance is one of the key features of object-oriented programming, it's not necessarily the best way to achieve behavior reuse. It'll get you started, and often it's the right design choice, but design patterns will help you see other more subtle and flexible options. If you don't know about design patterns, a good follow-on to this book would be *Head First Design Patterns*.

DO NOT use inheritance just so that you can reuse code from another class, if the relationship between the superclass and subclass violate either of the above two rules. For example, imagine you wrote special printing code in the Alarm class and now you need printing code in the Piano class, so you have Piano extend Alarm so that Piano inherits the printing code. That makes no sense! A Piano is *not* a more specific type of Alarm. (So the printing code should be in a Printer class, that all printable objects can take advantage of via a HAS-A relationship.)

DO NOT use inheritance if the subclass and superclass do not pass the IS-A test. Always ask yourself if the subclass IS-A more specific type of the superclass. Example: Tea IS-A Beverage makes sense. Beverage IS-A Tea does not.



BULLET POINTS

- A subclass *extends* a superclass.
- A subclass inherits all *public* instance variables and methods of the superclass, but does not inherit the *private* instance variables and methods of the superclass.
- Inherited methods *can* be overridden; instance variables *cannot* be overridden (although they *can* be *redefined* in the subclass, but that's not the same thing, and there's almost never a need to do it.)
- Use the IS-A test to verify that your inheritance hierarchy is valid. If X *extends* Y, then X IS-A Y must make sense.
- The IS-A relationship works in only one direction. A Hippo is an Animal, but not all Animals are Hippos.
- When a method is overridden in a subclass, and that method is invoked on an instance of the subclass, the overridden version of the method is called. (*The lowest one wins.*)
- If class B extends A, and C extends B, class B IS-A class A, and class C IS-A class B, and class C also IS-A class A.

So what does all this inheritance really buy you?

You get a lot of OO mileage by designing with inheritance. You can get rid of duplicate code by abstracting out the behavior common to a group of classes, and sticking that code in a superclass. That way, when you need to modify it, you have only one place to update, and *the change is magically reflected in all the classes that inherit that behavior*. Well, there's no magic involved, but it *is* pretty simple: make the change and compile the class again. That's it. **You don't have to touch the subclasses!**

Just deliver the newly-changed superclass, and all classes that extend it will automatically use the new version.

A Java program is nothing but a pile of classes, so the subclasses don't have to be recompiled in order to use the new version of the superclass. As long as the superclass doesn't *break* anything for the subclass, everything's fine. (We'll discuss what the word 'break' means in this context, later in the book. For now, think of it as modifying something in the superclass that the subclass is depending on, like a particular method's arguments or return type, or method name, etc.)

① You avoid duplicate code.

Put common code in one place, and let the subclasses inherit that code from a superclass. When you want to change that behavior, you have to modify it in only one place, and everybody else (i.e. all the subclasses) see the change.

② You define a common protocol for a group of classes.



Inheritance lets you guarantee that all classes grouped under a certain supertype have all the methods that the supertype has.*

In other words, you define a common protocol for a set of classes related through inheritance.

When you define methods in a superclass, that can be inherited by subclasses, you're announcing a kind of protocol to other code that says, "All my subtypes (i.e. subclasses) can do these things, with these methods that look like this..."

In other words, you establish a *contract*.

Class Animal establishes a common protocol for all Animal subtypes:

Animal
makeNoise() eat() sleep() roam()

You're telling the world that any Animal can do these four things. That includes the method arguments and return types.

And remember, when we say *any Animal*, we mean Animal *and any class that extends from Animal*. Which again means, *any class that has Animal somewhere above it in the inheritance hierarchy*.

But we're not even at the really cool part yet, because we saved the best—*polymorphism*—for last.

When you define a supertype for a group of classes, *any subclass of that supertype can be substituted where the supertype is expected*.

Say, what?

Don't worry, we're nowhere near done explaining it. Two pages from now, you'll be an expert.

*When we say "all the methods" we mean "all the *inheritable* methods", which for now actually means, "all the *public* methods", although later we'll refine that definition a bit more.

And I care because...

Because you get to take advantage of polymorphism.

Which matters to me because...

Because you get to refer to a subclass object using a reference declared as the supertype.

And that means to me...

You get to write really flexible code. Code that's cleaner (more efficient, simpler). Code that's not just easier to develop, but also much, much easier to extend, in ways you never imagined at the time you originally wrote your code.

That means you can take that tropical vacation while your co-workers update the program, and your co-workers might not even need your source code.

You'll see how it works on the next page.

We don't know about you, but personally, we find the whole tropical vacation thing particularly motivating.



the way polymorphism works

To see how polymorphism works, we have to step back and look at the way we normally declare a reference and create an object...

The 3 steps of object declaration and assignment

1 2
Dog myDog = new Dog();
 3

- 1 Declare a reference variable

Dog myDog = new Dog();

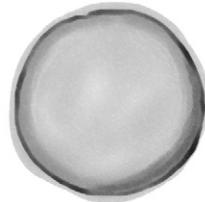
Tells the JVM to allocate space for a reference variable. The reference variable is, forever, of type Dog. In other words, a remote control that has buttons to control a Dog, but not a Cat or a Button or a Socket.



- 2 Create an object

Dog myDog = new Dog();

Tells the JVM to allocate space for a new Dog object on the garbage collectible heap.

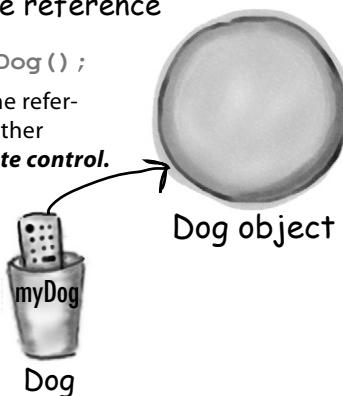


Dog object

- 3 Link the object and the reference

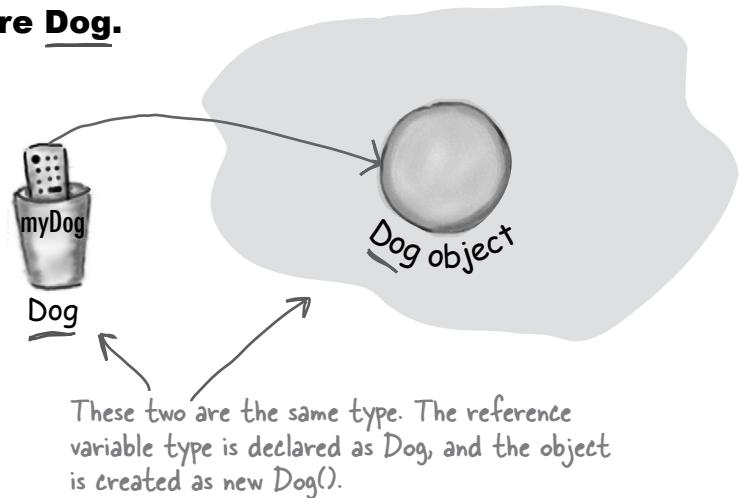
Dog myDog = new Dog();

Assigns the new Dog to the reference variable myDog. In other words, **program the remote control**.



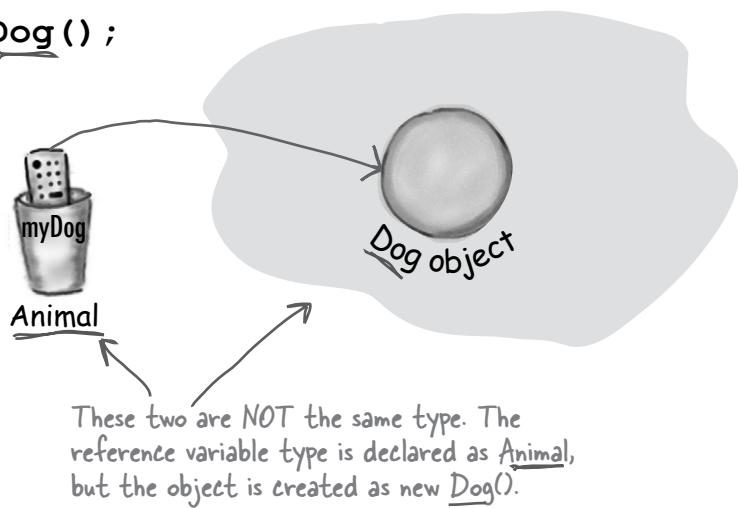
The important point is that the reference type AND the object type are the same.

In this example, both are Dog.



But with polymorphism, the reference and the object can be different.

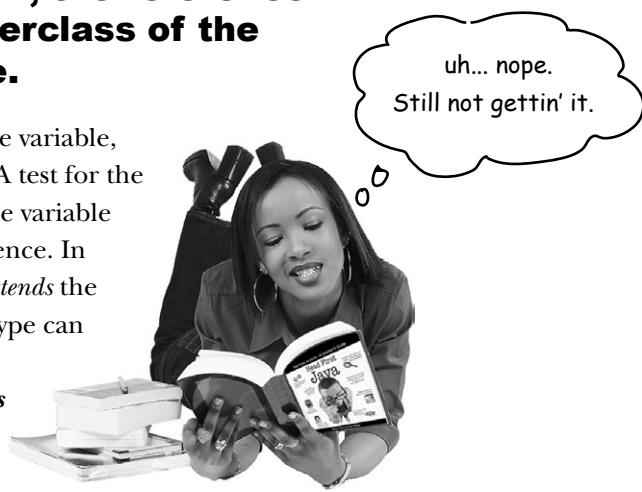
Animal myDog = new Dog();



polymorphism in action

With polymorphism, the reference type can be a superclass of the actual object type.

When you declare a reference variable, any object that passes the IS-A test for the declared type of the reference variable can be assigned to that reference. In other words, anything that *extends* the declared reference variable type can be *assigned* to the reference variable. *This lets you do things like make polymorphic arrays.*



OK, OK maybe an example will help.

```
Animal[] animals = new Animal[5];  
  
animals [0] = new Dog();  
animals [1] = new Cat();  
animals [2] = new Wolf();  
animals [3] = new Hippo();  
animals [4] = new Lion();  
  
for (int i = 0; i < animals.length; i++) {  
  
    animals[i].eat();  
    animals[i].roam();  
}
```

Declare an array of type Animal. In other words, an array that will hold objects of type Animal.

But look what you get to do... you can put ANY subclass of Animal in the Animal array!

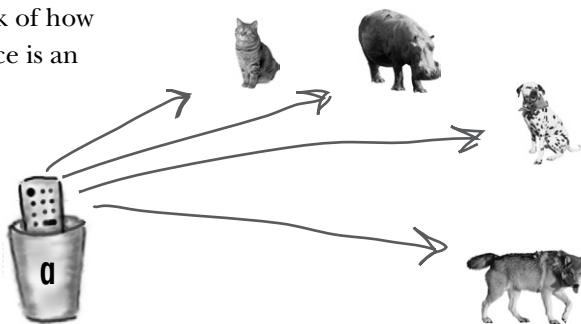
And here's the best polymorphic part (the raison d'être for the whole example), you get to loop through the array and call one of the Animal-class methods, and every object does the right thing!

When 'i' is 0, a Dog is at index 0 in the array, so you get the Dog's eat() method. When 'i' is 1, you get the Cat's eat() method
Same with roam().

But wait! There's more!

You can have polymorphic arguments and return types.

If you can declare a reference variable of a supertype, say, Animal, and assign a subclass object to it, say, Dog, think of how that might work when the reference is an argument to a method...



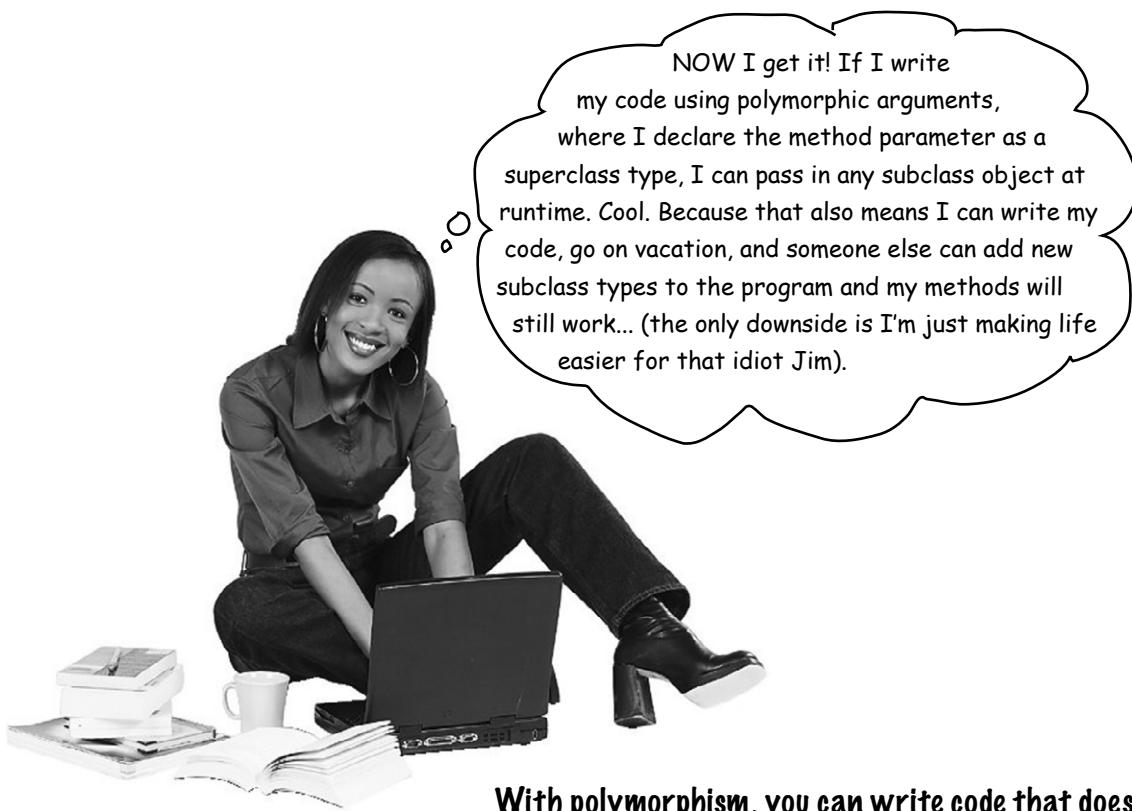
```
class Vet {
    public void giveShot(Animal a) {
        // do horrible things to the Animal at
        // the other end of the 'a' parameter
        a.makeNoise();
    }
}
```

The 'a' parameter can take ANY Animal type as the argument. And when the Vet is done giving the shot, it tells the Animal to makeNoise(), and whatever Animal is really out there on the heap, that's whose makeNoise() method will run.

```
class PetOwner {
    public void start() {
        Vet v = new Vet();
        Dog d = new Dog();
        Hippo h = new Hippo();
        v.giveShot(d);           ← Dog's makeNoise() runs
        v.giveShot(h);          ← Hippo's makeNoise() runs
    }
}
```

The Vet's giveShot() method can take any Animal you give it. As long as the object you pass in as the argument is a subclass of Animal, it will work.

exploiting the power of polymorphism



NOW I get it! If I write my code using polymorphic arguments, where I declare the method parameter as a superclass type, I can pass in any subclass object at runtime. Cool. Because that also means I can write my code, go on vacation, and someone else can add new subclass types to the program and my methods will still work... (the only downside is I'm just making life easier for that idiot Jim).

With polymorphism, you can write code that doesn't have to change when you introduce new subclass types into the program.

Remember that Vet class? If you write that Vet class using arguments declared as type *Animal*, your code can handle any *Animal subclass*. That means if others want to take advantage of your Vet class, all they have to do is make sure *their* new Animal types extend class *Animal*. The Vet methods will still work, even though the Vet class was written without any knowledge of the new Animal subtypes the Vet will be working on.



Why is polymorphism guaranteed to work this way? Why is it always safe to assume that any *subclass* type will have the methods you think you're calling on the *superclass* type (the superclass reference type you're using the dot operator on)?

there are no
Dumb Questions

Q: Are there any practical limits on the levels of subclassing? How deep can you go?

A: If you look in the Java API, you'll see that most inheritance hierarchies are wide but not deep. Most are no more than one or two levels deep, although there are exceptions (especially in the GUI classes). You'll come to realize that it usually makes more sense to keep your inheritance trees shallow, but there isn't a hard limit (well, not one that you'd ever run into).

Q: Hey, I just thought of something... if you don't have access to the source code for a class, but you want to change the way a method of that class works, could you use subclassing to do that? To extend the "bad" class and override the method with your own better code?

A: Yep. That's one cool feature of OO, and sometimes it saves you from having to rewrite the class from scratch, or track down the programmer who hid the source code.

Q: Can you extend *any* class? Or is it like class members where if the class is private you can't inherit it...

A: There's no such thing as a private class, except in a very special case called an *inner* class, that we haven't looked at yet. But there are three things that can prevent a class from being subclassed.

The first is access control. Even though a class *can't* be marked `private`, a class *can* be non-public (what you get if you don't declare the class as `public`). A non-public class can be subclassed only by classes in the same package as the class. Classes in a different package won't be able to subclass (or even use, for that matter) the non-public class.

The second thing that stops a class from being subclassed is the keyword modifier `final`. A final class means that it's the end of the inheritance line. Nobody, ever, can extend a final class.

The third issue is that if a class has only `private` constructors (we'll look at constructors in chapter 9), it can't be subclassed.

Q: Why would you ever want to make a final class? What advantage would there be in preventing a class from being subclassed?

A: Typically, you won't make your classes final. But if you need security — the security of knowing that the methods will always work the way that you wrote them (because they can't be overridden), a final class will give you that. A lot of classes in the Java API are final for that reason. The `String` class, for example, is final because, well, imagine the havoc if somebody came along and changed the way `Strings` behave!

Q: Can you make a *method* final, without making the whole *class* final?

A: If you want to protect a specific method from being overridden, mark the *method* with the `final` modifier. Mark the whole *class* as final if you want to guarantee that *none* of the methods in that class will ever be overridden.

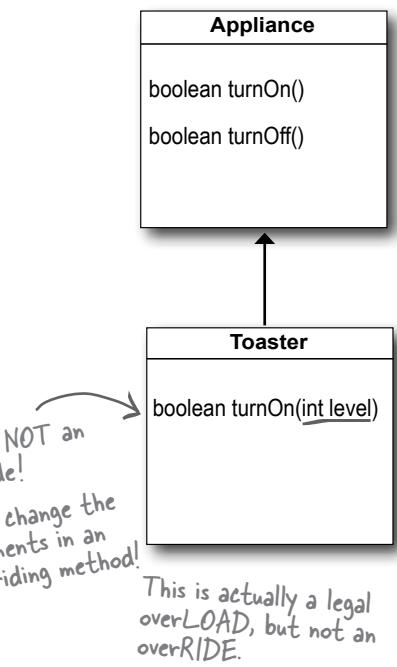
overriding methods

Keeping the contract: rules for overriding

When you override a method from a superclass, you're agreeing to fulfill the contract. The contract that says, for example, "I take no arguments and I return a boolean." In other words, the arguments and return types of your overriding method must look to the outside world *exactly* like the overridden method in the superclass.

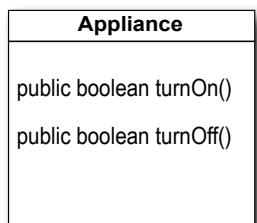
The methods *are* the contract.

If polymorphism is going to work, the Toaster's version of the overridden method from Appliance has to work at runtime. Remember, the compiler looks at the reference type to decide whether you can call a particular method on that reference. With an Appliance reference to a Toaster, the compiler cares only if class *Appliance* has the method you're invoking on an Appliance reference. But at runtime, the JVM looks not at the *reference* type (*Appliance*) but at the actual *Toaster* object on the heap. So if the compiler has already *approved* the method call, the only way it can work is if the overriding method has the same arguments and return types. Otherwise, someone with an Appliance reference will call *turnOn()* as a no-arg method, even though there's a version in Toaster that takes an int. Which one is called at runtime? The one in Appliance. In other words, *the turnOn(int level) method in Toaster is not an override!*



① Arguments must be the same, and return types must be compatible.

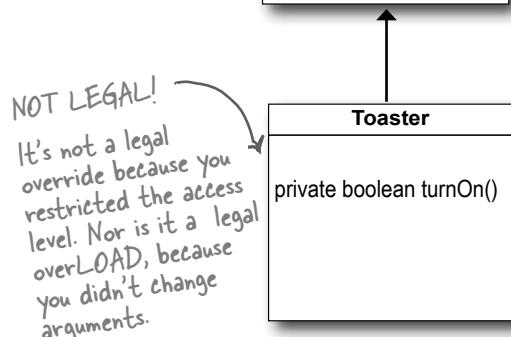
The contract of superclass defines how other code can use a method. Whatever the superclass takes as an argument, the subclass overriding the method must use that same argument. And whatever the superclass declares as a return type, the overriding method must declare either the same type, or a subclass type. Remember, a subclass object is guaranteed to be able to do anything its superclass declares, so it's safe to return a subclass where the superclass is expected.



② The method can't be less accessible.

That means the access level must be the same, or friendlier. That means you can't, for example, override a public method and make it private. What a shock that would be to the code invoking what it *thinks* (at compile time) is a public method, if suddenly at runtime the JVM slammed the door shut because the overriding version called at runtime is private!

So far we've learned about two access levels: private and public. The other two are in the deployment chapter (Release your Code) and appendix B. There's also another rule about overriding related to exception handling, but we'll wait until the chapter on exceptions (Risky Behavior) to cover that.



Overloading a method

Method overloading is nothing more than having two methods with the same name but different argument lists. Period. There's no polymorphism involved with overloaded methods!

Overloading lets you make multiple versions of a method, with different argument lists, for convenience to the callers. For example, if you have a method that takes only an int, the calling code has to convert, say, a double into an int before calling your method. But if you overloaded the method with another version that takes a double, then you've made things easier for the caller. You'll see more of this when we look into constructors in the object lifecycle chapter.

Since an overloading method isn't trying to fulfill the polymorphism contract defined by its superclass, overloaded methods have much more flexibility.

① The return types can be different.

You're free to change the return types in overloaded methods, as long as the argument lists are different.

② You can't change ONLY the return type.

If only the return type is different, it's not a valid *overload*—the compiler will assume you're trying to *override* the method. And even *that* won't be legal unless the return type is a subtype of the return type declared in the superclass. To overload a method, you **MUST** change the argument list, although you *can* change the return type to anything.

③ You can vary the access levels in any direction.

You're free to overload a method with a method that's more restrictive. It doesn't matter, since the new method isn't obligated to fulfill the contract of the overloaded method.

An overloaded method is just a different method that happens to have the same method name. It has nothing to do with inheritance and polymorphism. An overloaded method is NOT the same as an overridden method.

Legal examples of method overloading:

```
public class Overloads {
    String uniqueID;

    public int addNums(int a, int b) {
        return a + b;
    }

    public double addNums(double a, double b) {
        return a + b;
    }

    public void setUniqueID(String theID) {
        // lots of validation code, and then:
        uniqueID = theID;
    }

    public void setUniqueID(int ssNumber) {
        String numString = "" + ssNumber;
        setUniqueID(numString);
    }
}
```

exercise: Mixed Messages



Mixed Messages

```
a = 6; → 56
b = 5; → 11
a = 5; → 65
```

A short Java program is listed below. One block of the program is missing! Your challenge is to match the candidate block of code (on the left), with the output that you'd see if the block were inserted. Not all the lines of output will be used, and some of the lines of output might be used more than once. Draw lines connecting the candidate blocks of code with their matching command-line output.

the program:

```
class A {
    int ivar = 7;
    void m1() {
        System.out.print("A's m1, ");
    }
    void m2() {
        System.out.print("A's m2, ");
    }
    void m3() {
        System.out.print("A's m3, ");
    }
}

class B extends A {
    void m1() {
        System.out.print("B's m1, ");
    }
}
```

```
class C extends B {
    void m3() {
        System.out.print("C's m3, "+(ivar + 6));
    }
}

public class Mixed2 {
    public static void main(String [] args) {
        A a = new A();
        B b = new B();
        C c = new C();
        A a2 = new C();
    }
}
```

candidate code
goes here
(three lines)

**code
candidates:**

```
b.m1();
c.m2();
a.m3(); }
```

```
c.m1();
c.m2();
c.m3(); }
```

```
a.m1();
b.m2();
c.m3(); }
```

```
a2.m1();
a2.m2();
a2.m3(); }
```

output:

A's m1, A's m2, C's m3, 6

B's m1, A's m2, A's m3,

A's m1, B's m2, A's m3,

B's m1, A's m2, C's m3, 13

B's m1, C's m2, A's m3,

B's m1, A's m2, C's m3, 6

A's m1, A's m2, C's m3, 13



BE the Compiler

Which of the A-B pairs of methods listed on the right, if inserted into the classes on the left, would compile and produce the output shown? (The A method inserted into class Monster, the B method inserted into class Vampire.)

```
public class MonsterTestDrive {
    public static void main(String [] args) {
        Monster [] ma = new Monster[3];
        ma[0] = new Vampire();
        ma[1] = new Dragon();
        ma[2] = new Monster();
        for(int x = 0; x < 3; x++) {
            ma[x].frighten(x);
        }
    }
}
```

```
class Monster {
    A
```

```
class Vampire extends Monster {
    B
```

```
class Dragon extends Monster {
    boolean frighten(int degree) {
        System.out.println("breath fire");
        return true;
    }
}
```

```
File Edit Window Help SaveYourself
% java MonsterTestDrive
a bite?
breath fire
arrrgh
```

- 1


```
boolean frighten(int d) {
    System.out.println("arrrgh");
    return true;
}
```

A

```
boolean frighten(int x) {
    System.out.println("a bite?");
    return false;
}
```

B

```
boolean frighten(int x) {
    System.out.println("arrrgh");
    return true;
}
```

A

```
int frighten(int f) {
    System.out.println("a bite?");
    return 1;
}
```

B

```
boolean frighten(int x) {
    System.out.println("arrrgh");
    return false;
}
```

A

```
boolean scare(int x) {
    System.out.println("a bite?");
    return true;
}
```

B

```
boolean frighten(int z) {
    System.out.println("arrrgh");
    return true;
}
```

A

```
boolean frighten(byte b) {
    System.out.println("a bite?");
    return true;
}
```

puzzle: Pool Puzzle



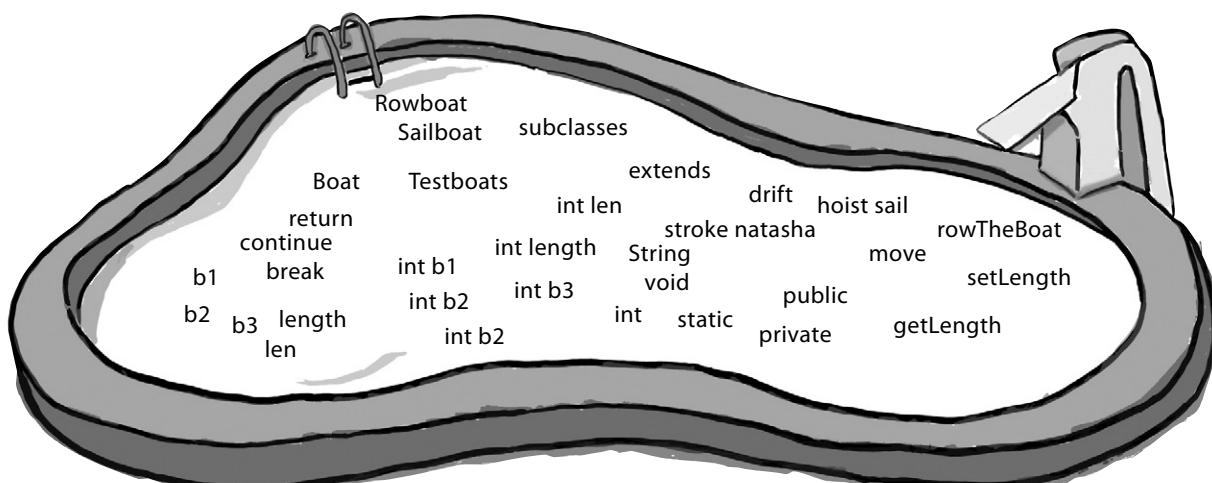
Pool Puzzle

Your **job** is to take code snippets from the pool and place them into the blank lines in the code. You may use the same snippet more than once, and you might not need to use all the snippets. Your **goal** is to make a set of classes that will compile and run together as a program. Don't be fooled – this one's harder than it looks.

```
public class Rowboat _____ {  
    public _____ rowTheBoat() {  
        System.out.print("stroke natasha");  
    }  
}  
  
public class _____ {  
    private int _____;  
    _____ void _____(_____) {  
        length = len;  
    }  
    public int getLength() {  
        _____ _____;  
    }  
    public _____ move() {  
        System.out.print("_____");  
    }  
}
```

```
public class TestBoats {  
    _____ main(String[] args) {  
        _____ b1 = new Boat();  
        Sailboat b2 = new _____();  
        Rowboat _____ = new Rowboat();  
        b2.setLength(32);  
        b1._____();  
        b3._____();  
        _____.move();  
    }  
}  
  
public class _____ Boat {  
    public _____() {  
        System.out.print("_____");  
    }  
}
```

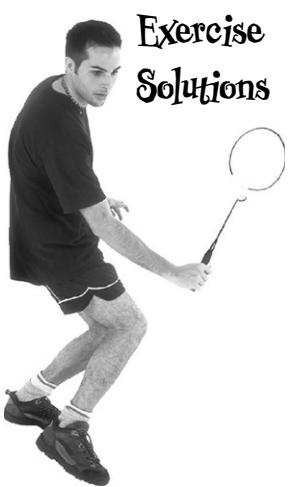
OUTPUT: drift drift hoist sail





BE the Compiler

Set 1 will work.



Set 2 will not compile because of Vampire's return type (int).

The Vampire's frighten() method (B) is not a legal override OR overload of Monster's frighten() method. Changing ONLY the return type is not enough to make a valid overload, and since an int is not compatible with a boolean, the method is not a valid override. (Remember, if you change ONLY the return type, it must be to a return type that is compatible with the superclass version's return type, and then it's an *override*.)

Sets 3 and 4 will compile, but produce:

```
arrrgh
breath fire
arrrgh
```

Remember, class Vampire did not *override* class Monster's frighten() method. (The frighten() method in Vampire's set 4 takes a byte, not an int.)

code candidates:

Mixed Messages

```
b.m1();
c.m2();
a.m3(); }
```

```
c.m1();
c.m2();
c.m3(); }
```

```
a.m1();
b.m2();
c.m3(); }
```

```
a2.m1();
a2.m2();
a2.m3(); }
```

output:

A's m1, A's m2, C's m3, 6

B's m1, A's m2, A's m3,

A's m1, B's m2, A's m3,

B's m1, A's m2, C's m3, 13

B's m1, C's m2, A's m3,

B's m1, A's m2, C's m3, 6

A's m1, A's m2, C's m3, 13

puzzle answers



```
public class Rowboat extends Boat {
    public void rowTheBoat() {
        System.out.print("stroke natasha");
    }
}

public class Boat {
    private int length;
    public void setLength( int len ) {
        length = len;
    }
    public int getLength() {
        return length;
    }
    public void move() {
        System.out.print("drift ");
    }
}
```

```
public class TestBoats {
    public static void main(String[] args) {
        Boat b1 = new Boat();
        Sailboat b2 = new Sailboat();
        Rowboat b3 = new Rowboat();
        b2.setLength(32);
        b1.move();
        b3.move();
        b2.move();
    }
}

public class Sailboat extends Boat {
    public void move() {
        System.out.print("hoist sail ");
    }
}
```

OUTPUT: drift drift hoist sail

8 interfaces and abstract classes

Serious Polymorphism

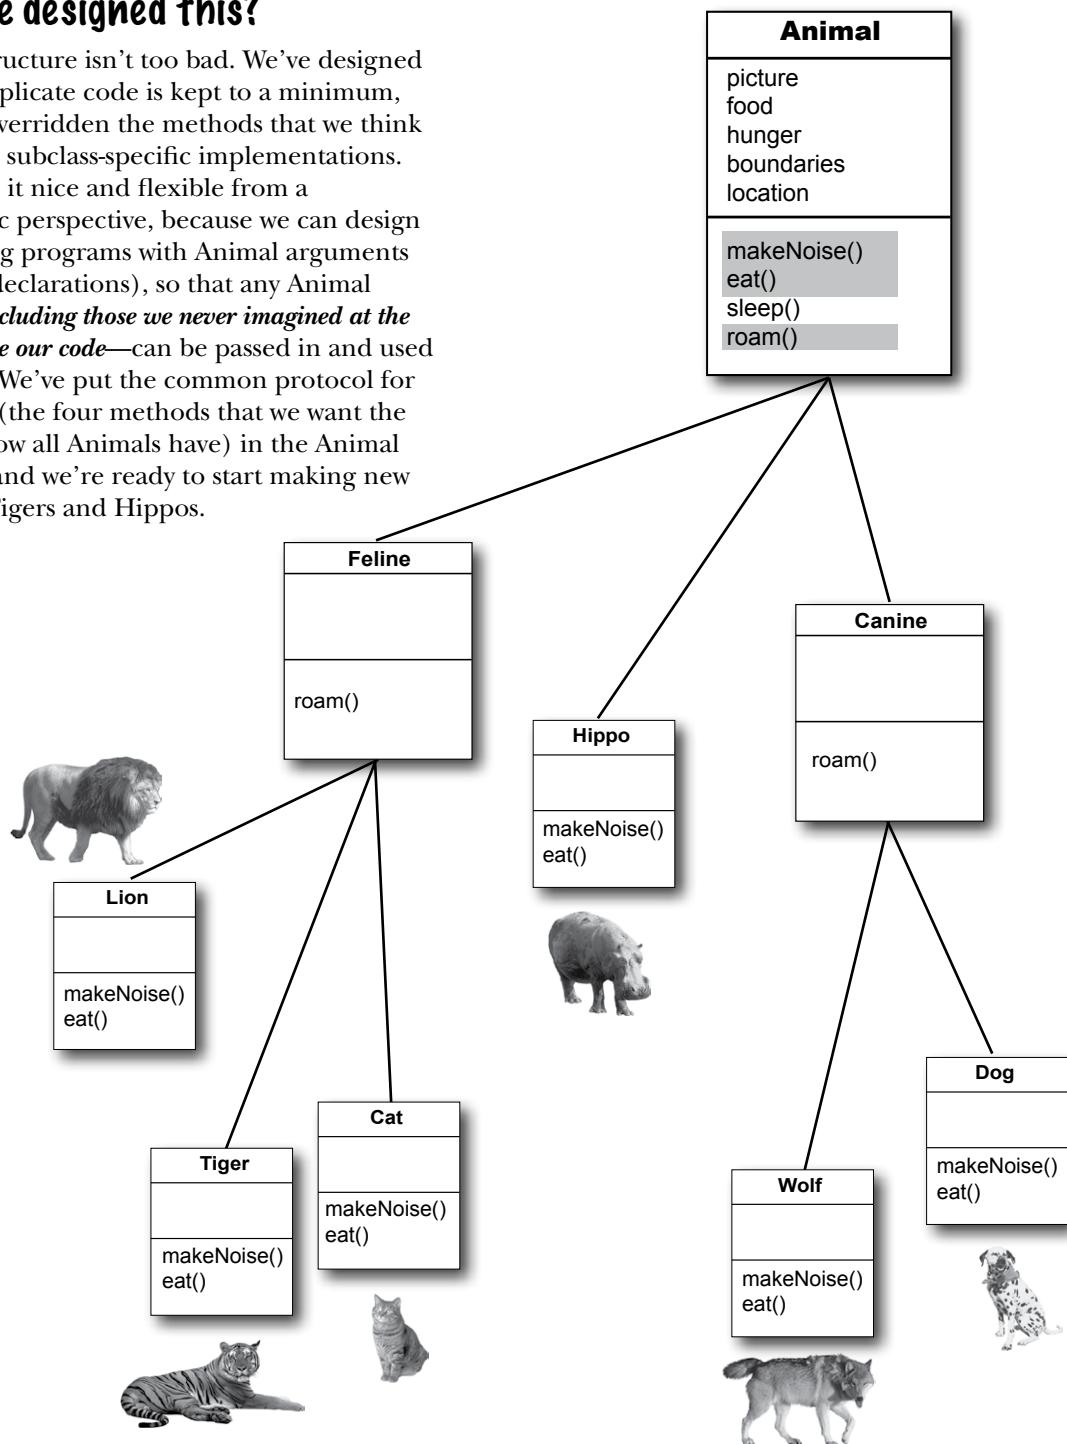


Inheritance is just the beginning. To exploit polymorphism, we need interfaces (and not the GUI kind). We need to go beyond simple inheritance to a level of flexibility and extensibility you can get only by designing and coding to interface specifications. Some of the coolest parts of Java wouldn't even be possible without interfaces, so even if you don't design with them yourself, you still have to use them. But you'll *want* to design with them. You'll *need* to design with them. **You'll wonder how you ever lived without them.** What's an interface? It's a 100% abstract class. What's an abstract class? It's a class that can't be instantiated. What's that good for? You'll see in just a few moments. But if you think about the end of the last chapter, and how we used polymorphic arguments so that a single Vet method could take Animal subclasses of all types, well, that was just scratching the surface. Interfaces are the **poly** in polymorphism. The **ab** in abstract. The **caffeine** in Java.

designing with inheritance

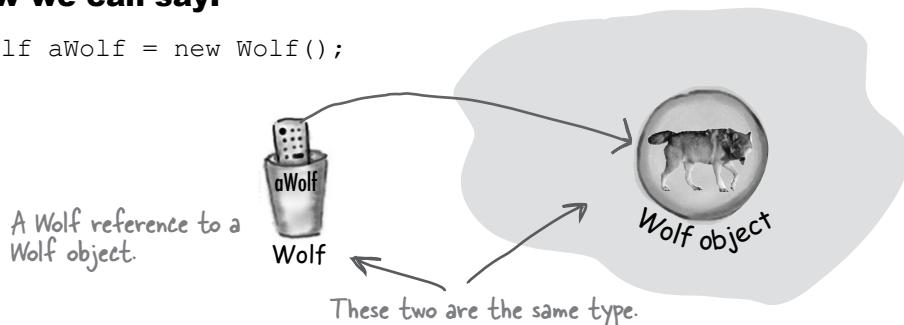
Did we forget about something when we designed this?

The class structure isn't too bad. We've designed it so that duplicate code is kept to a minimum, and we've overridden the methods that we think should have subclass-specific implementations. We've made it nice and flexible from a polymorphic perspective, because we can design Animal-using programs with Animal arguments (and array declarations), so that any Animal subtype—*including those we never imagined at the time we wrote our code*—can be passed in and used at runtime. We've put the common protocol for all Animals (the four methods that we want the world to know all Animals have) in the Animal superclass, and we're ready to start making new Lions and Tigers and Hippos.



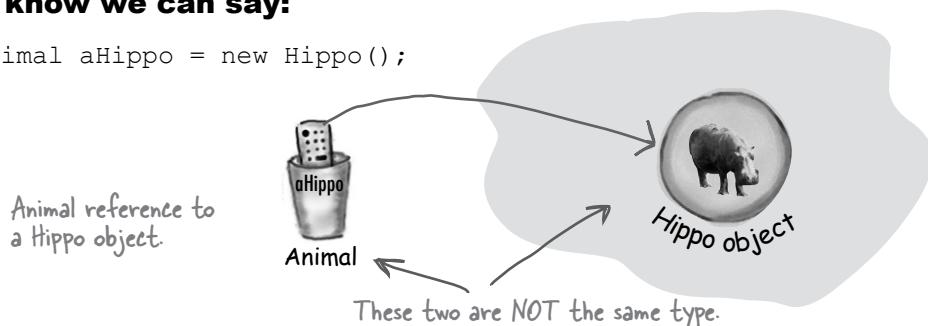
We know we can say:

```
Wolf aWolf = new Wolf();
```



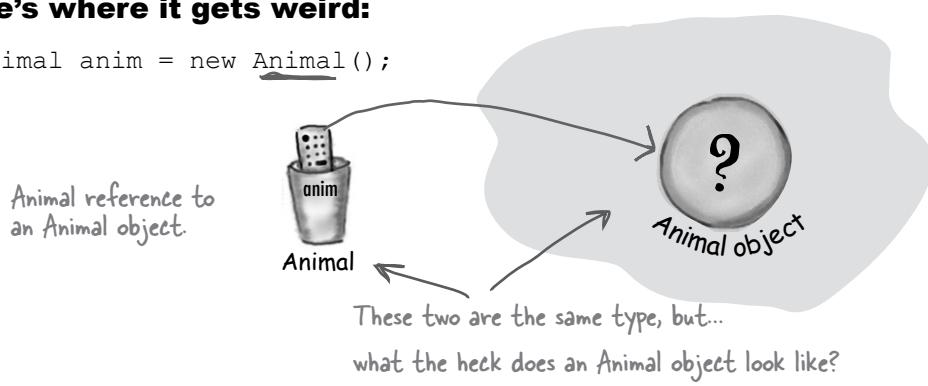
And we know we can say:

```
Animal aHippo = new Hippo();
```



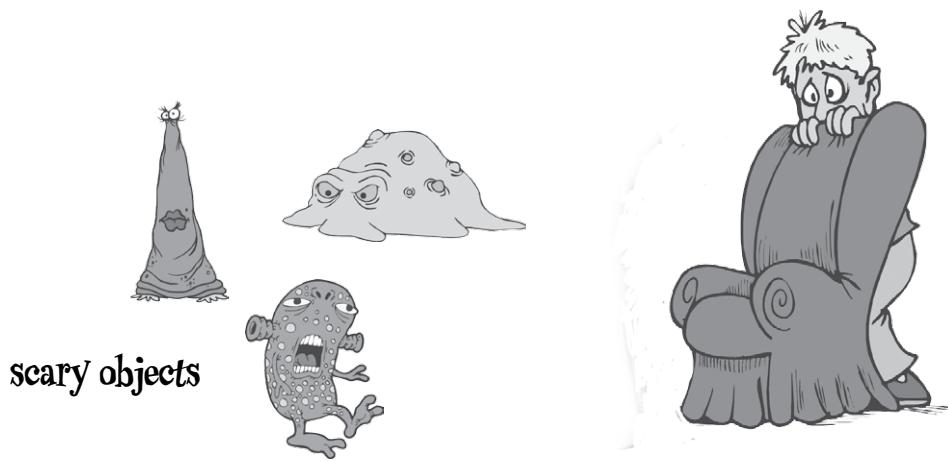
But here's where it gets weird:

```
Animal anim = new Animal();
```



when objects go bad

What does a new `Animal()` object look like?



What are the instance variable values?

Some classes just should not be instantiated!

It makes sense to create a Wolf object or a Hippo object or a Tiger object, but what exactly *is* an Animal object? What shape is it? What color, size, number of legs...

Trying to create an object of type Animal is like a **nightmare Star Trek™ transporter accident**. The one where somewhere in the beam-me-up process something bad happened to the buffer.

But how do we deal with this? We *need* an Animal class, for inheritance and polymorphism. But we want programmers to instantiate only the less abstract *subclasses* of class Animal, not Animal itself. We want Tiger objects and Lion objects, ***not* Animal objects**.

Fortunately, there's a simple way to prevent a class from ever being instantiated. In other words, to stop anyone from saying “`new`” on that type. By marking the class as **abstract**, the compiler will stop any code, anywhere, from ever creating an instance of that type.

You can still use that abstract type as a reference type. In fact, that's a big part of why you have that abstract class in the first place (to use it as a polymorphic argument or return type, or to make a polymorphic array).

When you're designing your class inheritance structure, you have to decide which classes are *abstract* and which are *concrete*. Concrete classes are those that are specific enough to be instantiated. A *concrete* class just means that it's OK to make objects of that type.

Making a class abstract is easy—put the keyword **abstract** before the class declaration:

```
abstract class Canine extends Animal {  
    public void roam() {}  
}
```

The compiler won't let you instantiate an abstract class

An abstract class means that nobody can ever make a new instance of that class. You can still use that abstract class as a declared reference type, for the purpose of polymorphism, but you don't have to worry about somebody making objects of that type. The compiler *guarantees* it.

```
abstract public class Canine extends Animal
{
    public void roam() { }

}

public class MakeCanine {
    public void go() {
        Canine c; } } ← This is OK, because you can always assign
        c = new Dog(); } ← a subclass object to a superclass reference,
        c = new Canine(); ← even if the superclass is abstract.
        c.roam();
    }
}
```

This is OK, because you can always assign a subclass object to a superclass reference, even if the superclass is abstract.

class Canine is marked abstract, so the compiler will NOT let you do this.

```
File Edit Window Help BeamMeUp
% javac MakeCanine.java
MakeCanine.java:5: Canine is abstract;
cannot be instantiated
    c = new Canine();
    ^
1 error
```

An **abstract class** has virtually* no use, no value, no purpose in life, unless it is **extended**.

With an abstract class, the guys doing the work at runtime are **instances of a subclass** of your abstract class.

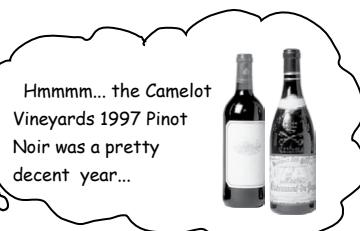
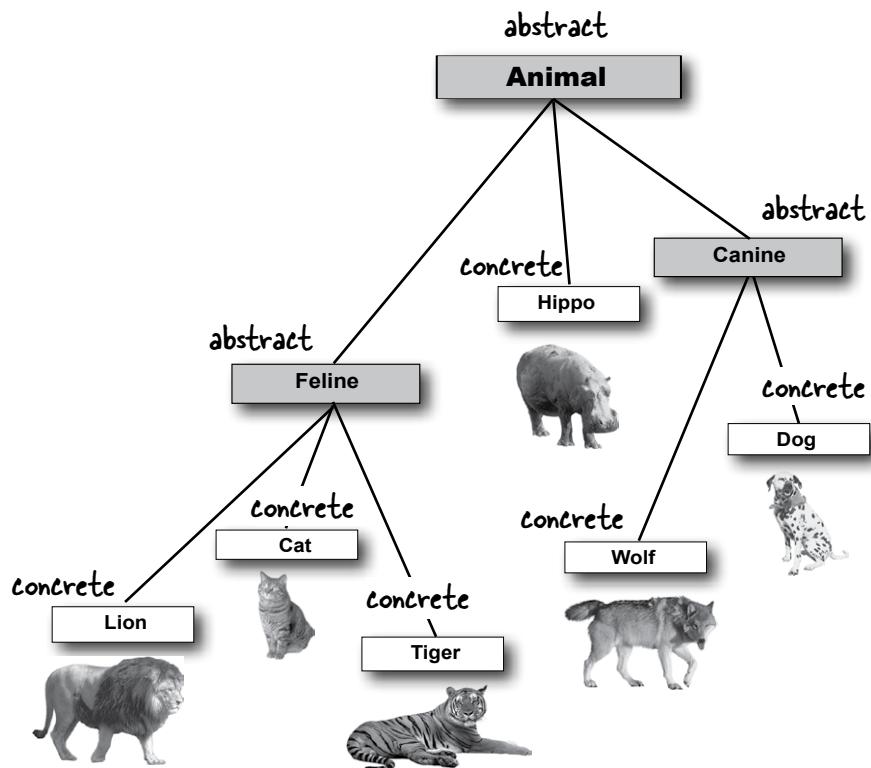
*There is an exception to this—an abstract class can have static members (see chapter 10).

abstract and concrete classes

Abstract vs. Concrete

A class that's not abstract is called a *concrete* class. In the Animal inheritance tree, if we make Animal, Canine, and Feline abstract, that leaves Hippo, Wolf, Dog, Tiger, Lion, and Cat as the concrete subclasses.

Flip through the Java API and you'll find a lot of abstract classes, especially in the GUI library. What does a GUI Component look like? The Component class is the superclass of GUI-related classes for things like buttons, text areas, scrollbars, dialog boxes, you name it. You don't make an instance of a generic *Component* and put it on the screen, you make a JButton. In other words, you instantiate only a *concrete subclass* of Component, but never Component itself.



abstract or concrete?

How do you know when a class should be abstract? **Wine** is probably abstract. But what about **Red** and **White**? Again probably abstract (for some of us, anyway). But at what point in the hierarchy do things become concrete?

Do you make **PinotNoir** concrete, or is it abstract too? It looks like the Camelot Vineyards 1997 Pinot Noir is probably concrete no matter what. But how do you know for sure?

Look at the Animal inheritance tree above. Do the choices we've made for which classes are abstract and which are concrete seem appropriate? Would you change anything about the Animal inheritance tree (other than adding more Animals, of course)?

Abstract methods

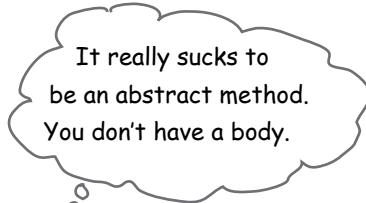
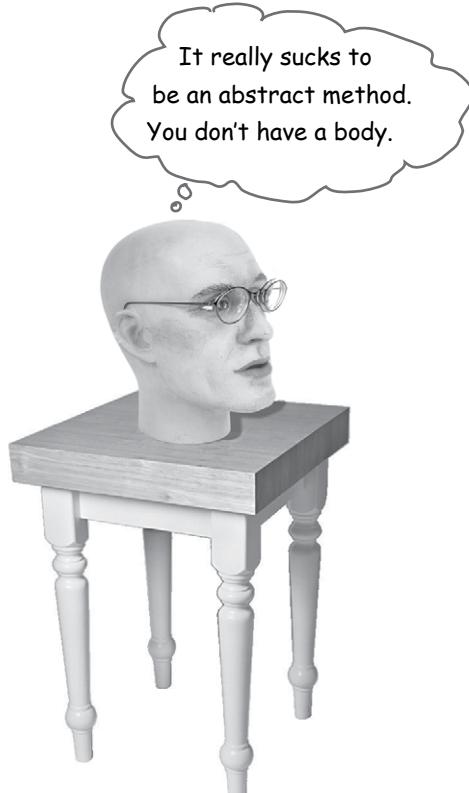
Besides classes, you can mark *methods* abstract, too. An abstract class means the class must be *extended*; an abstract method means the method must be *overridden*. You might decide that some (or all) behaviors in an abstract class don't make any sense unless they're implemented by a more specific subclass. In other words, you can't think of any generic method implementation that could possibly be useful for subclasses. What would a generic `eat()` method look like?

An abstract method has no body!

Because you've already decided there isn't any code that would make sense in the abstract method, you won't put in a method body. So no curly braces—just end the declaration with a semicolon.

```
public abstract void eat();
```

No method body!
End it with a semicolon.

If you declare an abstract method, you MUST mark the class abstract as well. You can't have an abstract method in a non-abstract class.

If you put even a single abstract method in a class, you have to make the class abstract. But you *can* mix both abstract and non-abstract methods in the abstract class.

there are no
Dumb Questions

Q: What is the *point* of an abstract method? I thought the whole point of an abstract class was to have common code that could be inherited by subclasses.

A: Inheritable method implementations (in other words, methods with actual *bodies*) are A Good Thing to put in a superclass. *When it makes sense*. And in an abstract class, it often *doesn't* make sense, because you can't come up with any generic code that subclasses would find useful. The point of an abstract method is that even though you haven't put in any actual method code, you've still defined part of the *protocol* for a group of subtypes (subclasses).

Q: Which is good because...

A: Polymorphism! Remember, what you want is the ability to use a superclass type (often abstract) as a method argument, return type, or array type. That way, you get to add new subtypes (like a new Animal subclass) to your program without having to rewrite (or add) new methods to deal with those new types. Imagine how you'd have to change the Vet class, if it didn't use Animal as its argument type for methods. You'd have to have a separate method for every single Animal subclass! One that takes a Lion, one that takes a Wolf, one that takes a... you get the idea. So with an abstract method, you're saying, "All subtypes of this type have THIS method." for the benefit of polymorphism.

you must implement abstract methods

You **MUST** implement all abstract methods



Implementing an abstract method is just like overriding a method.

Abstract methods don't have a body; they exist solely for polymorphism. That means the first concrete class in the inheritance tree must implement *all* abstract methods.

You can, however, pass the buck by being abstract yourself. If both Animal and Canine are abstract, for example, and both have abstract methods, class Canine does not have to implement the abstract methods from Animal. But as soon as we get to the first concrete subclass, like Dog, that subclass must implement *all* of the abstract methods from both Animal and Canine.

But remember that an abstract class can have both abstract and *non-abstract* methods, so Canine, for example, could implement an abstract method from Animal, so that Dog didn't have to. But if Canine says nothing about the abstract methods from Animal, Dog has to implement all of Animal's abstract methods.

When we say "you must implement the abstract method", that means you *must provide a body*. That means you must create a non-abstract method in your class with the same method signature (name and arguments) and a return type that is compatible with the declared return type of the abstract method. What you put *in* that method is up to you. All Java cares about is that the method is *there*, in your concrete subclass.



Abstract vs. Concrete Classes

Let's put all this abstract rhetoric into some concrete use. In the middle column we've listed some classes. Your job is to imagine applications where the listed class might be concrete, and applications where the listed class might be abstract. We took a shot at the first few to get you going. For example, class `Tree` would be abstract in a tree nursery program, where differences between an Oak and an Aspen matter. But in a golf simulation program, `Tree` might be a concrete class (perhaps a subclass of `Obstacle`), because the program doesn't care about or distinguish between different types of trees. (There's no one right answer; it depends on your design.)

Concrete

golf course simulation

satellite photo application

Sample class

Tree

House

Town

Football Player

Chair

Customer

Sales Order

Book

Store

Supplier

Golf Club

Carburetor

Oven

Abstract

tree nursery application

architect application

coaching application

polymorphism examples

Polymorphism in action

Let's say that we want to write our *own* kind of list class, one that will hold Dog objects, but pretend for a moment that we don't know about the ArrayList class. For the first pass, we'll give it just an *add()* method. We'll use a simple Dog array (Dog []) to keep the added Dog objects, and give it a length of 5. When we reach the limit of 5 Dog objects, you can still call the *add()* method but it won't do anything. If we're *not* at the limit, the *add()* method puts the Dog in the array at the next available index position, then increments that next available index (*nextIndex*).

Building our own Dog-specific list

(Perhaps the world's worst attempt at making our own ArrayList kind of class, from scratch.)

version
1

MyDogList
Dog[] dogs int nextIndex
add(Dog d)

```
public class MyDogList {
    private Dog [] dogs = new Dog[5];
    private int nextIndex = 0; ← We'll increment this each time a new Dog is added.

    public void add(Dog d) {
        if (nextIndex < dogs.length) {
            dogs[nextIndex] = d;
            System.out.println("Dog added at " + nextIndex);
            nextIndex++;
        }
    }
}
```

Use a plain old Dog array behind the scenes.

If we're not already at the limit of the dogs array, add the Dog and print a message.

increment, to give us the next index to use

Uh-oh, now we need to keep Cats, too.

We have a few options here:

- 1) Make a separate class, MyCatList, to hold Cat objects. Pretty clunky.
- 2) Make a single class, DogAndCatList, that keeps two different arrays as instance variables and has two different add() methods: addCat(Cat c) and addDog(Dog d). Another clunky solution.
- 3) Make heterogeneous AnimalList class, that takes *any* kind of Animal subclass (since we know that if the spec changed to add Cats, sooner or later we'll have some *other* kind of animal added as well). We like this option best, so let's change our class to make it more generic, to take Animals instead of just Dogs. We've highlighted the key changes (the logic is the same, of course, but the type has changed from Dog to Animal everywhere in the code).

Building our own Animal-specific list

MyAnimalList
Animal[] animals
int nextIndex
add(Animal a)

version 2

```

public class MyAnimalList {
    private Animal[] animals = new Animal[5];
    private int nextIndex = 0;

    public void add(Animal a) {
        if (nextIndex < animals.length) {
            animals[nextIndex] = a;
            System.out.println("Animal added at " + nextIndex);
            nextIndex++;
        }
    }
}

```

Don't panic. We're not making a new Animal object; we're making a new array object, of type Animal. (Remember, you cannot make a new instance of an abstract type, but you CAN make an array object declared to HOLD that type.)

```

public class AnimalTestDrive{
    public static void main (String[] args) {
        MyAnimalList list = new MyAnimalList();
        Dog a = new Dog();
        Cat c = new Cat();
        list.add(a);
        list.add(c);
    }
}

File Edit Window Help Harm
% java AnimalTestDrive
Animal added at 0
Animal added at 1

```

the ultimate superclass: Object

What about non-Animals? Why not make a class generic enough to take anything?

You know where this is heading. We want to change the type of the array, along with the `add()` method argument, to something *above* Animal. Something even *more* generic, *more* abstract than Animal. But how can we do it? We don't *have* a superclass for Animal.

Then again, maybe we do...

Remember those methods of ArrayList?
Look how the remove, contains, and
indexOf method all use an object of type...
Object!

Every class in Java extends class Object.

Class Object is the mother of all classes; it's the superclass of *everything*.

Even if you take advantage of polymorphism, you still have to create a class with methods that take and return *your* polymorphic type. Without a common superclass for everything in Java, there'd be no way for the developers of Java to create classes with methods that could take *your* custom types... *types they never knew about when they wrote the ArrayList class.*

So you were making subclasses of class Object from the very beginning and you didn't even know it. *Every class you write extends Object,* without your ever having to say it. But you can think of it as though a class you write looks like this:

```
public class Dog extends Object { }
```

But wait a minute, Dog *already* extends something, *Canine*. That's OK. The compiler will make *Canine* extend Object instead. Except *Canine* extends Animal. No problem, then the compiler will just make *Animal* extend Object.

Any class that doesn't explicitly extend another class, implicitly extends Object.

So, since Dog extends Canine, it doesn't *directly* extend Object (although it does extend it indirectly), and the same is true for Canine, but Animal *does* directly extend Object.

version
3

(These are just a few of the methods in ArrayList...there are many more.)

ArrayList

boolean remove(Object elem)

Removes one instance of the object specified in the parameter. Returns 'true' if the element was in the list.

boolean contains(Object elem)

Returns 'true' if there's a match for the object parameter.

boolean isEmpty()

Returns 'true' if the list has no elements.

int indexOf(Object elem)

Returns either the index of the object parameter, or -1.

Object get(int index)

Returns the element at this position in the list.

boolean add(Object elem)

Adds the element to the list (returns 'true').

// more

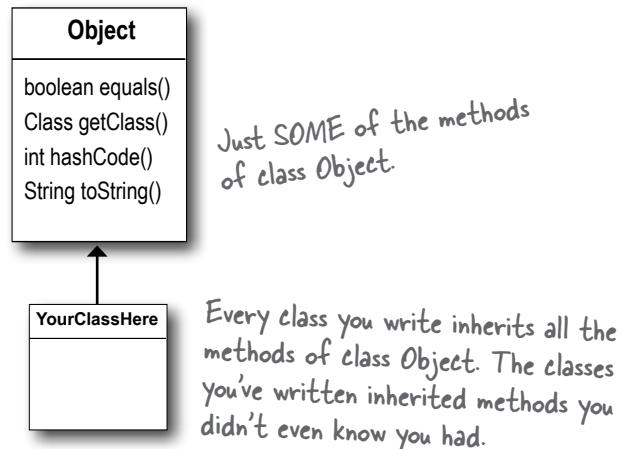
Many of the ArrayList methods use the ultimate polymorphic type, Object. Since every class in Java is a subclass of Object, these ArrayList methods can take anything!

(Note: as of Java 5.0, the `get()` and `add()` methods actually look a little different than the ones shown here, but for now this is the way to think about it. We'll get into the full story a little later.)

So what's in this ultra-super-mega-class Object?

If you were Java, what behavior would you want *every* object to have? Hmm... let's see... how about a method that lets you find out if one object is equal to another object? What about a method that can tell you the actual class type of that object? Maybe a method that gives you a hashCode for the object, so you can use the object in hashtables (we'll talk about Java's hashtables in chapter 16). Oh, here's a good one—a method that prints out a String message for that object.

And what do you know? As if by magic, class Object does indeed have methods for those four things. That's not all, though, but these are the ones we really care about.



① equals(Object o)

```

Dog a = new Dog();
Cat c = new Cat();

if (a.equals(c)) {
    System.out.println("true");
} else {
    System.out.println("false");
}
  
```

```

File Edit Window Help Stop
% java TestObject
false
  
```

Tells you if two objects are considered 'equal'.

③ hashCode()

```

Cat c = new Cat();
System.out.println(c.hashCode());
  
```

```

File Edit Window Help Drop
% java TestObject
8202111
  
```

Prints out a hashCode for the object (for now, think of it as a unique ID).

② getClass()

```

Cat c = new Cat();
System.out.println(c.getClass());
  
```

```

File Edit Window Help Faint
% java TestObject
class Cat
  
```

Gives you back the class that object was instantiated from.

④ toString()

```

Cat c = new Cat();
System.out.println(c.toString());
  
```

```

File Edit Window Help LapseIntoComa
% java TestObject
Cat@7d277f
  
```

Prints out a String message with the name of the class and some other number we rarely care about.

Object and abstract classes

^{there are no}
Dumb Questions

Q: Is class Object abstract?

A: No. Well, not in the formal Java sense anyway. Object is a non-abstract class because it's got method implementation code that all classes can inherit and use out-of-the-box, without having to override the methods.

Q: Then can you override the methods in Object?

A: Some of them. But some of them are marked `final`, which means you can't override them. You're encouraged (strongly) to override `hashCode()`, `equals()`, and `toString()` in your own classes, and you'll learn how to do that a little later in the book. But some of the methods, like `getClass()`, do things that must work in a specific, guaranteed way.

Q: If ArrayList methods are generic enough to use Object, then what does it mean to say `ArrayList<DotCom>`? I thought I was restricting the ArrayList to hold only DotCom objects?

A: You were restricting it. Prior to Java 5.0, ArrayLists couldn't be restricted. They were all essentially what you get in Java 5.0 today if you write `ArrayList<Object>`. In other words, **an ArrayList restricted to anything that's an Object**, which means *any* object in Java, instantiated from *any* class type! We'll cover the details of this new `<type>` syntax later in the book.

Q: OK, back to class Object being non-abstract (so I guess that means it's concrete), HOW can you let somebody make an Object object? Isn't that just as weird as making an Animal object?

A: Good question! Why is it acceptable to make a new Object instance? Because sometimes you just want a generic object to use as, well, as an object. A *lightweight* object. By far, the most common use of an instance of type Object is for thread synchronization (which you'll learn about in chapter 15). For now, just stick that on the back burner and assume that you will rarely make objects of type Object, even though you *can*.

Q: So is it fair to say that the main purpose for type Object is so that you can use it for a polymorphic argument and return type? Like in ArrayList?

A: The Object class serves two main purposes: to act as a polymorphic type for methods that need to work on any class that you or anyone else makes, and to provide *real* method code that all objects in Java need at runtime (and putting them in class Object means all other classes inherit them). Some of the most important methods in Object are related to threads, and we'll see those later in the book.

Q: If it's so good to use polymorphic types, why don't you just make ALL your methods take and return type Object?

A: Ahhhh... think about what would happen. For one thing, you would defeat the whole point of 'type-safety', one of Java's greatest protection mechanisms for your code. With type-safety, Java guarantees that you won't ask the wrong object to do something you *meant* to ask of another object type. Like, ask a *Ferrari* (which you think is a *Toaster*) to *cook itself*. But the truth is, you *don't* have to worry about that fiery Ferrari scenario, even if you *do* use Object references for everything. Because when objects are referred to by an Object reference type, Java *thinks* it's referring to an instance of type Object. And that means the only methods you're allowed to call on that object are the ones declared in class Object! So if you were to say:

```
Object o = new Ferrari();  
o.goFast(); //Not legal!
```

You wouldn't even make it past the compiler.

Because Java is a strongly-typed language, the compiler checks to make sure that you're calling a method on an object that's actually capable of *responding*. In other words, you can call a method on an object reference *only* if the class of the reference type actually *has* the method. We'll cover this in much greater detail a little later, so don't worry if the picture isn't crystal clear.

Using polymorphic references of type Object has a price...

Before you run off and start using type Object for all your ultra-flexible argument and return types, you need to consider a little issue of using type Object as a reference. And keep in mind that we're not talking about making instances of type Object; we're talking about making instances of some other type, but using a reference of type Object.

When you put an object into an `ArrayList<Dog>`, it goes in as a Dog, and comes out as a Dog:

```
ArrayList<Dog> myDogArrayList = new ArrayList<Dog>(); ← Make an ArrayList declared
to hold Dog objects.  

Dog aDog = new Dog(); ← Make a Dog.  

myDogArrayList.add(aDog); ← Add the Dog to the list.  

Dog d = myDogArrayList.get(0); ← Assign the Dog from the list to a new Dog reference variable.  

(Think of it as though the get() method declares a Dog return  

type because you used ArrayList<Dog>.)
```

But what happens when you declare it as `ArrayList<Object>`? If you want to make an `ArrayList` that will literally take *any* kind of Object, you declare it like this:

```
ArrayList<Object> myDogArrayList = new ArrayList<Object>(); ← Make an ArrayList declared
to hold any type of Object.  

Dog aDog = new Dog(); ← Make a Dog. (These two steps are the same.)  

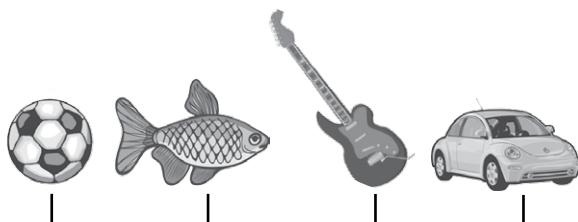
myDogArrayList.add(aDog); ← Add the Dog to the list.
```

But what happens when you try to get the Dog object and assign it to a Dog reference?

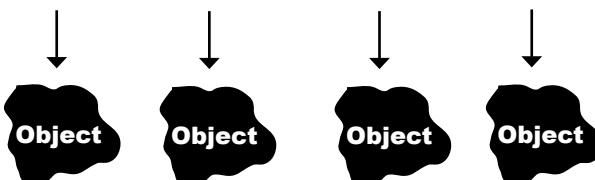
 `Dog d = myDogArrayList.get(0);` NO!! Won't compile!! When you use `ArrayList<Object>`, the `get()` method returns type Object. The Compiler knows only that the object inherits from Object (somewhere in its inheritance tree) but it doesn't know it's a Dog!!

Everything comes out of an `ArrayList<Object>` as a reference of type Object, regardless of what the actual object is, or what the reference type was when you added the object to the list.

The objects go IN as **SoccerBall**, **Fish**, **Guitar**, and **Car**.



But they come OUT as though they were of type **Object**.



Objects come out of an `ArrayList<Object>` acting like they're generic instances of class Object. The Compiler cannot assume the object that comes out is of any type other than Object.

When a Dog loses its Dogness

When a Dog won't act like a Dog

The problem with having everything treated polymorphically as an Object is that the objects *appear* to lose (but not permanently) their true essence. *The Dog appears to lose its dogness.* Let's see what happens when we pass a Dog to a method that returns a reference to the same Dog object, but declares the return type as type Object rather than Dog.



BAD ☹

```
public void go() {
    Dog aDog = new Dog();
    Dog sameDog = getobject(aDog);
}
```

This line won't work! Even though the method returned a reference to the very same Dog the argument referred to, the return type Object means the compiler won't let you assign the returned reference to anything but Object.

```
public Object getObject(Object o) {
    return o;
}
```

↑ We're returning a reference to the same Dog, but as a return type of Object. This part is perfectly legal. Note: this is similar to how the get() method works when you have an ArrayList<Object> rather than an ArrayList<Dog>.

```
File Edit Window Help Remember
DogPolyTest.java:10: incompatible types
found   : java.lang.Object
required: Dog
    Dog sameDog = getObject(aDog);
1 error
```

The compiler doesn't know that the thing returned from the method is actually a Dog, so it won't let you assign it to a Dog reference. (You'll see why on the next page.)

GOOD ☺

```
public void go() {
    Dog aDog = new Dog();
    Object sameDog = getObject(aDog);
}

public Object getObject(Object o) {
    return o;
}
```

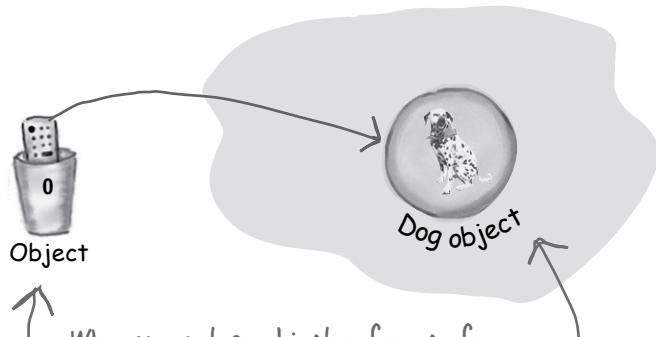
This works (although it may not be very useful, as you'll see in a moment) because you can assign ANYTHING to a reference of type Object, since every class passes the IS-A test for Object. Every object in Java is an instance of type Object, because every class in Java has Object at the top of its inheritance tree.

Objects don't bark.

So now we know that when an object is referenced by a variable declared as type Object, it can't be assigned to a variable declared with the actual object's type. And we know that this can happen when a return type or argument is declared as type Object, as would be the case, for example, when the object is put into an ArrayList of type Object using `ArrayList<Object>`. But what are the implications of this? Is it a problem to have to use an Object reference variable to refer to a Dog object? Let's try to call Dog methods on our Dog-That-Compiler-Thinks-Is-An-Object:

```
Object o = al.get(index);
int i = o.hashCode();
o.bark();
```

Won't compile!



When you get an object reference from an `ArrayList<Object>` (or any method that declares `Object` as the return type), it comes back as a polymorphic reference type of `Object`. So you have an Object reference to (in this case) a Dog instance.

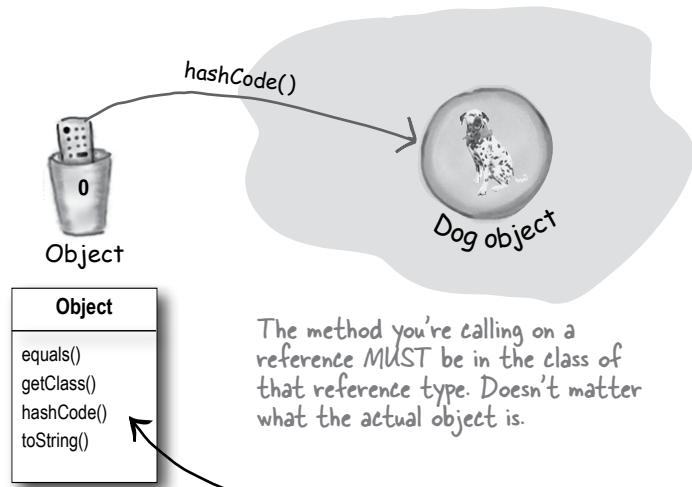
This is fine. Class `Object` has a `hashCode()` method, so you can call that method on ANY object in Java.

Can't do this!! The `Object` class has no idea what it means to `bark()`. Even though YOU know it's really a Dog at that index, the compiler doesn't..

The compiler decides whether you can call a method based on the reference type, not the actual object type.

Even if you *know* the object is capable ("...but it really *is* a Dog, honest..."), the compiler sees it only as a generic Object. For all the compiler knows, you put a Button object out there. Or a Microwave object. Or some other thing that really doesn't know how to bark.

The compiler checks the class of the *reference type*—not the *object type*—to see if you can call a method using that reference.

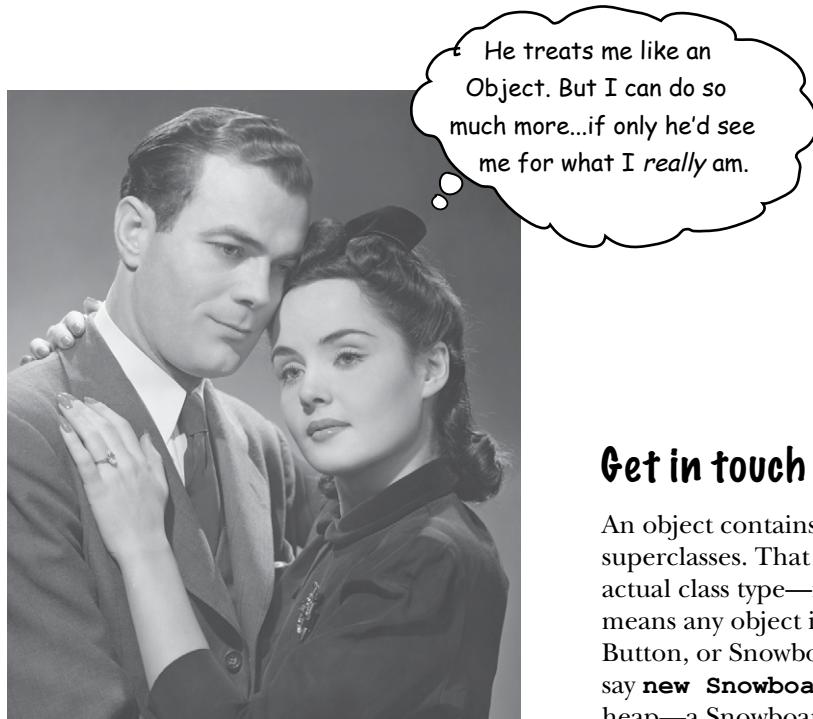


The method you're calling on a reference **MUST** be in the class of that reference type. Doesn't matter what the actual object is.

`o.hashCode();`

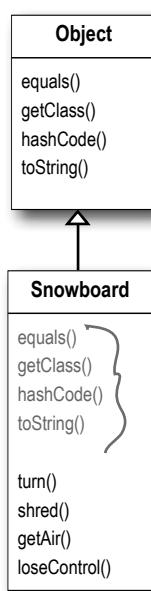
The "o" reference was declared as type `Object`, so you can call methods only if those methods are in class `Object`..

objects are Objects

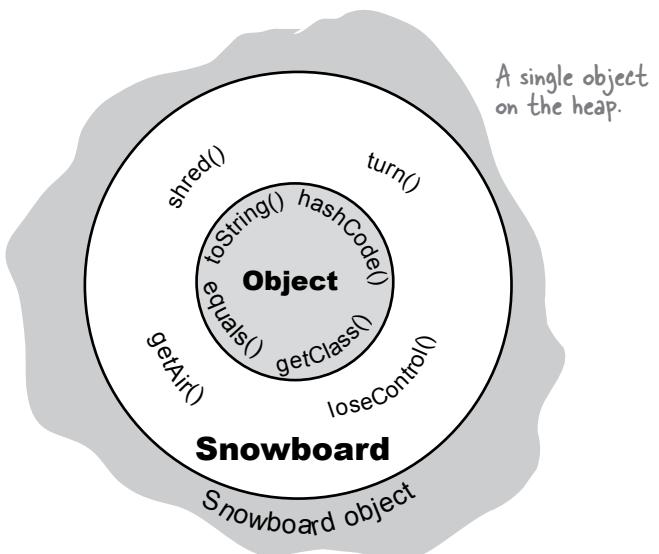


Get in touch with your inner Object.

An object contains *everything* it inherits from each of its superclasses. That means *every* object—regardless of its actual class type—is *also* an instance of class `Object`. That means any object in Java can be treated not just as a Dog, Button, or Snowboard, but also as an Object. When you say `new Snowboard()`, you get a single object on the heap—a Snowboard object—but that Snowboard wraps itself around an inner core representing the `Object` (capital “O”) portion of itself.



Snowboard inherits methods from superclass Object, and adds four more.



There is only ONE object on the heap here. A Snowboard object. But it contains both the Snowboard class parts of itself and the Object class parts of itself.

**'Polymorphism' means
'many forms'.**

**You can treat a Snowboard as a
Snowboard or as an Object.**

If a reference is like a remote control, the remote control takes on more and more buttons as you move down the inheritance tree. A remote control (reference) of type Object has only a few buttons—the buttons for the exposed methods of class Object. But a remote control of type Snowboard includes all the buttons from class Object, plus any new buttons (for new methods) of class Snowboard. The more specific the class, the more buttons it may have.

Of course that's not always true; a subclass might not add any new methods, but simply override the methods of its superclass. The key point is that even if the *object* is of type Snowboard, an Object *reference* to the Snowboard object can't see the Snowboard-specific methods.

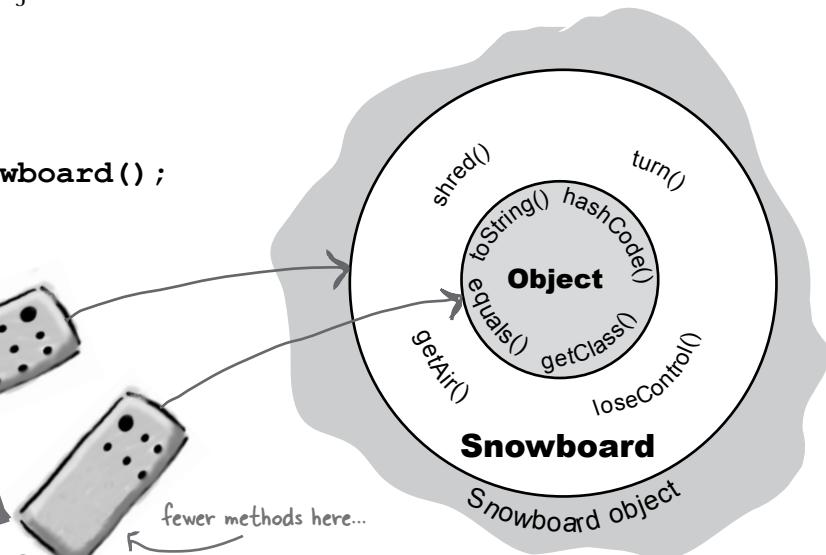
**When you put
an object in an
ArrayList<Object>, you
can treat it only as an
Object, regardless of
the type it was when
you put it in.**

**When you get a
reference from an
ArrayList<Object>, the
reference is always of
type Object.**

**That means you get an
Object remote control.**

```
Snowboard s = new Snowboard();
Object o = s;
```

The Snowboard remote control (reference) has more buttons than an Object remote control. The Snowboard remote can see the full Snowboardness of the Snowboard object. It can access all the methods in Snowboard, including both the inherited Object methods and the methods from class Snowboard.



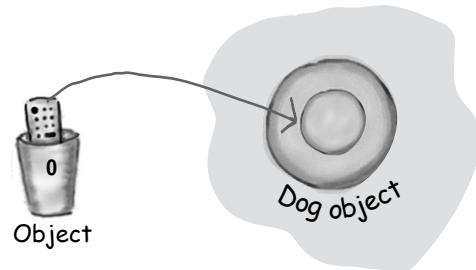
The Object reference can see only the Object parts of the Snowboard object. It can access only the methods of class Object. It has fewer buttons than the Snowboard remote control.

casting objects



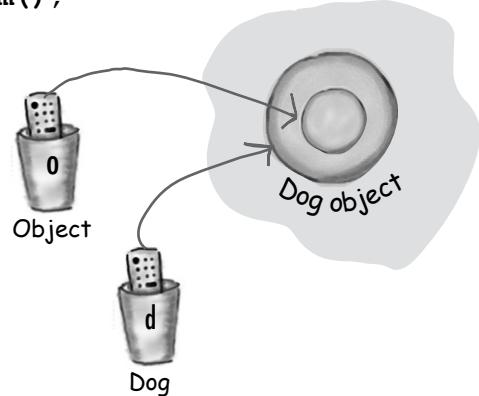
Cast the so-called 'Object' (but we know he's actually a Dog) to type Dog, so that you can treat him like the Dog he really is.

Casting an object reference back to its *real* type.



It's really still a Dog *object*, but if you want to call Dog-specific methods, you need a *reference* declared as type Dog. If you're *sure** the object is really a Dog, you can make a new Dog reference to it by copying the Object reference, and forcing that copy to go into a Dog reference variable, using a cast (Dog). You can use the new Dog reference to call Dog methods.

```
Object o = al.get(index);  
Dog d = (Dog) o; ← cast the Object back to  
a Dog we know is there.  
d.roam();
```



*If you're *not* sure it's a Dog, you can use the `instanceof` operator to check. Because if you're wrong when you do the cast, you'll get a ClassCastException at runtime and come to a grinding halt.

```
if (o instanceof Dog) {  
    Dog d = (Dog) o;  
}
```

So now you've seen how much Java cares about the methods in the class of the reference variable.

You can call a method on an object only if the class of the reference variable has that method.

Think of the public methods in your class as your contract, your promise to the outside world about the things you can do.



When you write a class, you almost always *expose* some of the methods to code outside the class. To *expose* a method means you make a method *accessible*, usually by marking it public.

Imagine this scenario: you're writing code for a small business accounting program. A custom application for "Simon's Surf Shop". The good re-user that you are, you found an Account class that appears to meet your needs perfectly, according to its documentation, anyway. Each account instance represents an individual customer's account with the store. So there you are minding your own business invoking the *credit()* and *debit()* methods on an account object when you realize you need to get a balance on an account. No problem—there's a *getBalance()* method that should do nicely.

Account
debit(double amt)
credit(double amt)
double getBalance()

Except... when you invoke the *getBalance()* method, the whole thing blows up at runtime. Forget the documentation, the class does not have that method. Yikes!

But that won't happen to you, because everytime you use the dot operator on a reference (*a.doStuff()*), the compiler looks at the *reference type* (the type 'a' was declared to be) and checks that class to guarantee the class has the method, and that the method does indeed take the argument you're passing and return the kind of value you're expecting to get back.

Just remember that the compiler checks the class of the reference variable, not the class of the actual *object* at the other end of the reference.

What if you need to change the contract?

OK, pretend you're a Dog. Your Dog class isn't the *only* contract that defines who you are. Remember, you inherit accessible (which usually means *public*) methods from all of your superclasses.

True, your Dog class defines a contract.

But not *all* of your contract.

Everything in class *Canine* is part of your contract.

Everything in class *Animal* is part of your contract.

Everything in class *Object* is part of your contract.

According to the IS-A test, you *are* each of those things—Canine, Animal, and Object.

But what if the person who designed your class had in mind the Animal simulation program, and now he wants to use you (class Dog) for a Science Fair Tutorial on Animal objects.

That's OK, you're probably reusable for that.

But what if later he wants to use you for a PetShop program? *You don't have any Pet behaviors.* A Pet needs methods like *beFriendly()* and *play()*.

OK, now pretend you're the Dog class programmer. No problem, right? Just add some more methods to the Dog class. You won't be breaking anyone else's code by *adding* methods, since you aren't touching the *existing* methods that someone else's code might be calling on Dog objects.

Can you see any drawbacks to that approach (adding Pet methods to the Dog class)?



Think about what **YOU** would do if **YOU** were the Dog class programmer and needed to modify the Dog so that it could do Pet things, too. We know that simply adding new Pet behaviors (methods) to the Dog class will work, and won't break anyone else's code.

But... this is a PetShop program. It has more than just Dogs! And what if someone wants to use your Dog class for a program that has *wild Dogs*? What do you think your options might be, and without worrying about how Java handles things, just try to imagine how you'd *like* to solve the problem of modifying some of your Animal classes to include Pet behaviors.

Stop right now and think about it, before you look at the next page where we begin to reveal *everything*.

(thus rendering the whole exercise completely useless, robbing you of your One Big Chance to burn some brain calories)

Let's explore some design options for reusing some of our existing classes in a PetShop program.

On the next few pages, we're going to walk through some possibilities. We're not yet worried about whether Java can actually *do* what we come up with. We'll cross that bridge once we have a good idea of some of the tradeoffs.

① Option one

We take the easy path, and put pet methods in class Animal.

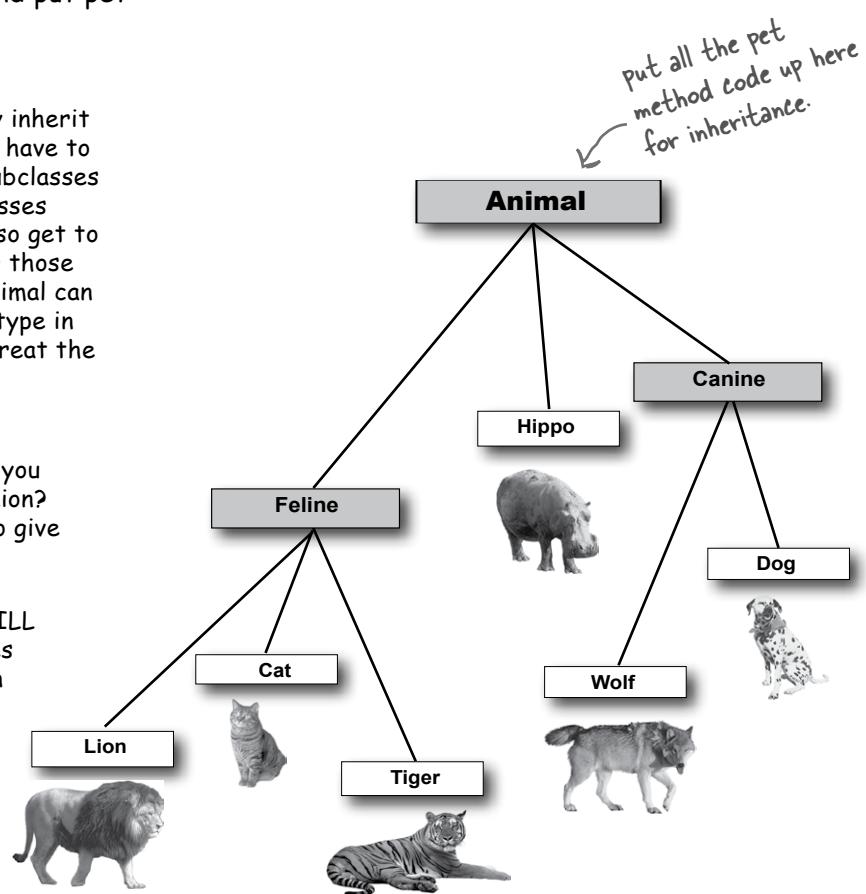
Pros:

All the Animals will instantly inherit the pet behaviors. We won't have to touch the existing Animal subclasses at all, and any Animal subclasses created in the future will also get to take advantage of inheriting those methods. That way, class Animal can be used as the polymorphic type in any program that wants to treat the Animals as pets

Cons:

So... when was the last time you saw a Hippo at a pet shop? Lion? Wolf? Could be dangerous to give non-pets pet methods.

Also, we almost certainly WILL have to touch the pet classes like Dog and Cat, because (in our house, anyway) Dogs and Cats tend to implement pet behaviors VERY differently.



modifying existing classes

② Option two

We start with Option One, putting the pet methods in class Animal, but we make the methods abstract, forcing the Animal subclasses to override them.

Pros:

That would give us all the benefits of Option One, but without the drawback of having non-pet Animals running around with pet methods (like `beFriendly()`). All Animal classes would have the method (because it's in class Animal), but because it's abstract the non-pet Animal classes won't inherit any functionality. All classes **MUST** override the methods, but they can make the methods "do-nothings".

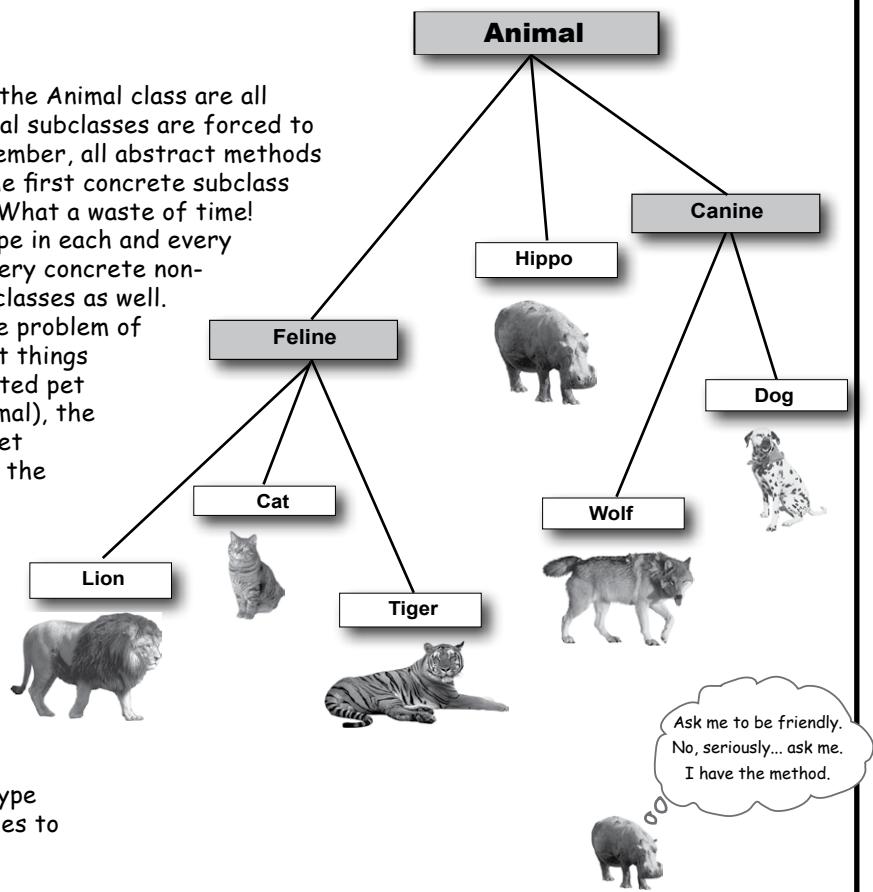
Put all the pet methods up here, but with no implementations. Make all pet methods abstract.

Cons:

Because the pet methods in the Animal class are all abstract, the concrete Animal subclasses are forced to implement all of them. (Remember, all abstract methods **MUST** be implemented by the first concrete subclass down the inheritance tree.) What a waste of time! You have to sit there and type in each and every pet method into each and every concrete non-pet class, and all future subclasses as well.

And while this does solve the problem of non-pets actually **DOING** pet things (as they would if they inherited pet functionality from class Animal), the contract is bad. Every *non-pet* class would be announcing to the world that it, too, has those pet methods, even though the methods wouldn't actually **DO** anything when called.

This approach doesn't look good at all. It just seems wrong to stuff everything into class Animal that more than one Animal type might need, UNLESS it applies to **ALL** Animal subclasses.



③ Option three

Put the pet methods ONLY in the classes where they belong.

Pros:

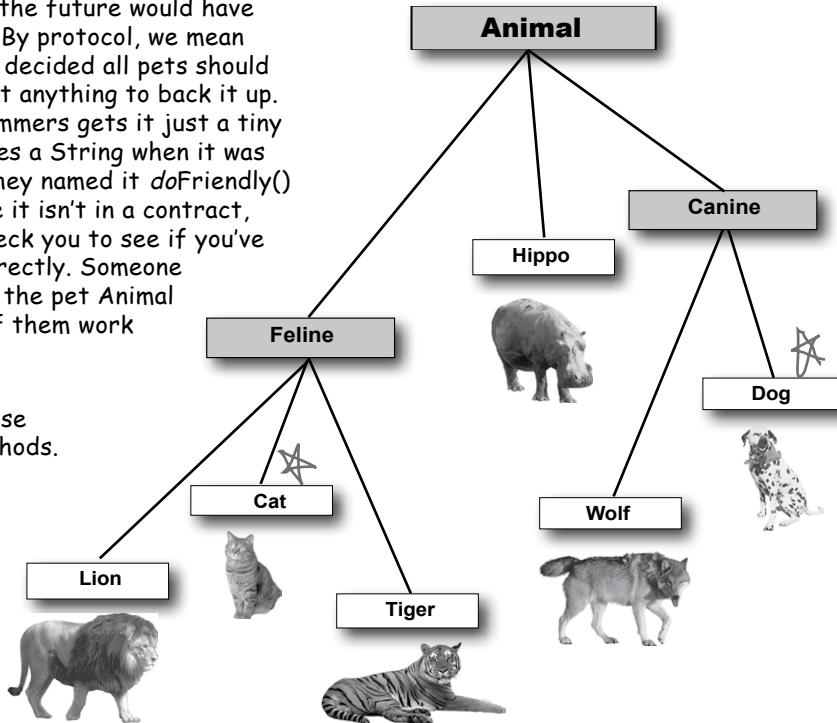
No more worries about Hippos greeting you at the door or licking your face. The methods are where they belong, and ONLY where they belong. Dogs can implement the methods and Cats can implement the methods, but nobody else has to know about them.

Cons:

Two Big Problems with this approach. First off, you'd have to agree to a protocol, and all programmers of pet Animal classes now and in the future would have to KNOW about the protocol. By protocol, we mean the exact methods that we've decided all pets should have. The pet contract without anything to back it up. But what if one of the programmers gets it just a tiny bit wrong? Like, a method takes a String when it was supposed to take an int? Or they named it `doFriendly()` instead of `beFriendly()`? Since it isn't in a contract, the compiler has no way to check you to see if you've implemented the methods correctly. Someone could easily come along to use the pet Animal classes and find that not all of them work quite right.

And second, you don't get to use polymorphism for the pet methods. Every class that needs to use pet behaviors would have to know about each and every class! In other words, you can't use `Animal` as the polymorphic type now, because the compiler won't let you call a Pet method on an `Animal` reference (even if it's really a Dog object) because class `Animal` doesn't have the method.

~~Put the pet methods ONLY in the Animal classes that can be pets, instead of in Animal.~~

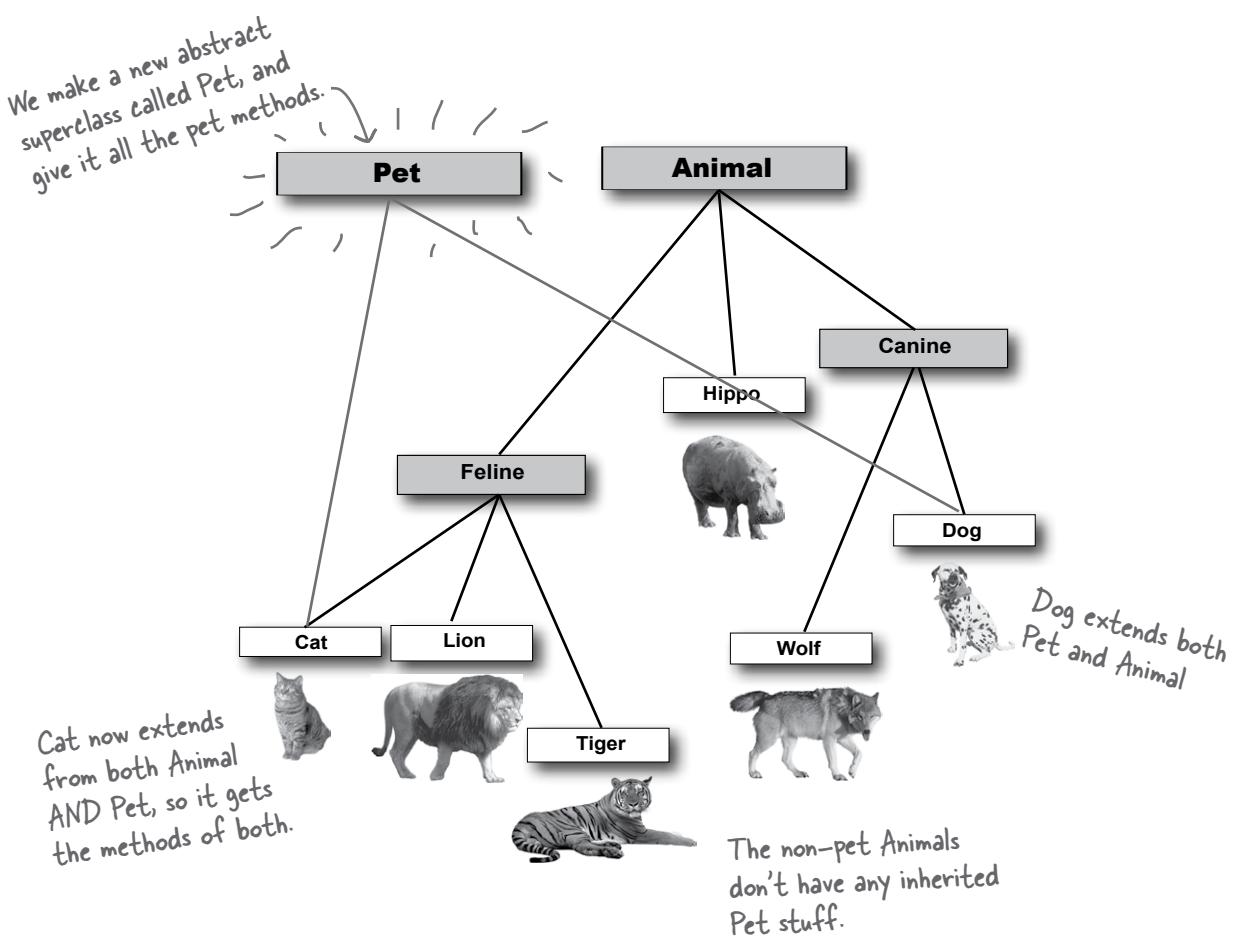


multiple inheritance?

So what we **REALLY** need is:

- ↗ A way to have pet behavior in **just** the pet classes
- ↗ A way to guarantee that all pet classes have all of the same methods defined (same name, same arguments, same return types, no missing methods, etc.), without having to cross your fingers and hope all the programmers get it right.
- ↗ A way to take advantage of polymorphism so that all pets can have their pet methods called, without having to use arguments, return types, and arrays for each and every pet class.

It looks like we need TWO superclasses at the top



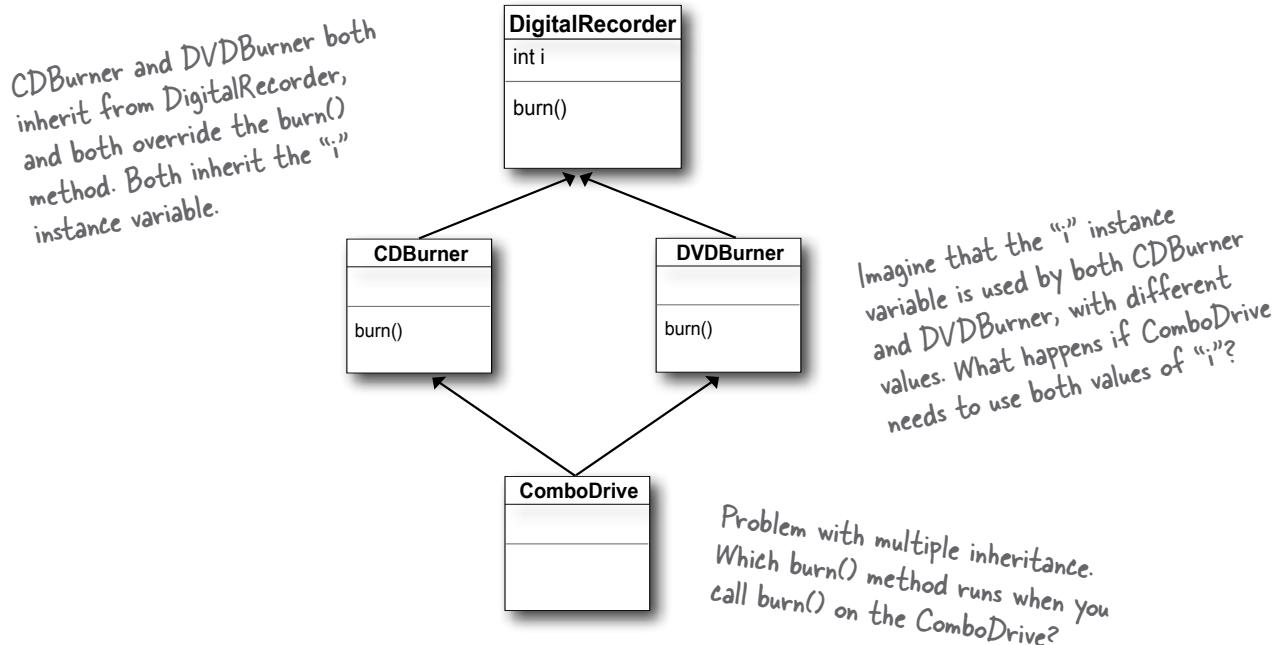
There's just one problem with the "two superclasses" approach...

It's called "multiple inheritance" and it can be a Really Bad Thing.

That is, if it were possible to do in Java.

But it isn't, because multiple inheritance has a problem known as The Deadly Diamond of Death.

Deadly Diamond of Death



A language that allows the Deadly Diamond of Death can lead to some ugly complexities, because you have to have special rules to deal with the potential ambiguities. And extra rules means extra work for you both in *learning* those rules and watching out for those "special cases". Java is supposed to be *simple*, with consistent rules that don't blow up under some scenarios. So Java (unlike C++) protects you from having to think about the Deadly Diamond of Death. But that brings us back to the original problem! *How do we handle the Animal/Pet thing?*

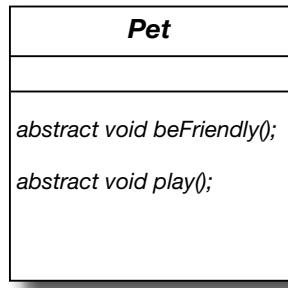
interfaces

Interface to the rescue!

Java gives you a solution. An *interface*. Not a *GUI* interface, not the generic use of the *word* interface as in, “That’s the public interface for the Button class API,” but the Java *keyword* **interface**.

A Java interface solves your multiple inheritance problem by giving you much of the polymorphic *benefits* of multiple inheritance without the pain and suffering from the Deadly Diamond of Death (DDD).

The way in which interfaces side-step the DDD is surprisingly simple: *make all the methods abstract!* That way, the subclass *must* implement the methods (remember, abstract methods *must* be implemented by the first concrete subclass), so at runtime the JVM isn’t confused about *which* of the two inherited versions it’s supposed to call.



A Java interface is like a 100% pure abstract class.

All methods in an interface are abstract, so any class that IS-A Pet MUST implement (i.e. override) the methods of Pet.

To DEFINE an interface:

```
public interface Pet { ... }
```

↑
Use the keyword “interface” instead of “class”

To IMPLEMENT an interface:

```
public class Dog extends Canine implements Pet { ... }
```

↑
Use the keyword “implements” followed by the interface name. Note that when you implement an interface you still get to extend a class

Making and Implementing the Pet interface

You say 'interface' instead of 'class' here

```
public interface Pet {
    public abstract void beFriendly();
    public abstract void play();
}
```

interface methods are implicitly public and is optional (in fact, it's not considered 'good style' to type the words in, but we did here just to reinforce it, and because we've never been slaves to fashion...)

All interface methods are abstract, so they MUST end in semicolons. Remember, they have no body!

Dog IS-A Animal
and Dog IS-A Pet

```
public class Dog extends Canine implements Pet {
    public void beFriendly() {...}
    public void play() {...}
    public void roam() {...}
    public void eat() {...}
}
```

You say 'implements' followed by the name of the interface.

You SAID you are a Pet, so you MUST implement the Pet methods. It's your contract. Notice the curly braces instead of semicolons.

These are just normal overriding methods.

there are no Dumb Questions

Q: Wait a minute, interfaces don't really give you multiple inheritance, because you can't put any implementation code in them. If all the methods are abstract, what does an interface really buy you?

A: Polymorphism, polymorphism, polymorphism. Interfaces are the ultimate in flexibility, because if you use interfaces instead of concrete subclasses (or even abstract superclass types) as arguments and return

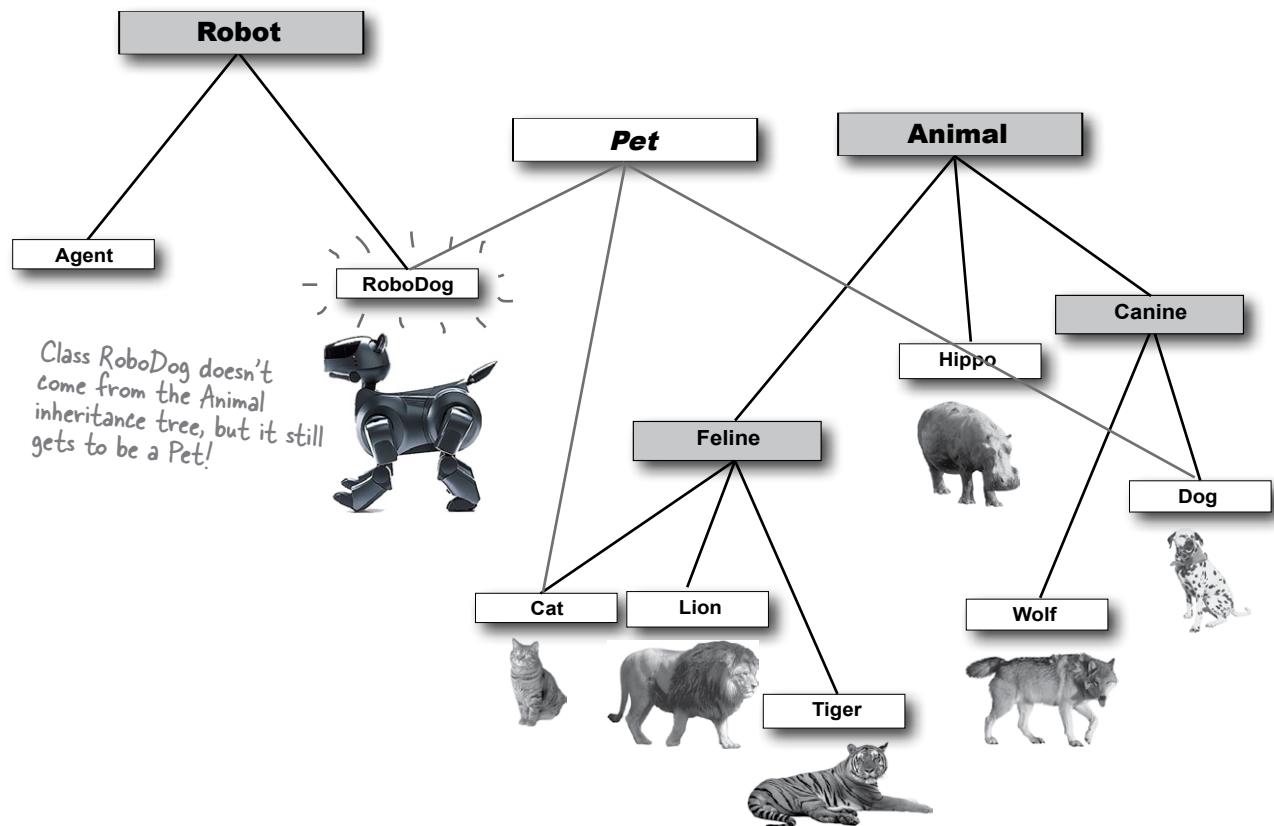
types, you can pass anything that implements that interface. And think about it—with an interface, a class doesn't have to come from just one inheritance tree. A class can extend one class, and implement an interface. But another class might implement the same interface, yet come from a completely different inheritance tree! So you get to treat an object by the role it plays, rather than by the class type from which it was instantiated. In fact, if you wrote your code to use interfaces, you wouldn't even have to give anyone a superclass

that they had to extend. You could just give them the interface and say, "Here, I don't care what kind of class inheritance structure you come from, just implement this interface and you'll be good to go."

The fact that you can't put in implementation code turns out not to be a problem for most good designs, because most interface methods wouldn't make sense if implemented in a generic way. In other words, most interface methods would need to be overridden even if the methods weren't forced to be abstract.

interface polymorphism

Classes from different inheritance trees can implement the same interface.



When you use a *class* as a polymorphic type (like an array of type Animal or a method that takes a Canine argument), the objects you can stick in that type must be from the same inheritance tree. But not just anywhere in the inheritance tree; the objects must be from a class that is a subclass of the polymorphic type. An argument of type Canine can accept a Wolf and a Dog, but not a Cat or a Hippo.

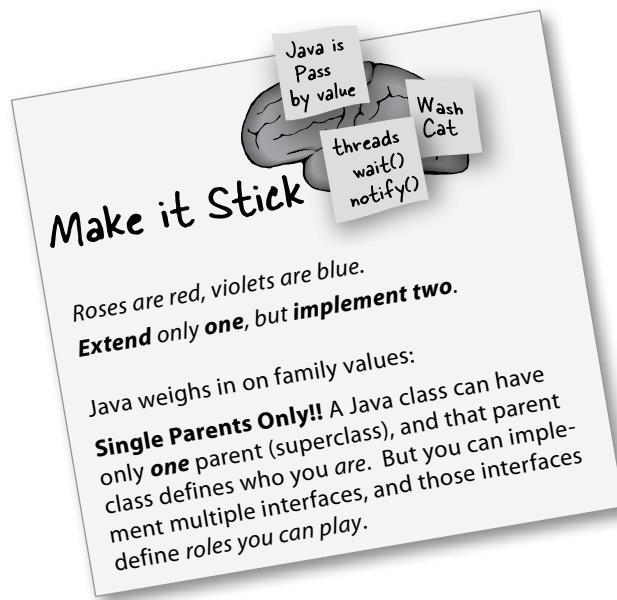
But when you use an *interface* as a polymorphic type (like an array of Pets), the objects can be from *anywhere* in the inheritance tree. The only requirement is that the objects are from a class that *implements* the interface. Allowing classes in different inheritance trees to implement a common interface is crucial in the Java API. Do you want an object to be able to save its state to a file? Implement the Serializable interface. Do you need objects to run

their methods in a separate thread of execution? Implement Runnable. You get the idea. You'll learn more about Serializable and Runnable in later chapters, but for now, remember that classes from *any* place in the inheritance tree might need to implement those interfaces. Nearly *any* class might want to be saveable or runnable.

Better still, a class can implement multiple interfaces!

A Dog object IS-A Canine, and IS-A Animal, and IS-A Object, all through inheritance. But a Dog IS-A Pet through interface implementation, and the Dog might implement other interfaces as well. You could say:

```
public class Dog extends Animal implements Pet, Saveable, Paintable { ... }
```



How do you know whether to make a class, a subclass, an **abstract class**, or an **interface**?

- Make a class that doesn't extend anything (other than Object) when your new class doesn't pass the IS-A test for any other type.
- Make a subclass (in other words, *extend a class*) only when you need to make a **more specific** version of a class and need to override or add new behaviors.
- Use an abstract class when you want to define a **template** for a group of subclasses, and you have at least *some* implementation code that all subclasses could use. Make the class abstract when you want to guarantee that nobody can make objects of that type.
- Use an interface when you want to define a **role** that other classes can play, regardless of where those classes are in the inheritance tree.

using super

Invoking the superclass version of a method

Q: What if you make a concrete subclass and you need to override a method, but you want the behavior in the superclass version of the method? In other words, what if you don't need to replace the method with an override, but you just want to add to it with some additional specific code.

A: Ahhh... think about the meaning of the word 'extends'. One area of good OO design looks at how to design concrete code that's *meant* to be overridden. In other words, you write method code in, say, an abstract class, that does work that's generic enough to support typical concrete implementations. But, the concrete code isn't enough to handle *all* of the subclass-specific work. So the subclass overrides the method and *extends* it by adding the rest of the code. The keyword super lets you invoke a superclass version of an overridden method, from within the subclass.

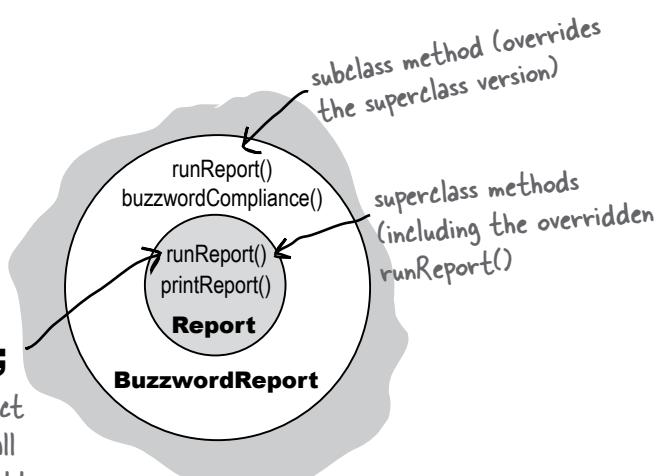
If method code inside a BuzzwordReport subclass says:
super.runReport();

the runReport() method inside the superclass Report will run

super.runReport();

A reference to the subclass object (BuzzwordReport) will always call the subclass version of an overridden method. That's polymorphism. But the subclass code can call super.runReport() to invoke the superclass version.

```
abstract class Report {  
    void runReport() {  
        // set-up report  
    }  
    void printReport() {  
        // generic printing  
    }  
}  
  
class BuzzwordsReport extends Report {  
  
    void runReport() {  
        super.runReport(); ← superclass version of the  
        buzzwordCompliance(); method does important stuff  
        printReport(); ← that subclasses could use  
    }  
    void buzzwordCompliance() {...}  
}
```



The super keyword is really a reference to the superclass portion of an object. When subclass code uses super, as in super.runReport(), the superclass version of the method will run.



BULLET POINTS

- When you don't want a class to be instantiated (in other words, you don't want anyone to make a new object of that class type) mark the class with the **abstract** keyword.
- An abstract class can have both abstract and non-abstract methods.
- If a class has even *one* abstract method, the class must be marked abstract.
- An abstract method has no body, and the declaration ends with a semicolon (no curly braces).
- All abstract methods must be implemented in the first concrete subclass in the inheritance tree.
- Every class in Java is either a direct or indirect subclass of class **Object** (`java.lang.Object`).
- Methods can be declared with **Object** arguments and/or return types.
- You can call methods on an object *only* if the methods are in the class (or interface) used as the *reference* variable type, regardless of the actual *object* type. So, a reference variable of type **Object** can be used only to call methods defined in class **Object**, regardless of the type of the object to which the reference refers.
- A reference variable of type **Object** can't be assigned to any other reference type without a cast. A cast can be used to assign a reference variable of one type to a reference variable of a subtype, but at runtime the cast will fail if the object on the heap is NOT of a type compatible with the cast.
Example: `Dog d = (Dog) x.getObject(aDog);`
- All objects come out of an `ArrayList<Object>` as type **Object** (meaning, they can be referenced only by an **Object** reference variable, unless you use a *cast*).
- Multiple inheritance is not allowed in Java, because of the problems associated with the "Deadly Diamond of Death". That means you can extend only one class (i.e. you can have only one immediate superclass).
- An interface is like a 100% pure abstract class. It defines *only* abstract methods.
- Create an interface using the **interface** keyword instead of the word **class**.
- Implement an interface using the keyword **implements**
Example: `Dog implements Pet`
- Your class can implement multiple interfaces.
- A class that implements an interface *must* implement all the methods of the interface, since **all interface methods are implicitly public and abstract**.
- To invoke the superclass version of a method from a subclass that's overridden the method, use the **super** keyword. Example: `super.runReport();`

Q: There's still something strange here... you never explained how it is that `ArrayList<Dog>` gives back **Dog** references that don't need to be cast, yet the `ArrayList` class uses **Object** in its methods, not **Dog** (or **DotCom** or anything else). What's the special trick going on when you say `ArrayList<Dog>?`

A: You're right for calling it a special trick. In fact it is a special trick that `ArrayList<Dog>` gives back Dogs without you having to do any cast, since it looks like `ArrayList` methods don't know anything about Dogs, or any type besides **Object**.

The short answer is that *the compiler puts in the cast for you!* When you say `ArrayList<Dog>`, there is no special class that has methods to take and return Dog objects, but instead the `<Dog>` is a signal to the compiler that you want the compiler to let you put ONLY Dog objects in and to stop you if you try to add any other type to the list. And since the compiler stops you from adding anything but Dogs to the `ArrayList`, the compiler also knows that it's safe to cast anything that comes out of that `ArrayList` to a Dog reference. In other words, using `ArrayList<Dog>` saves you from having to cast the Dog you get back. But it's much more important than that... because remember, a cast can fail at runtime, and wouldn't you rather have your errors happen at compile time rather than, say, when your customer is using it for something critical?

But there's a lot more to this story, and we'll get into all the details in the Collections chapter.

exercise: What's the Picture?



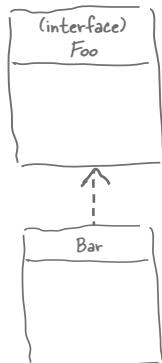
Here's your chance to demonstrate your artistic abilities. On the left you'll find sets of class and interface declarations. Your job is to draw the associated class diagrams on the right. We did the first one for you. Use a dashed line for "implements" and a solid line for "extends".

Given:

What's the Picture?

1) public interface Foo { }
public class Bar implements Foo { }

1)



2) public interface Vinn { }
public abstract class Vout implements Vinn { }

2)

3) public abstract class Muffie implements Whuffle { }
public class Fluffie extends Muffie { }
public interface Whuffle { }

3)

4) public class Zoop { }
public class Boop extends Zoop { }
public class Goop extends Boop { }

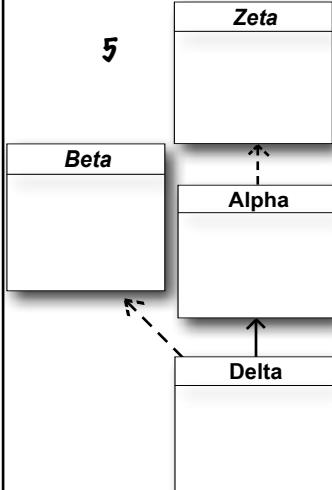
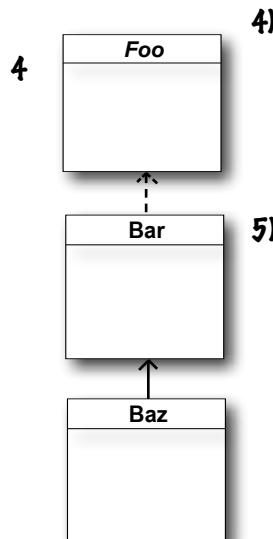
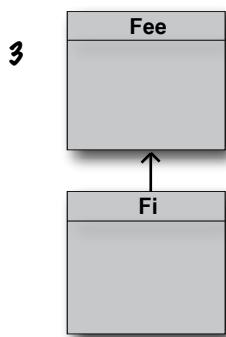
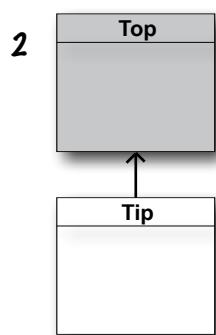
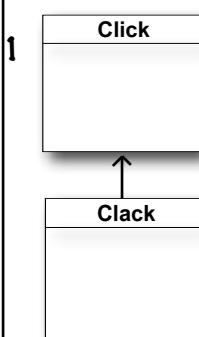
4)

5) public class Gamma extends Delta implements Epsilon { }
public interface Epsilon { }
public interface Beta { }
public class Alpha extends Gamma implements Beta { }
public class Delta { }

5)



Given:



What's the Declaration ?

1) public class Click { }
public class Clack extends Click { }

2)

3)

4)

5)

KEY

	extends
	implements
	class
	interface
	abstract class

puzzle: Pool Puzzle



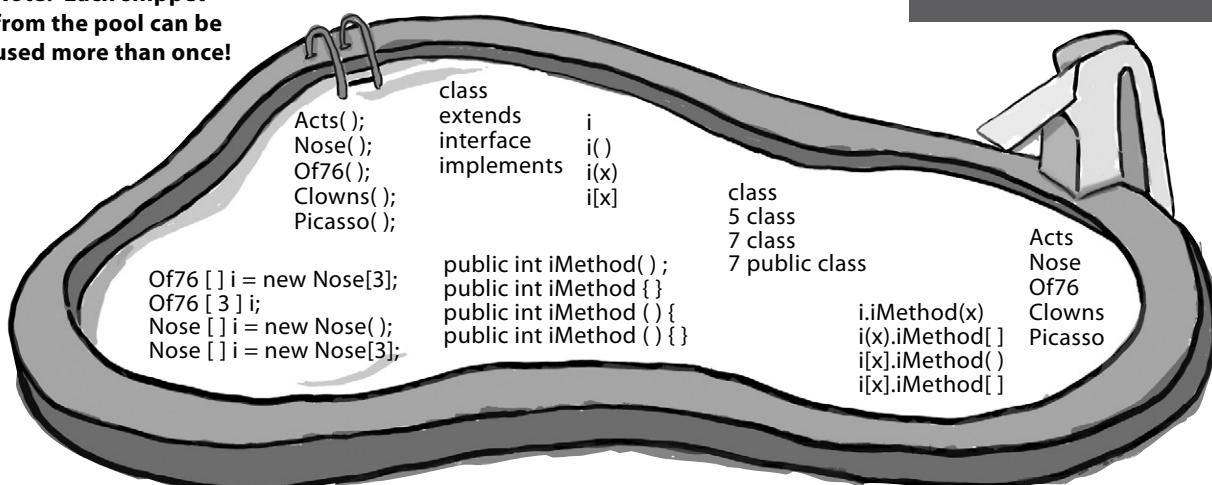
Poo Puzzle



Your **job** is to take code snippets from the pool and place them into the blank lines in the code and output. You **may** use the same snippet more than once, and you won't need to use all the snippets. Your **goal** is to make a set of classes that will compile and run and produce the output listed.

```
_____ Nose {  
_____ }  
  
abstract class Picasso implements _____ {  
_____ return 7;  
}  
  
class _____ { }  
  
class _____ {  
_____ return 5;  
}  
}
```

Note: Each snippet from the pool can be used more than once!



```
public _____ extends Clowns {  
  
public static void main(String [] args) {  
  
i[0] = new _____  
i[1] = new _____  
i[2] = new _____  
for(int x = 0; x < 3; x++) {  
System.out.println(_____  
+ " " + _____.getClass( ) );  
}  
}  
}
```

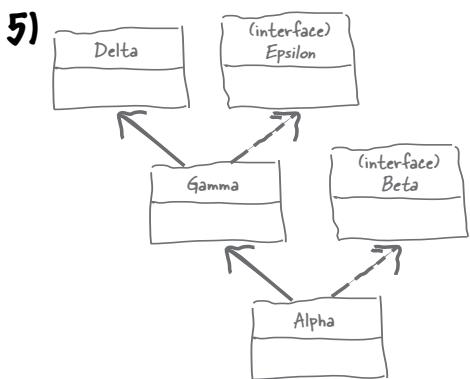
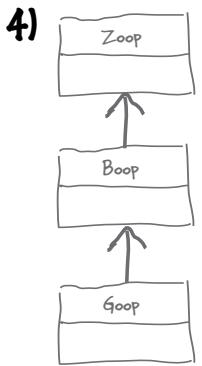
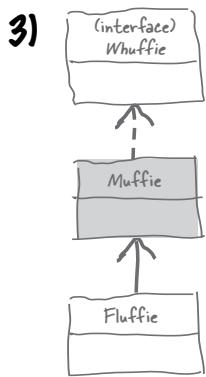
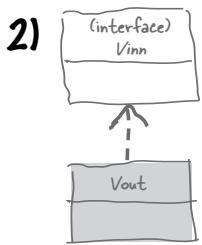
Output

```
File Edit Window Help BeAfraid  
%java _____  
5 class Acts  
7 class Clowns  
Of76
```



Exercise Solutions

What's the Picture?



What's the Declaration?

2) `public abstract class Top { }`
`public class Tip extends Top { }`

3) `public abstract class Fee { }`
`public abstract class Fi extends Fee { }`

4) `public interface Foo { }`
`public class Bar implements Foo { }`
`public class Baz extends Bar { }`

5) `public interface Zeta { }`
`public class Alpha implements Zeta { }`
`public interface Beta { }`
`public class Delta extends Alpha implements Beta { }`

puzzle solution



```
interface Nose {  
    public int iMethod();  
}  
  
abstract class Picasso implements Nose {  
    public int iMethod(){  
        return 7;  
    }  
}  
  
class Clowns extends Picasso {}  
  
class Acts extends Picasso {  
    public int iMethod(){  
        return 5;  
    }  
}
```

```
public class Of76 extends Clowns {  
    public static void main(String [] args) {  
        Nose [ ]i = new Nose [3];  
        i[0] = new Acts();  
        i[1] = new Clowns();  
        i[2] = new Of76();  
        for(int x = 0; x < 3; x++) {  
            System.out.println( i[x].iMethod()  
                + " " + i[x].getClass() );  
        }  
    }  
}
```

Output

```
File Edit Window Help KillTheMime  
%java Of76  
5 class Acts  
7 class Clowns  
7 class Of76
```

9 constructors and garbage collection

Life and Death of an Object



...then he said,
"I can't feel my legs!" and
I said "Joe! Stay with me Joe!"
But it was... too late. The garbage
collector came and... he was gone.
Best object I ever had.

Objects are born and objects die. You're in charge of an object's lifecycle.

You decide when and how to **construct** it. You decide when to **destroy** it. Except you don't actually *destroy* the object yourself, you simply *abandon* it. But once it's abandoned, the heartless **Garbage Collector (gc)** can vaporize it, reclaiming the memory that object was using. If you're gonna write Java, you're gonna create objects. Sooner or later, you're gonna have to let some of them go, or risk running out of RAM. In this chapter we look at how objects are created, where they live while they're alive, and how to keep or abandon them efficiently. That means we'll talk about the heap, the stack, scope, constructors, super constructors, null references, and more. Warning: this chapter contains material about object death that some may find disturbing. Best not to get too attached.

the stack and the heap

The Stack and the Heap: where things live

Before we can understand what really happens when you create an object, we have to step back a bit. We need to learn more about where everything lives (and for how long) in Java. That means we need to learn more about the Stack and the Heap. In Java, we (programmers) care about two areas of memory—the one where objects live (the heap), and the one where method invocations and local variables live (the stack). When a JVM starts up, it gets a chunk of memory from the underlying OS, and uses it to run your Java program. How *much* memory, and whether or not you can tweak it, is dependent on which version of the JVM (and on which platform) you're

running. But usually you *won't* have anything to say about it. And with good programming, you probably won't care (more on that a little later).

We know that all *objects* live on the garbage-collectible heap, but we haven't yet looked at where *variables* live. And where a variable lives depends on what *kind* of variable it is. And by "kind", we don't mean *type* (i.e. primitive or object reference). The two *kinds* of variables whose lives we care about now are *instance* variables and *local* variables. Local variables are also known as *stack* variables, which is a big clue for where they live.

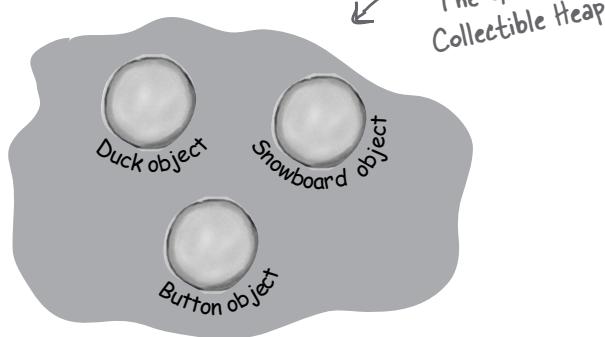
The Stack

Where method invocations and local variables live



The Heap

Where **ALL** objects live



Instance Variables

Instance variables are declared inside a class but not inside a method. They represent the "fields" that each individual object has (which can be filled with different values for each instance of the class). Instance variables live inside the object they belong to.

```
public class Duck {  
    int size;  
}  
  
Every Duck has a "size"  
instance variable.
```

Local Variables

Local variables are declared inside a method, including method parameters. They're temporary, and live only as long as the method is on the stack (in other words, as long as the method has not reached the closing curly brace).

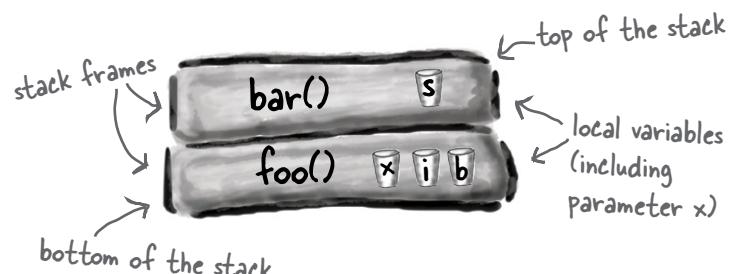
```
public void foo(int x) {  
    int i = x + 3;      The parameter x and  
    boolean b = true;  the variables i and b  
}  
  
are all local variables.
```

Methods are stacked

When you call a method, the method lands on the top of a call stack. That new thing that's actually pushed onto the stack is the *stack frame*, and it holds the state of the method including which line of code is executing, and the values of all local variables.

The method at the *top* of the stack is always the currently-running method for that stack (for now, assume there's only one stack, but in chapter 14 we'll add more.) A method stays on the stack until the method hits its closing curly brace (which means the method's done). If method *foo()* calls method *bar()*, method *bar()* is stacked on top of method *foo()*.

A call stack with two methods



The method on the top of the stack is always the currently-executing method.

```
public void doStuff() {
    boolean b = true;
    go(4);
}

public void go(int x) {
    int z = x + 24;
    crazy();
    // imagine more code here
}

public void crazy() {
    char c = 'a';
}
```

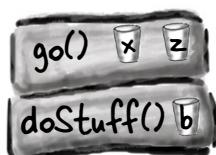
A stack scenario

The code on the left is a snippet (we don't care what the rest of the class looks like) with three methods. The first method (*doStuff()*) calls the second method (*go()*), and the second method calls the third (*crazy()*). Each method declares one local variable within the body of the method, and method *go()* also declares a parameter variable (which means *go()* has two local variables).

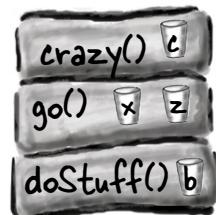
- ① Code from another class calls *doStuff()*, and *doStuff()* goes into a stack frame at the top of the stack. The boolean variable named 'b' goes on the *doStuff()* stack frame.



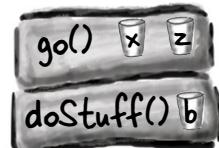
- ② *doStuff()* calls *go()*, *go()* is pushed on top of the stack. Variables 'x' and 'z' are in the *go()* stack frame.



- ③ *go()* calls *crazy()*, *crazy()* is now on the top of the stack, with variable 'c' in the frame.



- ④ *crazy()* completes, and its stack frame is popped off the stack. Execution goes back to the *go()* method, and picks up at the line following the call to *crazy()*.



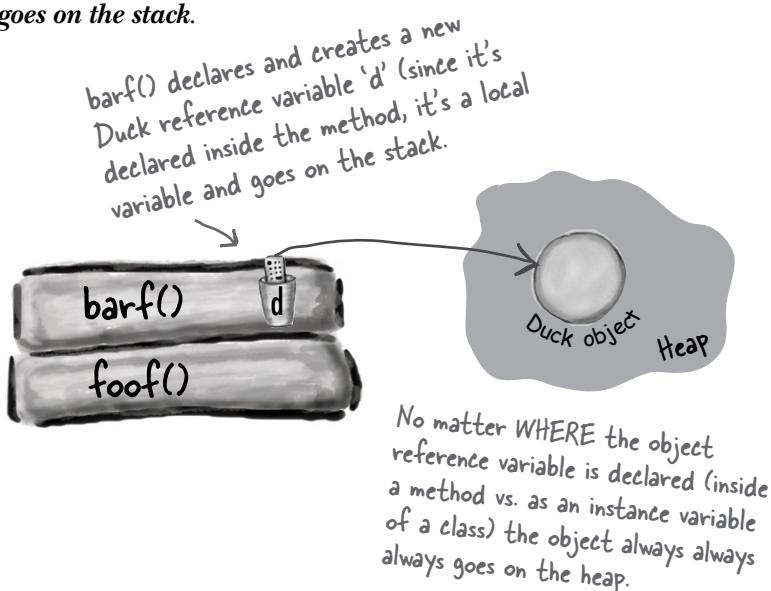
object references on the stack

What about local variables that are objects?

Remember, a non-primitive variable holds a *reference* to an object, not the object itself. You already know where objects live—on the heap. It doesn’t matter where they’re declared or created. *If the local variable is a reference to an object, only the variable (the reference/remote control) goes on the stack.*

The object itself still goes in the heap.

```
public class StackRef {  
    public void foof() {  
        barf();  
    }  
  
    public void barf() {  
        Duck d = new Duck(24);  
    }  
}
```



there are no
Dumb Questions

Q: One more time, WHY are we learning the whole stack/heap thing? How does this help me? Do I really need to learn about it?

A: Knowing the fundamentals of the Java Stack and Heap is crucial if you want to understand variable scope, object creation issues, memory management, threads, and exception handling. We cover threads and exception handling in later chapters but the others you’ll learn in this one. You do not need to know anything about *how* the Stack and Heap are implemented in any particular JVM and/or platform. Everything you need to know about the Stack and Heap is on this page and the previous one. If you nail these pages, all the other topics that depend on your knowing this stuff will go much, much, much easier. Once again, some day you will SO thank us for shoving Stacks and Heaps down your throat.

BULLET POINTS

- ▶ Java has two areas of memory we care about: the Stack and the Heap.
- ▶ Instance variables are variables declared inside a class but outside any method.
- ▶ Local variables are variables declared inside a method or method parameter.
- ▶ All local variables live on the stack, in the frame corresponding to the method where the variables are declared.
- ▶ Object reference variables work just like primitive variables—if the reference is declared as a local variable, it goes on the stack.
- ▶ All objects live in the heap, regardless of whether the reference is a local or instance variable.

If local variables live on the stack, where do instance variables live?

When you say `new CellPhone()`, Java has to make space on the Heap for that CellPhone. But how *much* space? Enough for the object, which means enough to house all of the object's instance variables. That's right, instance variables live on the Heap, inside the object they belong to.

Remember that the *values* of an object's instance variables live inside the object. If the instance variables are all primitives, Java makes space for the instance variables based on the primitive type. An int needs 32 bits, a long 64 bits, etc. Java doesn't care about the value inside primitive variables; the bit-size of an int variable is the same (32 bits) whether the value of the int is 32,000,000 or 32.

But what if the instance variables are *objects*? What if CellPhone HAS-A Antenna? In other words, CellPhone has a reference variable of type Antenna.

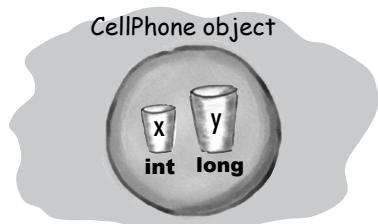
When the new object has instance variables that are object references rather than primitives, the real question is: does the object need space for all of the objects it holds references to? The answer is, *not exactly*. No matter what, Java has to make space for the instance variable *values*. But remember that a reference variable value is not the whole *object*, but merely a *remote control* to the object. So if CellPhone has an instance variable declared as the non-primitive type Antenna, Java makes space within the CellPhone object only for the Antenna's *remote control* (i.e. reference variable) but not the Antenna *object*.

Well then when does the Antenna *object* get space on the Heap? First we have to find out *when* the Antenna object itself is created. That depends on the instance variable declaration. If the instance variable is declared but no object is assigned to it, then only the space for the reference variable (the remote control) is created.

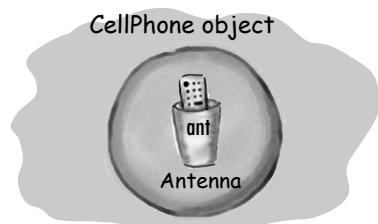
```
private Antenna ant;
```

No actual Antenna object is made on the heap unless or until the reference variable is assigned a new Antenna object.

```
private Antenna ant = new Antenna();
```

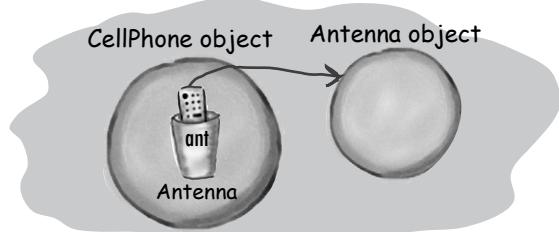


Object with two primitive instance variables.
Space for the variables lives in the object.



Object with one non-primitive instance variable—a reference to an Antenna object, but no actual Antenna object. This is what you get if you declare the variable but don't initialize it with an actual Antenna object.

```
public class CellPhone {  
    private Antenna ant;  
}
```



Object with one non-primitive instance variable, and the Antenna variable is assigned a new Antenna object.

```
public class CellPhone {  
    private Antenna ant = new Antenna();  
}
```

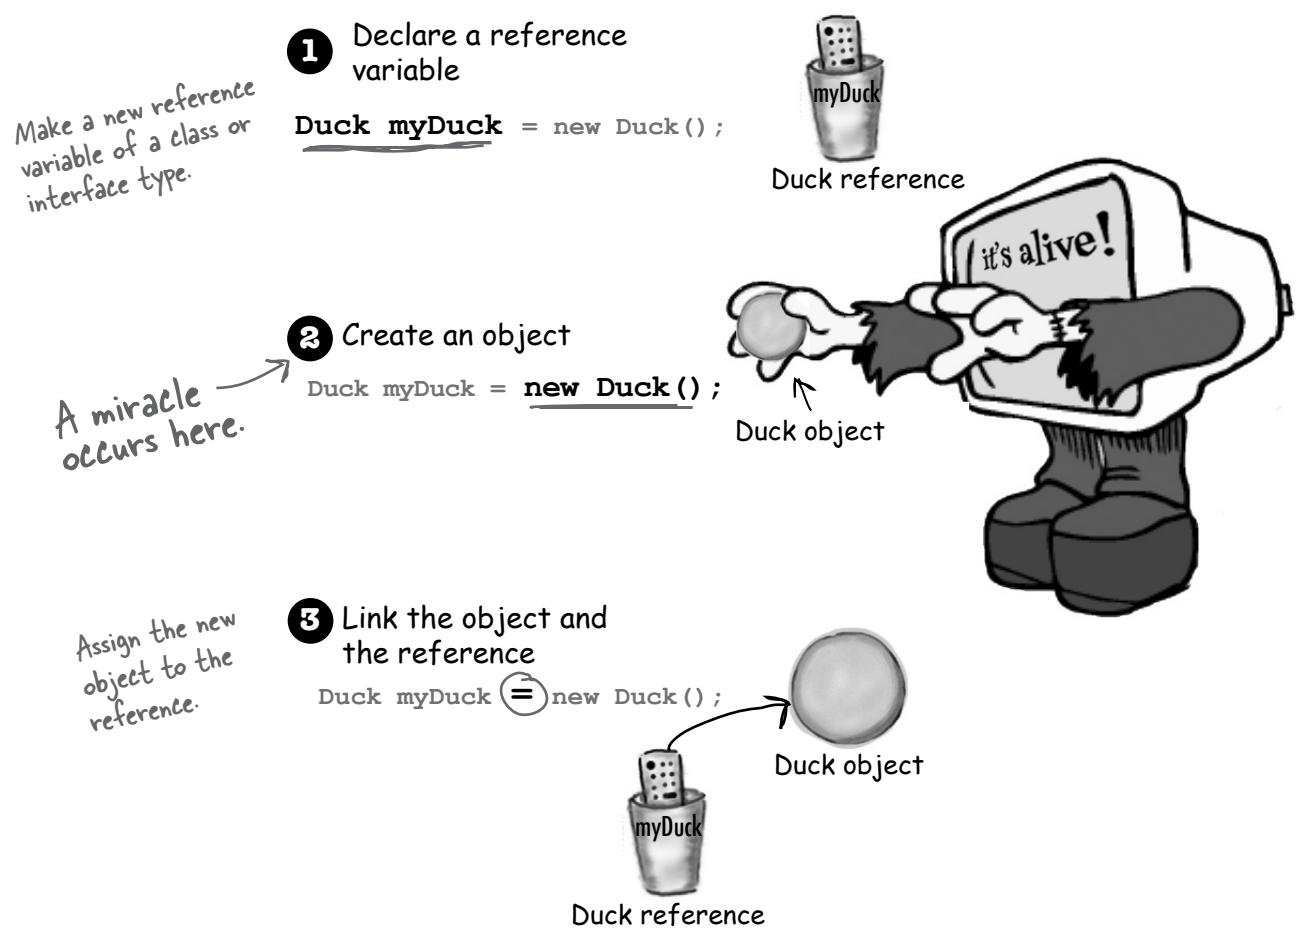
object creation

The miracle of object creation

Now that you know where variables and objects live, we can dive into the mysterious world of object creation. Remember the three steps of object declaration and assignment: declare a reference variable, create an object, and assign the object to the reference.

But until now, step two—where a miracle occurs and the new object is “born”—has remained a Big Mystery. Prepare to learn the facts of object life. *Hope you’re not squeamish.*

Review the 3 steps of object declaration, creation and assignment:



Are we calling a method named Duck()?

Because it sure *looks* like it.

Duck myDuck = new Duck();

No.

We're calling the **Duck constructor**.

A constructor *does* look and feel a lot like a method, but it's not a method. It's got the code that runs when you say `new`. In other words, *the code that runs when you instantiate an object*.

The only way to invoke a constructor is with the keyword `new` followed by the class name. The JVM finds that class and invokes the constructor in that class. (OK, technically this isn't the *only* way to invoke a constructor. But it's the only way to do it from *outside* a constructor. You *can* call a constructor from within another constructor, with restrictions, but we'll get into all that later in the chapter.)

But where is the constructor?

If we didn't write it, who did?

A constructor has the code that runs when you instantiate an object. In other words, the code that runs when you say `new` on a class type.

Every class you create has a constructor, even if you don't write it yourself.

You can write a constructor for your class (we're about to do that), but if you don't, *the compiler writes one for you!*

Here's what the compiler's default constructor looks like:

```
public Duck() {  
}
```

Notice something missing? How is this different from a method?

Where's the return type?
If this were a method,
you'd need a return type
between "public" and
"Duck()".

Its name is the same as the class name. That's mandatory.

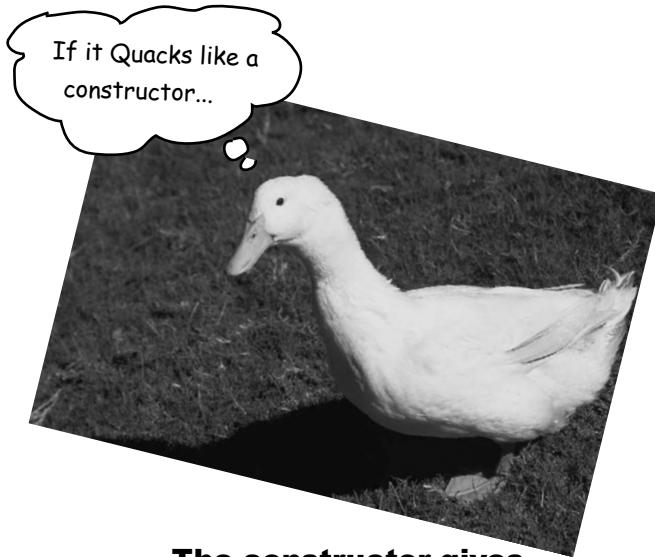
public Duck() {
 // constructor code goes here
}

constructing a new Duck

Construct a Duck

The key feature of a constructor is that it runs *before* the object can be assigned to a reference. That means you get a chance to step in and do things to get the object ready for use. In other words, before anyone can use the remote control for an object, the object has a chance to help construct itself. In our Duck constructor, we're not doing anything useful, but it still demonstrates the sequence of events.

```
public class Duck {  
  
    public Duck() {  
        System.out.println("Quack");  
    }  
  
    }  
  
    Constructor code.  
    ↗
```



The constructor gives you a chance to step into the middle of `new`.

```
public class UseADuck {  
  
    public static void main (String[] args) {  
        Duck d = new Duck();  
    }  
  
    }  
  
    This calls the Duck  
    constructor.  
    ↗
```

A screenshot of a terminal window. The window title is "File Edit Window Help Quack". The command "% java UseADuck" is entered, followed by the output "Quack".

Sharpen your pencil

A constructor lets you jump into the middle of the object creation step—into the middle of `new`. Can you imagine conditions where that would be useful? Which of these might be useful in a Car class constructor, if the Car is part of a Racing Game? Check off the ones that you came up with a scenario for.

- Increment a counter to track how many objects of this class type have been made.
- Assign runtime-specific state (data about what's happening NOW).
- Assign values to the object's important instance variables.
- Get and save a reference to the object that's *creating* the new object.
- Add the object to an ArrayList.
- Create HAS-A objects.
- _____ (your idea here)

Initializing the state of a new Duck

Most people use constructors to initialize the state of an object. In other words, to make and assign values to the object's instance variables.

```
public Duck() {
    size = 34;
}
```

That's all well and good when the Duck class *developer* knows how big the Duck object should be. But what if we want the programmer who is *using* Duck to decide how big a particular Duck should be?

Imagine the Duck has a size instance variable, and you want the programmer using your Duck class to set the size of the new Duck. How could you do it?

Well, you could add a setSize() setter method to the class. But that leaves the Duck temporarily without a size*, and forces the Duck user to write *two* statements—one to create the Duck, and one to call the setSize() method. The code below uses a setter method to set the initial size of the new Duck.

```
public class Duck {
    int size;           ← instance variable

    public Duck() {
        System.out.println("Quack");   ← constructor
    }

    public void setSize(int newSize) { ← setter method
        size = newSize;
    }
}
```

```
public class UseADuck {

    public static void main (String[] args) {
        Duck d = new Duck();
        d.setSize(42);
    }
}
```

*Instance variables do have a default value. 0 or 0.0 for numeric primitives, false for booleans, and null for references.

there are no
Dumb Questions

Q: Why do you need to write a constructor if the compiler writes one for you?

A: If you need code to help initialize your object and get it ready for use, you'll have to write your own constructor. You might, for example, be dependent on input from the user before you can finish making the object ready. There's another reason you might have to write a constructor, even if you don't need any constructor code yourself. It has to do with your superclass constructor, and we'll talk about that in a few minutes.

Q: How can you tell a constructor from a method? Can you also have a method that's the same name as the class?

A: Java lets you declare a method with the same name as your class. That doesn't make it a constructor, though. The thing that separates a method from a constructor is the return type. Methods *must* have a return type, but constructors *cannot* have a return type.

Q: Are constructors inherited? If you don't provide a constructor but your superclass does, do you get the superclass constructor instead of the default?

A: Nope. Constructors are not inherited. We'll look at that in just a few pages.

initializing object state

Using the constructor to initialize important Duck state*

If an object shouldn't be used until one or more parts of its state (instance variables) have been initialized, don't let anyone get ahold of a Duck object until you're finished initializing! It's usually way too risky to let someone make—and get a reference to—a new Duck object that isn't quite ready for use until that someone turns around and calls the `setSize()` method. How will the Duck-user even *know* that he's required to call the setter method after making the new Duck?

The best place to put initialization code is in the constructor. And all you need to do is make a constructor with arguments.

Let the user make a new Duck and set the Duck's size all in one call. The call to new. The call to the Duck constructor.

```
public class Duck {  
    int size;  
  
    public Duck(int duckSize) {  
        System.out.println("Quack");  
        size = duckSize;  
        System.out.println("size is " + size);  
    }  
}
```

Add an int parameter to the Duck constructor.

Use the argument value to set the size instance variable.

```
public class UseADuck {  
  
    public static void main (String[] args) {  
        Duck d = new Duck(42);  
    }  
}  
  
This time there's only one statement. We make the new Duck and set its size in one statement.
```

Pass a value to the constructor.

```
File Edit Window Help Honk  
% java UseADuck  
Quack  
size is 42
```

*Not to imply that not all Duck state is not unimportant.

Make it easy to make a Duck

Be sure you have a no-arg constructor

What happens if the Duck constructor takes an argument? Think about it. On the previous page, there's only *one* Duck constructor—and it takes an int argument for the *size* of the Duck. That might not be a big problem, but it does make it harder for a programmer to create a new Duck object, especially if the programmer doesn't *know* what the size of a Duck should be. Wouldn't it be helpful to have a default size for a Duck, so that if the user doesn't know an appropriate size, he can still make a Duck that works?

Imagine that you want Duck users to have TWO options for making a Duck—one where they supply the Duck size (as the constructor argument) and one where they don't specify a size and thus get your default Duck size.

You can't do this cleanly with just a single constructor. Remember, if a method (or constructor—same rules) has a parameter, you *must* pass an appropriate argument when you invoke that method or constructor. You can't just say, "If someone doesn't pass anything to the constructor, then use the default size", because they won't even be able to compile without sending an int argument to the constructor call. You *could* do something clunky like this:

```
public class Duck {
    int size;

    public Duck(int newSize) {
        if (newSize == 0) {
            size = 27;
        } else {
            size = newSize;
        }
    }
}
```

If the parameter value is zero, give the new Duck a default size, otherwise use the parameter value for the size. NOT a very good solution.

But that means the programmer making a new Duck object has to *know* that passing a "0" is the protocol for getting the default Duck size. Pretty ugly. What if the other programmer doesn't know that? Or what if he really *does* want a zero-size Duck? (Assuming a zero-sized Duck is allowed. If you don't want zero-sized Duck objects, put validation code in the constructor to prevent it.) The point is, it might not always be possible to distinguish between a genuine "I want zero for the size" constructor argument and a "I'm sending zero so you'll give me the default size, whatever that is" constructor argument.

You really want TWO ways to make a new Duck:

```
public class Duck2 {
    int size;

    public Duck2() {
        // supply default size
        size = 27;
    }

    public Duck2(int duckSize) {
        // use duckSize parameter
        size = duckSize;
    }
}
```

To make a Duck when you know the size:

```
Duck2 d = new Duck2(15);
```

To make a Duck when you do not know the size:

```
Duck2 d2 = new Duck2();
```

So this two-options-to-make-a-Duck idea needs two constructors. One that takes an int and one that doesn't. **If you have more than one constructor in a class, it means you have Overloaded constructors.**

Doesn't the compiler always make a no-arg constructor for you? No!

You might think that if you write *only* a constructor with arguments, the compiler will see that you don't have a no-arg constructor, and stick one in for you. But that's not how it works. The compiler gets involved with constructor-making *only if you don't say anything at all about constructors.*

If you write a constructor that takes arguments, and you still want a no-arg constructor, you'll have to build the no-arg constructor yourself!

As soon as *you* provide a constructor, ANY kind of constructor, the compiler backs off and says, "OK Buddy, looks like you're in charge of constructors now."

If you have more than one constructor in a class, the constructors MUST have different argument lists.

The argument list includes the order and types of the arguments. As long as they're different, you can have more than one constructor. You can do this with methods as well, but we'll get to that in another chapter.

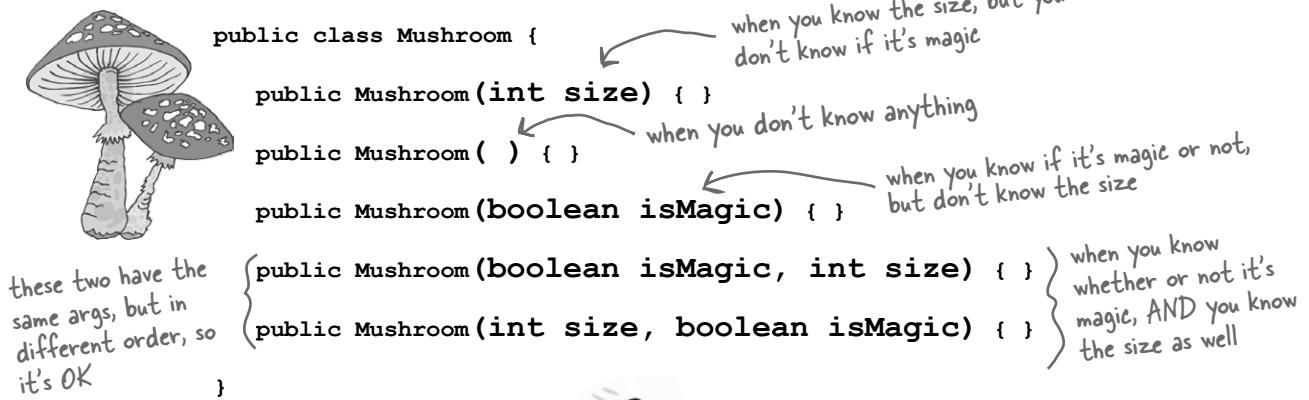


Overloaded constructors means you have more than one constructor in your class.

To compile, each constructor must have a different argument list!

The class below is legal because all five constructors have different argument lists. If you had two constructors that took only an int, for example, the class wouldn't compile. What you name the parameter variable doesn't count. It's the variable *type* (int, Dog, etc.) and *order* that matters. You *can* have two constructors that have identical types, *as long as the order is different*. A constructor that takes a String followed by an int, is *not* the same as one that takes an int followed by a String.

Five different constructors
means five different ways to
make a new mushroom.



BULLET POINTS

- ▶ Instance variables live within the object they belong to, on the Heap.
- ▶ If the instance variable is a reference to an object, both the reference and the object it refers to are on the Heap.
- ▶ A constructor is the code that runs when you say `new` on a class type.
- ▶ A constructor must have the same name as the class, and must *not* have a return type.
- ▶ You can use a constructor to initialize the state (i.e. the instance variables) of the object being constructed.
- ▶ If you don't put a constructor in your class, the compiler will put in a default constructor.
- ▶ The default constructor is always a no-arg constructor.
- ▶ If you put a constructor—any constructor—in your class, the compiler will not build the default constructor.
- ▶ If you want a no-arg constructor, and you've already put in a constructor with arguments, you'll have to build the no-arg constructor yourself.
- ▶ Always provide a no-arg constructor if you can, to make it easy for programmers to make a working object. Supply default values.
- ▶ Overloaded constructors means you have more than one constructor in your class.
- ▶ Overloaded constructors must have different argument lists.
- ▶ You cannot have two constructors with the same argument lists. An argument list includes the order and/or type of arguments.
- ▶ Instance variables are assigned a default value, even when you don't explicitly assign one. The default values are 0/0/false for primitives, and null for references.

overloaded constructors

Sharpen your pencil

Match the `new Duck()` call with the constructor that runs when that Duck is instantiated. We did the easy one to get you started.

```
public class TestDuck {  
  
    public static void main(String[] args){  
  
        int weight = 8;  
        float density = 2.3F;  
        String name = "Donald";  
        long[] feathers = {1,2,3,4,5,6};  
        boolean canFly = true;  
        int airspeed = 22;  
  
        Duck[] d = new Duck[7];  
  
        d[0] = new Duck();  
  
        d[1] = new Duck(density, weight);  
  
        d[2] = new Duck(name, feathers);  
  
        d[3] = new Duck(canFly);  
  
        d[4] = new Duck(3.3F, airspeed);  
  
        d[5] = new Duck(false);  
  
        d[6] = new Duck(airspeed, density);  
    }  
}
```

```
class Duck {  
  
    int pounds = 6;  
    float floatability = 2.1F;  
    String name = "Generic";  
    long[] feathers = {1,2,3,4,5,6,7};  
    boolean canFly = true;  
    int maxSpeed = 25;  
  
    public Duck() {  
        System.out.println("type 1 duck");  
    }  
  
    public Duck(boolean fly) {  
        canFly = fly;  
        System.out.println("type 2 duck");  
    }  
  
    public Duck(String n, long[] f) {  
        name = n;  
        feathers = f;  
        System.out.println("type 3 duck");  
    }  
  
    public Duck(int w, float f) {  
        pounds = w;  
        floatability = f;  
        System.out.println("type 4 duck");  
    }  
  
    public Duck(float density, int max) {  
        floatability = density;  
        maxSpeed = max;  
        System.out.println("type 5 duck");  
    }  
}
```

Q: Earlier you said that it's good to have a no-argument constructor so that if people call the no-arg constructor, we can supply default values for the "missing" arguments. But aren't there times when it's impossible to come up with defaults? Are there times when you should not have a no-arg constructor in your class?

A: You're right. There are times when a no-arg constructor doesn't make sense. You'll see this in the Java API—some classes don't have a no-arg constructor. The Color class, for example, represents a... color. Color objects are used to, for example, set or change the color of a screen font or GUI button. When you make a Color instance, that instance is of a particular color (you know, Death-by-Chocolate Brown, Blue-Screen-of-Death Blue, Scandalous Red, etc.). If you make a Color object, you must specify the color in some way.

```
Color c = new Color(3,45,200);
```

248 chapter 9

(We're using three ints for RGB values here. We'll get into using Color later, in the Swing chapters.) Otherwise, what would you get? The Java API programmers could have decided that if you call a no-arg Color constructor you'll get a lovely shade of mauve. But good taste prevailed. If you try to make a Color without supplying an argument:

```
Color c = new Color();
```

The compiler freaks out because it can't find a matching no-arg constructor in the Color class.

```
File Edit Window Help StopBeingStupid  
cannot resolve symbol  
:constructor Color()  
location: class  
java.awt.Color  
Color c = new Color();  
^  
1 error
```

Nanoreview: four things to remember about constructors

- ① A constructor is the code that runs when somebody says `new` on a class type

```
Duck d = new Duck();
```

- ② A constructor must have the same name as the class, and **no** return type

```
public Duck(int size) { }
```

- ③ If you don't put a constructor in your class, the compiler puts in a default constructor. The default constructor is always a no-arg constructor.

```
public Duck() { }
```

- ④ You can have more than one constructor in your class, as long as the argument lists are different. Having more than one constructor in a class means you have overloaded constructors.

```
public Duck() { }

public Duck(int size) { }

public Duck(String name) { }

public Duck(String name, int size) { }
```

Doing all the Brain Barbells has been shown to produce a 42% increase in neuron size. And you know what they say, "Big neurons..."



What about superclasses?

When you make a Dog, should the Canine constructor run too?

If the superclass is abstract, should it even have a constructor?

We'll look at this on the next few pages, so stop now and think about the implications of constructors and superclasses.

*there are no
Dumb Questions*

Q: Do constructors have to be `public`?

A: No. Constructors can be `public`, `protected`, `private`, or `default` (which means no access modifier at all). We'll look more at `default` access in chapter 16 and appendix B.

Q: How could a `private` constructor ever be useful? Nobody could ever call it, so nobody could ever make a new object!

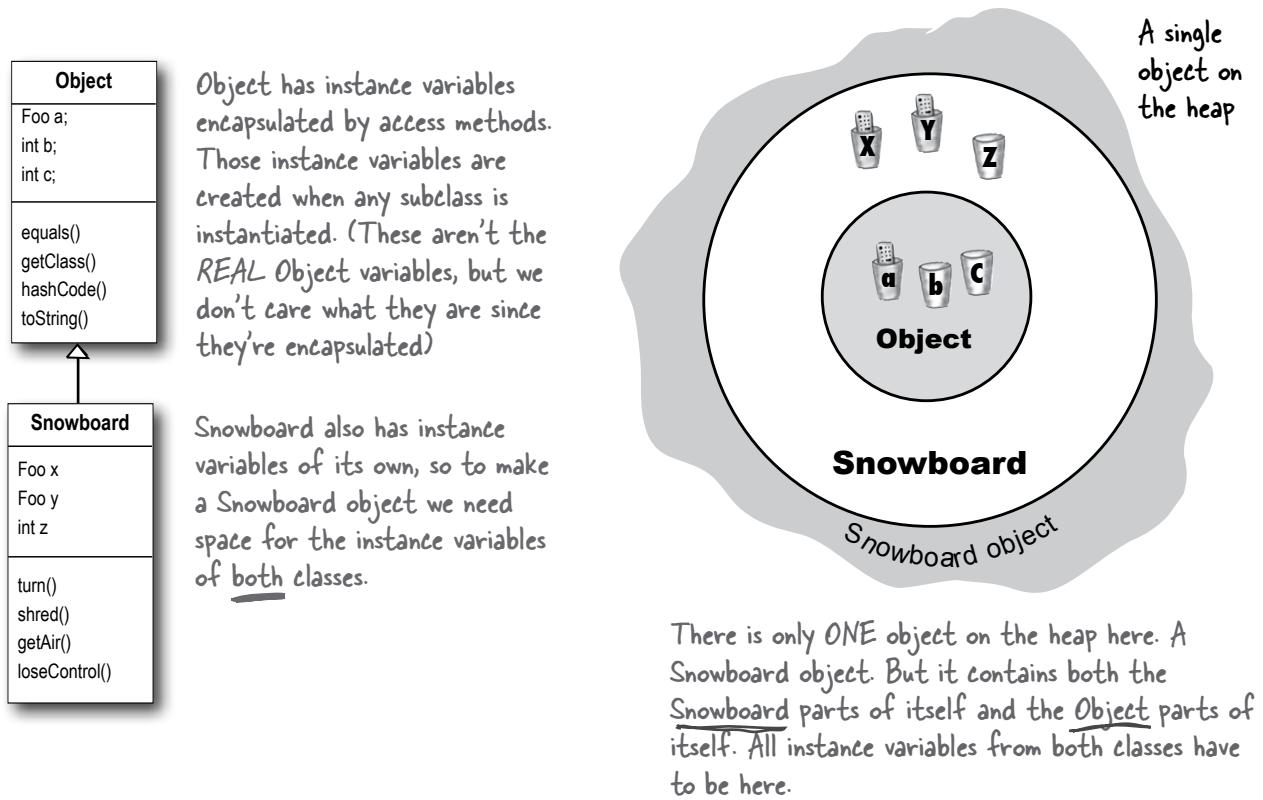
A: But that's not exactly right. Marking something `private` doesn't mean *nobody* can access it, it just means that *nobody outside the class* can access it. Bet you're thinking "Catch 22". Only code from the *same* class as the class-with-private-constructor can make a new object from that class, but without first making an object, how do you ever get to run code from that class in the first place? How do you ever get to anything in that class? *Patience grasshopper*. We'll get there in the next chapter.

space for an object's superclass parts

Wait a minute... we never DID talk about superclasses and inheritance and how that all fits in with constructors.

Here's where it gets fun. Remember from the last chapter, the part where we looked at the Snowboard object wrapping around an inner core representing the Object portion of the Snowboard class? The Big Point there was that every object holds not just its *own* declared instance variables, but also *everything from its superclasses* (which, at a minimum, means class Object, since *every* class extends Object).

So when an object is created (because somebody said `new`; there is *no other way* to create an object other than someone, somewhere saying `new` on the class type), the object gets space for *all* the instance variables, from all the way up the inheritance tree. Think about it for a moment... a superclass might have setter methods encapsulating a private variable. But that variable has to live *somewhere*. When an object is created, it's almost as though *multiple* objects materialize—the object being new'd and one object per each superclass. Conceptually, though, it's much better to think of it like the picture below, where the object being created has *layers* of itself representing each superclass.



The role of superclass constructors in an object's life.

All the constructors in an object's inheritance tree must run when you make a new object.

Let that sink in.

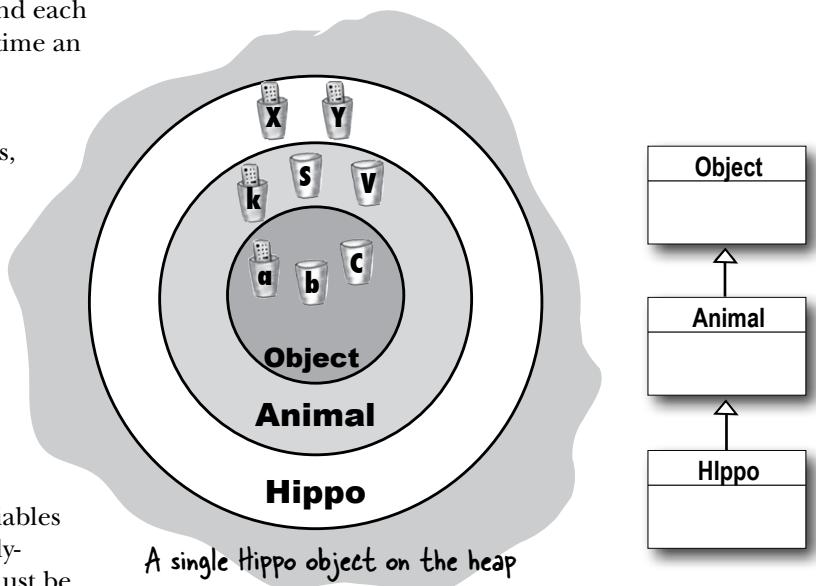
That means every superclass has a constructor (because every class has a constructor), and each constructor up the hierarchy runs at the time an object of a subclass is created.

Saying `new` is a Big Deal. It starts the whole constructor chain reaction. And yes, even abstract classes have constructors. Although you can never say `new` on an abstract class, an abstract class is still a superclass, so its constructor runs when someone makes an instance of a concrete subclass.

The super constructors run to build out the superclass parts of the object. Remember, a subclass might inherit methods that depend on superclass state (in other words, the value of instance variables in the superclass). For an object to be fully-formed, all the superclass parts of itself must be fully-formed, and that's why the super constructor *must* run. All instance variables from every class in the inheritance tree have to be declared and initialized. Even if `Animal` has instance variables that `Hippo` doesn't inherit (if the variables are private, for example), the `Hippo` still depends on the `Animal` methods that *use* those variables.

When a constructor runs, it immediately calls its superclass constructor, all the way up the chain until you get to the class `Object` constructor.

On the next few pages, you'll learn how superclass constructors are called, and how you can call them yourself. You'll also learn what to do if your superclass constructor has arguments!



A new Hippo object also IS-A Animal and IS-A Object. If you want to make a Hippo, you must also make the Animal and Object parts of the Hippo.

This all happens in a process called Constructor Chaining.

object construction

Making a Hippo means making the Animal and Object parts too...

```
public class Animal {  
    public Animal() {  
        System.out.println("Making an Animal");  
    }  
}
```

```
public class Hippo extends Animal {  
    public Hippo() {  
        System.out.println("Making a Hippo");  
    }  
}
```

```
public class TestHippo {  
    public static void main (String[] args) {  
        System.out.println("Starting...");  
        Hippo h = new Hippo();  
    }  
}
```



What's the real output? Given the code on the left, what prints out when you run TestHippo? A or B?

(the answer is at the bottom of the page)

A
File Edit Window Help Swear
% java TestHippo
Starting...
Making an Animal
Making a Hippo

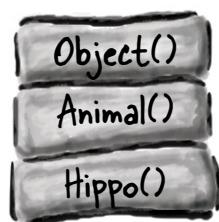
B
File Edit Window Help Swear
% java TestHippo
Starting...
Making a Hippo
Making an Animal

- ① Code from another class says `new Hippo()` and the `Hippo()` constructor goes into a stack frame at the top of the stack.

- ② `Hippo()` invokes the superclass constructor which pushes the `Animal()` constructor onto the top of the stack.

- ③ `Animal()` invokes the superclass constructor which pushes the `Object()` constructor onto the top of the stack, since `Object` is the superclass of `Animal`.

- ④ `Object()` completes, and its stack frame is popped off the stack. Execution goes back to the `Animal()` constructor, and picks up at the line following `Animal`'s call to its superclass constructor



The first one, A. The `Hippo()` constructor is invoked first, but it's the `Animal` constructor that finishes first.

How do you invoke a superclass constructor?

You might think that somewhere in, say, a Duck constructor, if Duck extends Animal you'd call Animal(). But that's not how it works:

```
public class Duck extends Animal {
    int size;

    public Duck(int newSize) {
        BAD! → Animal(); ← NO! This is not legal!
        size = newSize;
    }
}
```

The only way to call a super constructor is by calling *super()*. That's right—*super()* calls the *super constructor*.

What are the odds?

```
public class Duck extends Animal {
    int size;

    public Duck(int newSize) {
        super(); ← You just say super()
        size = newSize;
    }
}
```

A call to *super()* in your constructor puts the superclass constructor on the top of the Stack. And what do you think that superclass constructor does? *Calls its superclass constructor*. And so it goes until the Object constructor is on the top of the Stack. Once *Object()* finishes, it's popped off the Stack and the next thing down the Stack (the subclass constructor that called *Object()*) is now on top. *That constructor* finishes and so it goes until the original constructor is on the top of the Stack, where *it* can now finish.

And how is it that we've gotten away without doing it?

You probably figured that out.

Our good friend the compiler puts in a call to *super()* if you don't.

So the compiler gets involved in constructor-making in two ways:

① If you don't provide a constructor

The compiler puts one in that looks like:

```
public ClassName() {
    super();
}
```

② If you do provide a constructor but you do not put in the call to *super()*

The compiler will put a call to *super()* in each of your overloaded constructors.* The compiler-supplied call looks like:

```
super();
```

It always looks like that. The compiler-inserted call to *super()* is always a no-arg call. If the superclass has overloaded constructors, only the no-arg one is called.

*Unless the constructor calls another overloaded constructor (you'll see that in a few pages).

object lifecycle

Can the child exist before the parents?

If you think of a superclass as the parent to the subclass child, you can figure out which has to exist first. *The superclass parts of an object have to be fully-formed (completely built) before the subclass parts can be constructed.* Remember, the subclass object might depend on things it inherits from the superclass, so it's important that those inherited things be finished. No way around it. The superclass constructor must finish before its subclass constructor.

Look at the Stack series on page 252 again, and you can see that while the Hippo constructor is the *first* to be invoked (it's the first thing on the Stack), it's the *last* one to complete! Each subclass constructor immediately invokes its own superclass constructor, until the Object constructor is on the top of the Stack. Then Object's constructor completes and we bounce back down the Stack to Animal's constructor. Only after Animal's constructor completes do we finally come back down to finish the rest of the Hippo constructor. For that reason:

The call to super() must be the *first* statement in each constructor!*



Possible constructors for class Boop	
<input checked="" type="checkbox"/> public Boop() { super(); }	These are OK because the programmer explicitly coded the call to super(), as the first statement.
<input checked="" type="checkbox"/> public Boop(int i) { super(); size = i; }	
<input type="checkbox"/> public Boop(int i) { size = i; super(); }	BAD!! This won't compile! You can't explicitly put the call to super() below anything else.

*There's an exception to this rule; you'll learn it on page 256.

Superclass constructors with arguments

What if the superclass constructor has arguments? Can you pass something in to the `super()` call? Of course. If you couldn't, you'd never be able to extend a class that didn't have a no-arg constructor. Imagine this scenario: all animals have a name. There's a `getName()` method in class `Animal` that returns the value of the `name` instance variable. The instance variable is marked private, but the subclass (in this case, `Hippo`) inherits the `getName()` method. So here's the problem: `Hippo` has a `getName()` method (through inheritance), but does not have the `name` instance variable. `Hippo` has to depend on the `Animal` part of himself to keep the `name` instance variable, and return it when someone calls `getName()` on a `Hippo` object. But... how does the `Animal` part get the name? The only reference `Hippo` has to the `Animal` part of himself is through `super()`, so that's the place where `Hippo` sends the `Hippo`'s name up to the `Animal` part of himself, so that the `Animal` part can store it in the private `name` instance variable.

```
public abstract class Animal {
    private String name; ← All animals (including
                           subclasses) have a name

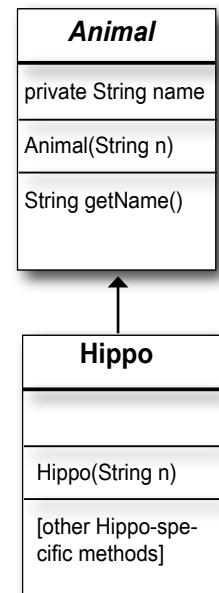
    public String getName() { ← A getter method that
        return name;
    }

    public Animal(String theName) {
        name = theName;
    }
}

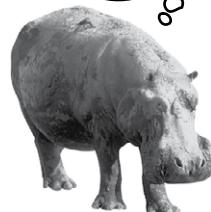
public class Hippo extends Animal {

    public Hippo(String name) {
        super(name); ← Hippo constructor takes a name
    }
}

public class MakeHippo {
    public static void main(String[] args) {
        Hippo h = new Hippo("Buffy"); ← Make a Hippo, passing the
                                       name "Buffy" to the Hippo
                                       constructor. Then call the
                                       Hippo's inherited getName()
        System.out.println(h.getName());
    }
}
```



The Animal part of me needs to know my name, so I take a name in my own Hippo constructor, then pass the name to super()



```

File Edit Window Help Hide
%java MakeHippo
Buffy
  
```

calling overloaded constructors

Invoking one overloaded constructor from another

What if you have overloaded constructors that, with the exception of handling different argument types, all do the same thing? You know that you don't want *duplicate* code sitting in each of the constructors (pain to maintain, etc.), so you'd like to put the bulk of the constructor code (including the call to `super()`) in only *one* of the overloaded constructors. You want whichever constructor is first invoked to call The Real Constructor and let The Real Constructor finish the job of construction. It's simple: just say `this()`. Or `this(aString)`. Or `this(27, x)`. In other words, just imagine that the keyword `this` is a reference to **the current object**.

You can say `this()` only within a constructor, and it must be the first statement in the constructor!

But that's a problem, isn't it? Earlier we said that `super()` must be the first statement in the constructor. Well, that means you get a choice.

Every constructor can have a call to super() or this(), but never both!

```
import java.awt.Color;
class Mini extends Car {
```

```
    Color color;
    public Mini() {
        this(Color.RED); ←
    }
```

The no-arg constructor supplies a default Color and calls the overloaded Real Constructor (the one that calls `super()`).

```
    public Mini(Color c) {
        super("Mini"); ←
        color = c;
        // more initialization
    }
```

This is The Real Constructor that does The Real Work of initializing the object (including the call to `super()`)

```
    public Mini(int size) {
        this(Color.RED);
        super(size); ←
    }
}
```

Won't work!! Can't have `super()` and `this()` in the same constructor, because they each must be the first statement!

Use this() to call a constructor from another overloaded constructor in the same class.

The call to `this()` can be used only in a constructor, and must be the first statement in a constructor.

A constructor can have a call to `super()` OR `this()`, but never both!

```
File Edit Window Help Drive
javac Mini.java
Mini.java:16: call to super must
be first statement in constructor
    super(); ^
```

Sharpen your pencil

Some of the constructors in the SonOfBoo class will not compile. See if you can recognize which constructors are not legal. Match the compiler errors with the SonOfBoo constructors that caused them, by drawing a line from the compiler error to the "bad" constructor.

```
public class Boo {
    public Boo(int i) { }
    public Boo(String s) { }
    public Boo(String s, int i) { }
}
```

```
class SonOfBoo extends Boo {

    public SonOfBoo() {
        super("boo");
    }

    public SonOfBoo(int i) {
        super("Fred");
    }

    public SonOfBoo(String s) {
        super(42);
    }

    public SonOfBoo(int i, String s) {
    }

    public SonOfBoo(String a, String b, String c) {
        super(a,b);
    }

    public SonOfBoo(int i, int j) {
        super("man", j);
    }

    public SonOfBoo(int i, int x, int y) {
        super(i, "star");
    }
}
```



```
File Edit Window Help
%javac SonOfBoo.java
cannot resolve symbol
symbol : constructor Boo
(java.lang.String,java.lang.String)
```

```
File Edit Window Help Yadayadaya
%javac SonOfBoo.java
cannot resolve symbol
symbol : constructor Boo
(int,java.lang.String)
```

```
File Edit Window Help ImNotListening
%javac SonOfBoo.java
cannot resolve symbol
symbol:constructor Boo()
```

object lifespan

Now we know how an object is born, but how long does an object live?

An *object*'s life depends entirely on the life of references referring to it. If the reference is considered "alive", the object is still alive on the Heap. If the reference dies (and we'll look at what that means in just a moment), the object will die.

So if an object's life depends on the reference variable's life, how long does a *variable* live?

That depends on whether the variable is a *local* variable or an *instance* variable. The code below shows the life of a local variable. In the example, the variable is a primitive, but variable lifetime is the same whether it's a primitive or reference variable.

```
public class TestLifeOne {  
  
    public void read() {  
        int s = 42; ← 's' is scoped to the read()  
        sleep();     method, so it can't be used  
    }  
  
    public void sleep() {  
        s = 7;  
    } ← BAD!! Not legal to  
         use 's' here!  
  
    sleep() can't see the 's' variable. Since  
    it's not in sleep()'s own Stack frame,  
    sleep() doesn't know anything about it.  
    sleep() doesn't know anything about it.  
  
    The variable 's' is alive, but in scope only within the  
    read() method. When sleep() completes and read() is  
    still see 's'. When read() completes and is popped off  
    the Stack, 's' is dead. Pushing up digital daisies.  
}
```

① A local variable lives only within the method that declared the variable.

```
public void read() {  
    int s = 42;  
    // 's' can be used only  
    // within this method.  
    // When this method ends,  
    // 's' disappears completely.  
}
```

Variable 's' can be used *only* within the *read()* method. In other words, **the variable is in scope only within its own method**. No other code in the class (or any other class) can see 's'.

② An instance variable lives as long as the object does. If the object is still alive, so are its instance variables.

```
public class Life {  
    int size;  
  
    public void setSize(int s) {  
        size = s;  
        // 's' disappears at the  
        // end of this method,  
        // but 'size' can be used  
        // anywhere in the class  
    }  
}
```

Variable 's' (this time a method parameter) is in scope only within the *setSize()* method. But instance variable *size* is scoped to the life of the *object* as opposed to the life of the *method*.

The difference between **life** and **scope** for local variables:

Life

A local variable is *alive* as long as its Stack frame is on the Stack. In other words, *until the method completes*.

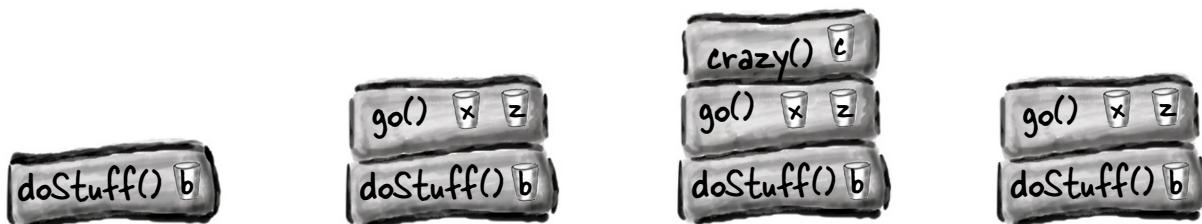
Scope

A local variable is in *scope* only within the method in which the variable was declared. When its own method calls another, the variable is alive, but not in scope until its method resumes. *You can use a variable only when it is in scope.*

```
public void doStuff() {
    boolean b = true;
    go(4);
}

public void go(int x) {
    int z = x + 24;
    crazy();
    // imagine more code here
}

public void crazy() {
    char c = 'a';
}
```



- 1 *doStuff()* goes on the Stack. Variable 'b' is alive and in scope.

- 2 *go()* plops on top of the Stack. 'x' and 'z' are alive and in scope, and 'b' is alive but *not* in scope.

- 3 *crazy()* is pushed onto the Stack, with 'c' now alive and in scope. The other three variables are alive but out of scope.

- 4 *crazy()* completes and is popped off the Stack, so 'c' is out of scope and dead. When *go()* resumes where it left off, 'x' and 'z' are both alive and back in scope. Variable 'b' is still alive but out of scope (until *go()* completes).

While a local variable is alive, its state persists. As long as method *doStuff()* is on the Stack, for example, the 'b' variable keeps its value. But the 'b' variable can be used only while *doStuff()*'s Stack frame is at the top. In other words, you can use a local variable *only* while that local variable's method is actually running (as opposed to waiting for higher Stack frames to complete).

What about reference variables?

The rules are the same for primitives and references. A reference variable can be used only when it's in scope, which means you can't use an object's remote control unless you've got a reference variable that's in scope. The *real* question is,

"How does *variable* life affect *object* life?"

An object is alive as long as there are live references to it. If a reference variable goes out of scope but is still alive, the object it *refers* to is still alive on the Heap. And then you have to ask... "What happens when the Stack frame holding the reference gets popped off the Stack at the end of the method?"

If that was the *only* live reference to the object, the object is now abandoned on the Heap. The reference variable disintegrated with the Stack frame, so the abandoned object is now, *officially*, toast. The trick is to know the point at which an object becomes *eligible for garbage collection*.

Once an object is eligible for garbage collection (GC), you don't have to worry about reclaiming the memory that object was using. If your program gets low on memory, GC will destroy some or all of the eligible objects, to keep you from running out of RAM. You can still run out of memory, but *not* before all eligible objects have been hauled off to the dump. Your job is to make sure that you abandon objects (i.e., make them eligible for GC) when you're done with them, so that the garbage collector has something to reclaim. If you hang on to objects, GC can't help you and you run the risk of your program dying a painful out-of-memory death.

An object's life has no value, no meaning, no point, unless somebody has a reference to it.

If you can't get to it, you can't ask it to do anything and it's just a big fat waste of bits.

But if an object is unreachable, the Garbage Collector will figure that out. Sooner or later, that object's goin' down.



An object becomes eligible for GC when its last live reference disappears.

Three ways to get rid of an object's reference:

- ① The reference goes out of scope, permanently

```
void go() {  
    Life z = new Life();  
}
```

reference 'z' dies at
end of method

- ② The reference is assigned another object

```
Life z = new Life();  
z = new Life();
```

the first object is abandoned
when z is 'reprogrammed' to
a new object.

- ③ The reference is explicitly set to null

```
Life z = new Life();  
z = null;
```

the first object is abandoned
when z is 'deprogrammed'.

Object-killer #1

Reference goes out of scope, permanently.



```
public class StackRef {
    public void foof() {
        barf();
    }

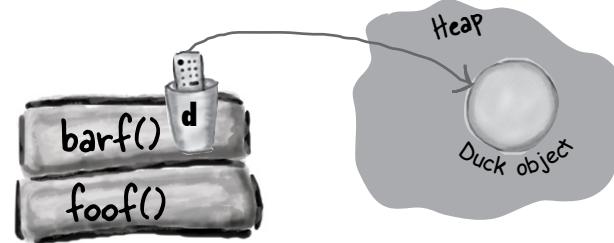
    public void barf() {
        Duck d = new Duck();
    }
}
```



- 1 *foof()* is pushed onto the Stack, no variables are declared.

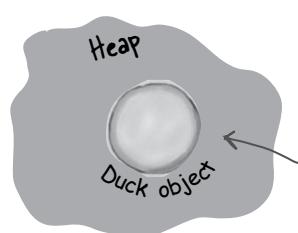


- 2 *barf()* is pushed onto the Stack, where it declares a reference variable, and creates a new object assigned to that reference. The object is created on the Heap, and the reference is alive and in scope.



The new Duck goes on the Heap, and as long as *barf()* is running, the 'd' reference is alive and in scope, so the Duck is considered alive.

- 3 *barf()* completes and pops off the Stack. Its frame disintegrates, so 'd' is now dead and gone. Execution returns to *foof()*, but *foof()* can't use 'd'.



Uh-oh. The 'd' variable went away when the *barf()* Stack frame was blown off the stack, so the Duck is abandoned. Garbage-collector bait.

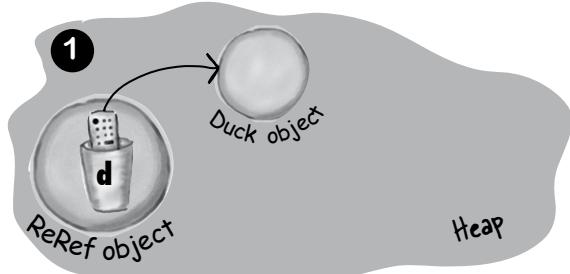
object lifecycle

Object-killer #2

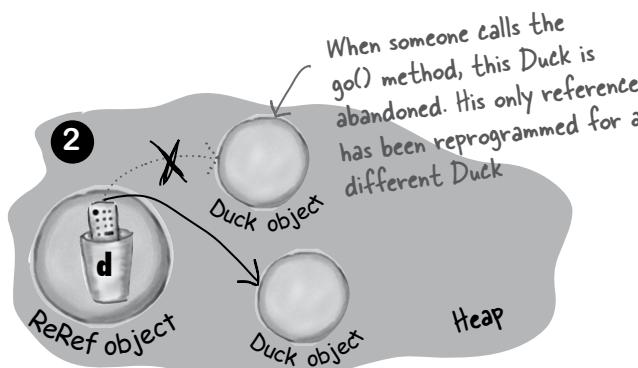
Assign the reference
to another object



```
public class ReRef {  
  
    Duck d = new Duck();  
  
    public void go() {  
        d = new Duck();  
    }  
}
```



The new Duck goes on the Heap, referenced by 'd'. Since 'd' is an instance variable, the Duck will live as long as the ReRef object that instantiated it is alive. Unless...



'd' is assigned a new Duck object, leaving the original (first) Duck object abandoned. That first Duck is now as good as dead.



Object-killer #3

Explicitly set the reference to null



```
public class ReRef {
    Duck d = new Duck();
    public void go() {
        d = null;
    }
}
```

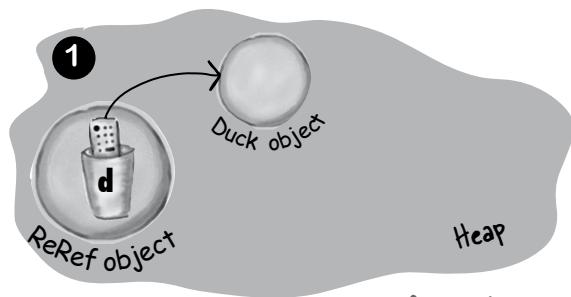
The meaning of null

When you set a reference to **null**, you're deprogramming the remote control.

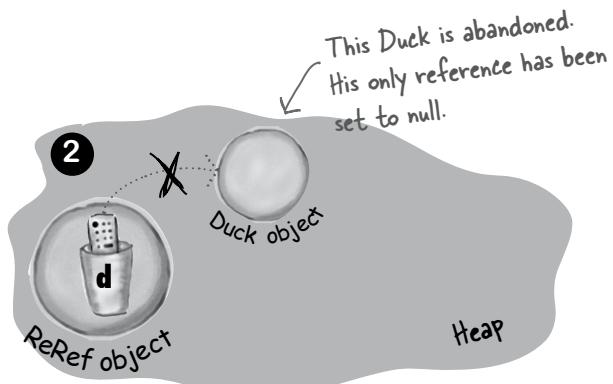
In other words, you've got a remote control, but no TV at the other end. A null reference has bits representing 'null' (we don't know or care what those bits are, as long as the JVM knows).

If you have an unprogrammed remote control, in the real world, the buttons don't do anything when you press them. But in Java, you can't press the buttons (i.e. use the dot operator) on a null reference, because the JVM knows (this is a runtime issue, not a compiler error) that you're expecting a bark but there's no Dog there to do it!

If you use the dot operator on a null reference, you'll get a `NullPointerException` at runtime. You'll learn all about Exceptions in the Risky Behavior chapter.

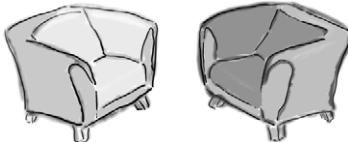


The new Duck goes on the Heap, referenced by 'd'. Since 'd' is an instance variable, the Duck will live as long as the ReRef object that instantiated it is alive. Unless...



'd' is set to null, which is just like having a remote control that isn't programmed to anything. You're not even allowed to use the dot operator on 'd' until it's reprogrammed (assigned an object).

Fireside Chats



Tonight's Talk: **An instance variable and a local variable discuss life and death (with remarkable civility)**

Instance Variable

I'd like to go first, because I tend to be more important to a program than a local variable. I'm there to support an object, usually throughout the object's entire life. After all, what's an object without *state*? And what is state? Values kept in *instance variables*.

No, don't get me wrong, I do understand your role in a method, it's just that your life is so short. So temporary. That's why they call you guys "temporary variables".

My apologies. I understand completely.

I never really thought about it like that. What are you doing while the other methods are running and you're waiting for your frame to be the top of the Stack again?

Local Variable

I appreciate your point of view, and I certainly appreciate the value of object state and all, but I don't want folks to be misled. Local variables are *really* important. To use your phrase, "After all, what's an object without *behavior*?" And what is behavior? Algorithms in methods. And you can bet your bits there'll be some *local variables* in there to make those algorithms work.

Within the local-variable community, the phrase "temporary variable" is considered derogatory. We prefer "local", "stack", "automatic", or "Scope-challenged".

Anyway, it's true that we don't have a long life, and it's not a particularly *good* life either. First, we're shoved into a Stack frame with all the other local variables. And then, if the method we're part of calls another method, another frame is pushed on top of us. And if *that* method calls *another* method... and so on. Sometimes we have to wait forever for all the other methods on top of the Stack to complete so that our method can run again.

Nothing. Nothing at all. It's like being in stasis—that thing they do to people in science fiction movies when they have to travel long distances. Suspended animation, really. We just sit there on hold. As long as our frame is still there, we're safe and the value we hold is secure, but it's a mixed blessing when our

Instance Variable

We saw an educational video about it once. Looks like a pretty brutal ending. I mean, when that method hits its ending curly brace, the frame is literally *blown* off the Stack! Now *that's* gotta hurt.

I live on the Heap, with the objects. Well, not *with* the objects, actually *in* an object. The object whose state I store. I have to admit life can be pretty luxurious on the Heap. A lot of us feel guilty, especially around the holidays.

OK, hypothetically, yes, if I'm an instance variable of the Collar and the Collar gets GC'd, then the Collar's instance variables would indeed be tossed out like so many pizza boxes. But I was told that this almost never happens.

They let us *drink*?

Local Variable

frame gets to run again. On the one hand, we get to be active again. On the other hand, the clock starts ticking again on our short lives. The more time our method spends running, the closer we get to the end of the method. We *all* know what happens then.

Tell me about it. In computer science they use the term *popped* as in “the frame was popped off the Stack”. That makes it sound fun, or maybe like an extreme sport. But, well, you saw the footage. So why don't we talk about you? I know what my little Stack frame looks like, but where do *you* live?

But you don't *always* live as long as the object who declared you, right? Say there's a Dog object with a Collar instance variable. Imagine *you're* an instance variable of the Collar object, maybe a reference to a Buckle or something, sitting there all happy inside the Collar object who's all happy inside the Dog object. But... what happens if the Dog wants a new Collar, or *nulls* out its Collar instance variable? That makes the Collar object eligible for GC. So... if *you're* an instance variable inside the Collar, and the whole Collar is abandoned, what happens to *you*?

And you believed it? That's what they say to keep us motivated and productive. But aren't you forgetting something else? What if you're an instance variable inside an object, and that object is referenced *only* by a *local* variable? If I'm the only reference to the object you're in, when I go, you're coming with me. Like it or not, our fates may be connected. So I say we forget about all this and go get drunk while we still can. Carpe RAM and all that.

exercise: Be the Garbage Collector



BE the Garbage Collector

Which of the lines of code on the right, if added to the class on the left at point A, would cause exactly one additional object to be eligible for the Garbage Collector? (Assume that point A (`//call more methods`) will execute for a long time, giving the Garbage Collector time to do its stuff.)

```
public class GC {  
    public static GC doStuff() {  
        GC newGC = new GC();  
        doStuff2(newGC);  
        return newGC;  
    }  
  
    public static void main(String [] args) {  
        GC gc1;  
        GC gc2 = new GC();  
        GC gc3 = new GC();  
        GC gc4 = gc3;  
        gc1 = doStuff();  
  
        // call more methods  
    }  
  
    public static void doStuff2(GC copyGC) {  
        GC localGC = copyGC;  
    }  
}  
  
1   copyGC = null;  
2   gc2 = null;  
3   newGC = gc3;  
4   gc1 = null;  
5   newGC = null;  
6   gc4 = null;  
7   gc3 = gc2;  
8   gc1 = gc4;  
9   gc3 = null;
```



Popular Objects

```

class Bees {
    Honey [] beeHA;
}

class Raccoon {
    Kit k;
    Honey rh;
}

class Kit {
    Honey kh;
}

class Bear {
    Honey hunny;
}

public class Honey {
    public static void main(String [] args) {
        Honey honeyPot = new Honey();
        Honey [] ha = {honeyPot, honeyPot, honeyPot, honeyPot};
        Bees b1 = new Bees();
        b1.beeHA = ha;
        Bear [] ba = new Bear[5];
        for (int x=0; x < 5; x++) {
            ba[x] = new Bear();
            ba[x].hunny = honeyPot;
        }
        Kit k = new Kit();
        k.kh = honeyPot;
        Raccoon r = new Raccoon(); ←
        r.rh = honeyPot;
        r.k = k;
        k = null;
    } // end of main
}

```

Here's a new Raccoon object!

Here's its reference variable 'r'.

In this code example, several new objects are created. Your challenge is to find the object that is 'most popular', i.e. the one that has the most reference variables referring to it. Then list how *many* total references there are for that object, and what they are! We'll start by pointing out one of the new objects, and its reference variable.

Good Luck!

puzzle: Five Minute Mystery



Five-Minute Mystery



“We’ve run the simulation four times, and the main module’s temperature consistently drifts out of nominal towards cold”, Sarah said, exasperated. “We installed the new temp-bots last week. The readings on the radiator bots, designed to cool the living quarters, seem to be within spec, so we’ve focused our analysis on the heat retention bots, the bots that help to warm the quarters.” Tom sighed, at first it had seemed that nano-technology was going to really put them ahead of schedule. Now, with only five weeks left until launch, some of the orbiter’s key life support systems were still not passing the simulation gauntlet.

“What ratios are you simulating?”, Tom asked.

“Well if I see where you’re going, we already thought of that”, Sarah replied. “Mission control will not sign off on critical systems if we run them out of spec. We are required to run the v3 radiator bot’s SimUnits in a 2:1 ratio with the v2 radiator’s SimUnits”, Sarah continued. “Overall, the ratio of retention bots to radiator bots is supposed to run 4:3.”

“How’s power consumption Sarah?”, Tom asked. Sarah paused, “Well that’s another thing, power consumption is running higher than anticipated. We’ve got a team tracking that down too, but because the nanos are wireless it’s been hard to isolate the power consumption of the radiators from the retention bots.” “Overall power consumption ratios”, Sarah continued, “are designed to run 3:2 with the radiators pulling more power from the wireless grid.”

“OK Sarah”, Tom said “Let’s take a look at some of the simulation initiation code. We’ve got to find this problem, and find it quick!”

```
import java.util.*;
class V2Radiator {
    V2Radiator(ArrayList list) {
        for(int x=0; x<5; x++) {
            list.add(new SimUnit("V2Radiator"));
        }
    }
}

class V3Radiator extends V2Radiator {
    V3Radiator(ArrayList llist) {
        super(llist);
        for(int g=0; g<10; g++) {
            llist.add(new SimUnit("V3Radiator"));
        }
    }
}

class RetentionBot {
    RetentionBot(ArrayList rlist) {
        rlist.add(new SimUnit("Retention"));
    }
}
```

Five-Minute Mystery continued...

```
public class TestLifeSupportSim {
    public static void main(String [] args) {
        ArrayList aList = new ArrayList();
        V2Radiator v2 = new V2Radiator(aList);
        V3Radiator v3 = new V3Radiator(aList);
        for(int z=0; z<20; z++) {
            RetentionBot ret = new RetentionBot(aList);
        }
    }
}

class SimUnit {
    String botType;
    SimUnit(String type) {
        botType = type;
    }
    int powerUse() {
        if ("Retention".equals(botType)) {
            return 2;
        } else {
            return 4;
        }
    }
}
```

Tom gave the code a quick look and a small smile crept across his lips. I think I've found the problem Sarah, and I bet I know by what percentage your power usage readings are off too!

What did Tom suspect? How could he guess the power readings errors, and what few lines of code could you add to help debug this program?

object lifecycle

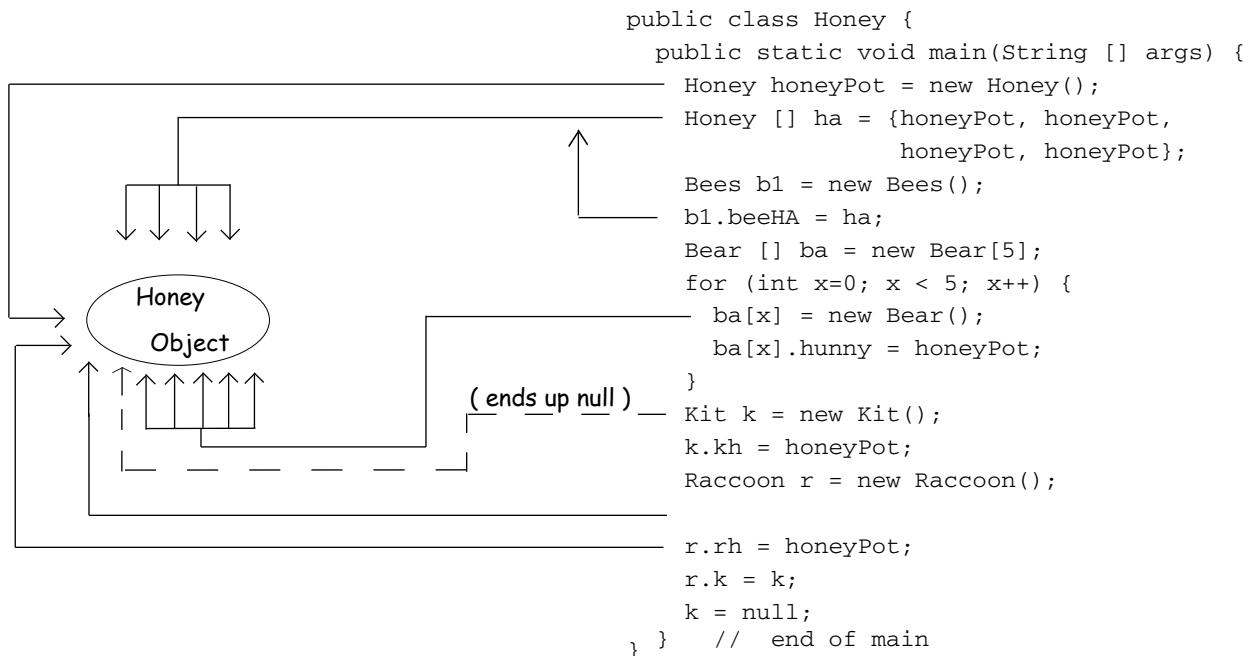


G.C.

- 1 copyGC = null; No - this line attempts to access a variable that is out of scope.
- 2 gc2 = null; OK - gc2 was the only reference variable referring to that object.
- 3 newGC = gc3; No - another out of scope variable.
- 4 gc1 = null; OK - gc1 had the only reference because newGC is out of scope.
- 5 newGC = null; No - newGC is out of scope.
- 6 gc4 = null; No - gc3 is still referring to that object.
- 7 gc3 = gc2; No - gc4 is still referring to that object.
- 8 gc1 = gc4; OK - Reassigning the only reference to that object.
- 9 gc3 = null; No - gc4 is still referring to that object.

Popular Objects

It probably wasn't too hard to figure out that the Honey object first referred to by the honeyPot variable is by far the most 'popular' object in this class. But maybe it was a little trickier to see that all of the variables that point from the code to the Honey object refer to the **same object!** There are a total of 12 active references to this object right before the main() method completes. The k.kh variable is valid for a while, but k gets nulled at the end. Since r.k still refers to the Kit object, r.k.kh (although never explicitly declared), refers to the object!





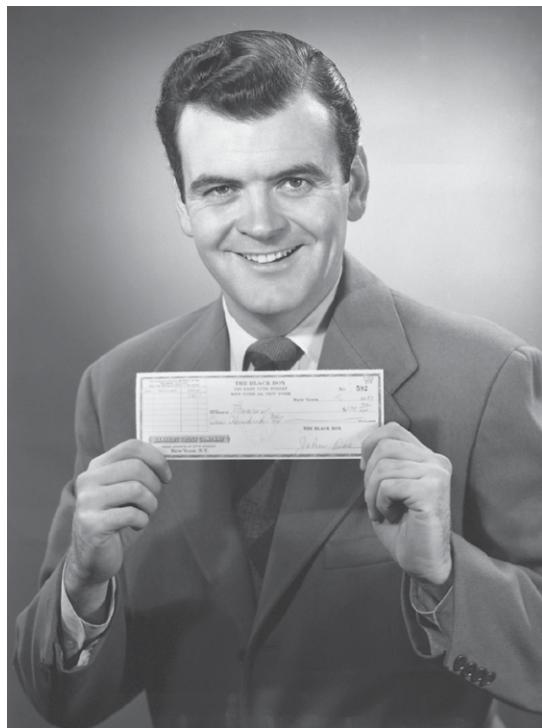
Five-Minute Mystery Solution

Tom noticed that the constructor for the V2Radiator class took an ArrayList. That meant that every time the V3Radiator constructor was called, it passed an ArrayList in its super() call to the V2Radiator constructor. That meant that an extra five V2Radiator SimUnits were created. If Tom was right, total power use would have been 120, not the 100 that Sarah's expected ratios predicted.

Since all the Bot classes create SimUnits, writing a constructor for the SimUnit class, that printed out a line everytime a SimUnit was created, would have quickly highlighted the problem!

10 numbers and statics

Numbers Matter



Do the Math. But there's more to working with numbers than just doing primitive arithmetic. You might want to get the absolute value of a number, or round a number, or find the larger of two numbers. You might want your numbers to print with exactly two decimal places, or you might want to put commas into your large numbers to make them easier to read. And what about working with dates? You might want to print dates in a variety of ways, or even *manipulate* dates to say things like, "add three weeks to today's date". And what about parsing a String into a number? Or turning a number into a String? You're in luck. The Java API is full of handy number-tweaking methods ready and easy to use. But most of them are **static**, so we'll start by learning what it means for a variable or method to be static, including constants in Java—static *final* variables.

Math methods

MATH methods: as close as you'll ever get to a *global* method

Except there's no global *anything* in Java. But think about this: what if you have a method whose behavior doesn't depend on an instance variable value. Take the round() method in the Math class, for example. It does the same thing every time—rounds a floating point number (the argument to the method) to the nearest integer. Every time. If you had 10,000 instances of class Math, and ran the round(42.2) method, you'd get an integer value of 42. Every time. In other words, the method acts on the argument, but is never affected by an instance variable state. The only value that changes the way the round() method runs is the argument passed to the method!

Doesn't it seem like a waste of perfectly good heap space to make an instance of class Math simply to run the round() method? And what about *other* Math methods like min(), which takes two numerical primitives and returns the smaller of the two. Or max(). Or abs(), which returns the absolute value of a number.

These methods never use instance variable values. In fact the Math class doesn't *have* any instance variables. So there's nothing to be gained by making an instance of class Math. So guess what? You don't have to. As a matter of fact, you can't.

If you try to make an instance of class Math:

```
Math mathObject = new Math();
```

You'll get this error:

```
File Edit Window Help IwasToldThereWouldBeNoMath
% javac TestMath
TestMath.java:3: Math() has private
access in java.lang.Math
    Math mathObject = new Math();
                           ^
1 error
```

← This error shows that the Math constructor is marked private! That means you can NEVER say 'new' on the Math class to make a new Math object.

Methods in the Math class don't use any instance variable values. And because the methods are 'static', you don't need to have an instance of Math. All you need is the Math class.

```
int x = Math.round(42.2);
int y = Math.min(56,12);
int z = Math.abs(-343);
```



These methods never use instance variables, so their behavior doesn't need to know about a specific object.

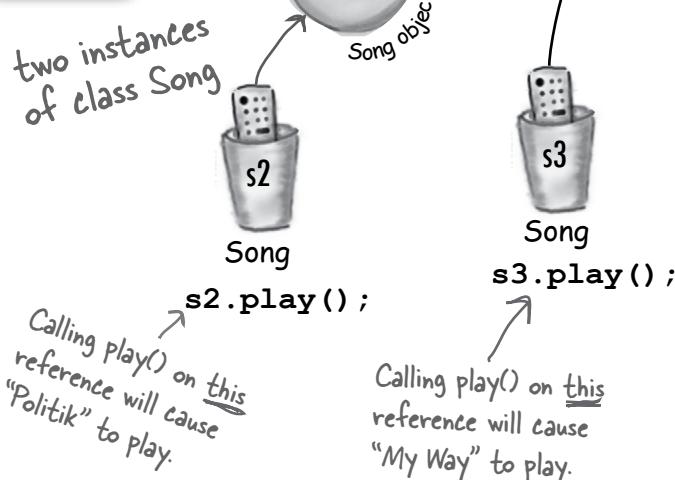
The difference between regular (non-static) and static methods

Java is object-oriented, but once in a while you have a special case, typically a utility method (like the Math methods), where there is no need to have an instance of the class. The keyword **static** lets a method run *without any instance of the class*. A static method means “behavior not dependent on an instance variable, so no instance/object is required. Just the class.”

regular (non-static) method

```
public class Song {
    String title;           ← Instance variable value affects
                           the behavior of the play()
    public Song(String t) { method.
        title = t;
    }
    public void play() {
        SoundPlayer player = new SoundPlayer();
        player.playSound(title);
    }
}
```

Song
title
play()



static method

```
public static int min(int a, int b) {
    //returns the lesser of a and b
}
```

Math
min()
max()
abs()
...

No instance variables.
The method behavior
doesn't change with
instance variable state.

Math.min(42, 36);

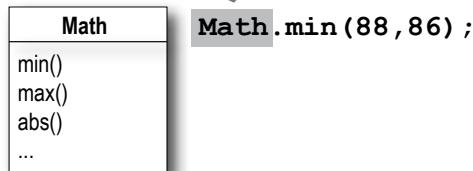
Use the Class name, rather
than a reference variable
name.



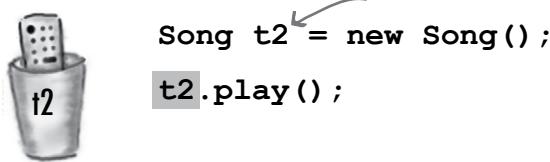
NO OBJECTS!!
Absolutely NO OBJECTS
anywhere in this picture!

static methods

Call a static method using a class name



Call a non-static method using a reference variable name



What it means to have a class with static methods.

Often (although not always), a class with static methods is not meant to be instantiated. In Chapter 8 we talked about abstract classes, and how marking a class with the **abstract** modifier makes it impossible for anyone to say ‘new’ on that class type. In other words, *it’s impossible to instantiate an abstract class*.

But you can restrict other code from instantiating a *non-abstract* class by marking the constructor **private**. Remember, a *method* marked private means that only code from within the class can invoke the method. A *constructor* marked private means essentially the same thing—only code from within the class can invoke the constructor. Nobody can say ‘new’ from *outside* the class. That’s how it works with the Math class, for example. The constructor is private, you cannot make a new instance of Math. The compiler knows that your code doesn’t have access to that private constructor.

This does *not* mean that a class with one or more static methods should never be instantiated. In fact, every class you put a main() method in is a class with a static method in it!

Typically, you make a main() method so that you can launch or test another class, nearly always by instantiating a class in main, and then invoking a method on that new instance.

So you’re free to combine static and non-static methods in a class, although even a single non-static method means there must be *some* way to make an instance of the class. The only ways to get a new object are through ‘new’ or deserialization (or something called the Java Reflection API that we don’t go into). No other way. But exactly *who* says new can be an interesting question, and one we’ll look at a little later in this chapter.

Static methods can't use non-static (instance) variables!

Static methods run without knowing about any particular instance of the static method's class. And as you saw on the previous pages, there might not even *be* any instances of that class. Since a static method is called using the *class* (*Math.random()*) as opposed to an *instance reference* (*t2.play()*), a static method can't refer to any instance variables of the class. The static method doesn't know *which* instance's variable value to use.

If you try to compile this code:

```
public class Duck {
    private int size;

    public static void main (String[] args) {
        System.out.println("Size of duck is " + size);
    }

    public void setSize(int s) {
        size = s;
    }

    public int getSize() {
        return size;
    }
}
```

Which Duck?
Whose size?

If there's a Duck on
the heap somewhere, we
don't know about it.

You'll get this error:

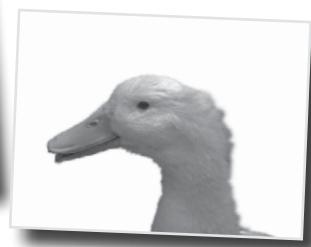
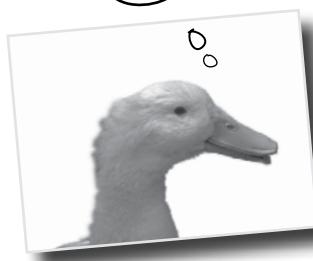
```
File Edit Window Help Quack
% javac Duck.java
Duck.java:6: non-static variable
size cannot be referenced from a
static context

    System.out.println("Size
of duck is " + size);
^
```

If you try to use an instance variable from inside a static method, the compiler thinks, "I don't know which object's instance variable you're talking about!" If you have ten Duck objects on the heap, a static method doesn't know about any of them.

I'm sure they're talking about MY size variable.

No, I'm pretty sure they're talking about MY size variable.



static methods

Static methods can't use non-static methods, either!

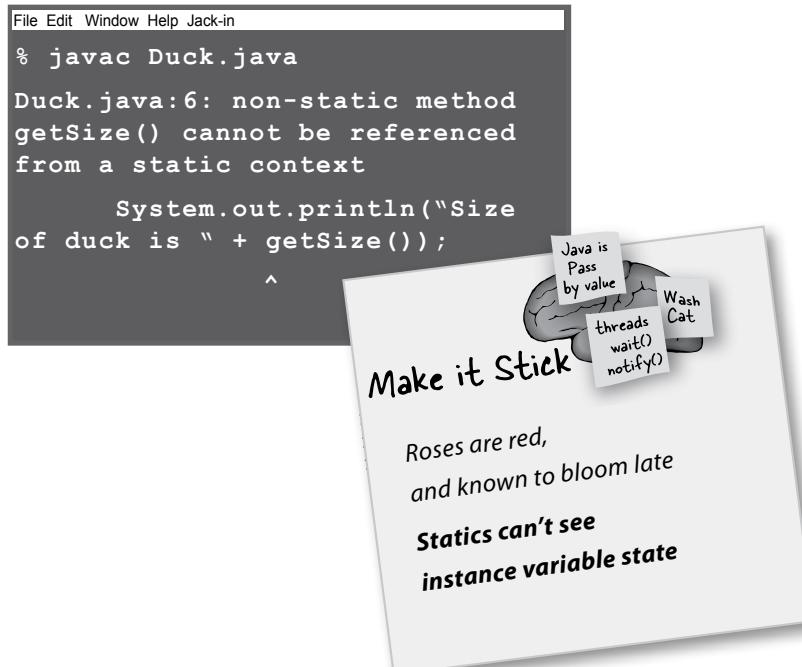
What do non-static methods do? *They usually use instance variable state to affect the behavior of the method.* A getName() method returns the value of the name variable. Whose name? The object used to invoke the getName() method.

This won't compile:

```
public class Duck {  
    private int size;  
  
    public static void main (String[] args) {  
        System.out.println("Size is " + getSize());  
    }  
  
    public void setSize(int s) {  
        size = s;  
    }  
    public int getSize() {  
        return size;  
    }  
}
```

Calling getSize() just postpones the inevitable—getSize() uses the size instance variable.

Back to the same problem... whose size?



there are no
Dumb Questions

Q: What if you try to call a non-static method from a static method, but the non-static method doesn't use any instance variables. Will the compiler allow that?

A: No. The compiler knows that whether you do or do not use instance variables in a non-static method, you *can*. And think about the implications... if you were allowed to compile a scenario like that, then what happens if in the future you want to change the implementation of that non-static method so that one day it *does* use an instance variable? Or worse, what happens if a subclass *overrides* the method and uses an instance variable in the overriding version?

Q: I could swear I've seen code that calls a static method using a reference variable instead of the class name.

A: You *can* do that, but as your mother always told you, "Just because it's legal doesn't mean it's good." Although it *works* to call a static method using any instance of the class, it makes for misleading (less-readable) code. You *can* say,

```
Duck d = new Duck();  
String[] s = {};  
d.main(s);
```

This code is legal, but the compiler just resolves it back to the real class anyway ("OK, *d* is of type Duck, and main() is static, so I'll call the static main() in class Duck"). In other words, using *d* to invoke main() doesn't imply that main() will have any special knowledge of the object that *d* is referencing. It's just an alternate way to invoke a static method, but the method is still static!

Static variable: value is the same for ALL instances of the class

Imagine you wanted to count how many Duck instances are being created while your program is running. How would you do it? Maybe an instance variable that you increment in the constructor?

```
class Duck {
    int duckCount = 0;
    public Duck() {
        duckCount++;
    }
}
```

this would always set
duckCount to 1 each time
a Duck was made

No, that wouldn't work because `duckCount` is an instance variable, and starts at 0 for each Duck. You could try calling a method in some other class, but that's kludgy. You need a class that's got only a single copy of the variable, and all instances share that one copy.

That's what a static variable gives you: a value shared by all instances of a class. In other words, one value per *class*, instead of one value per *instance*.

```
public class Duck {
    private int size;
    private static int duckCount = 0;

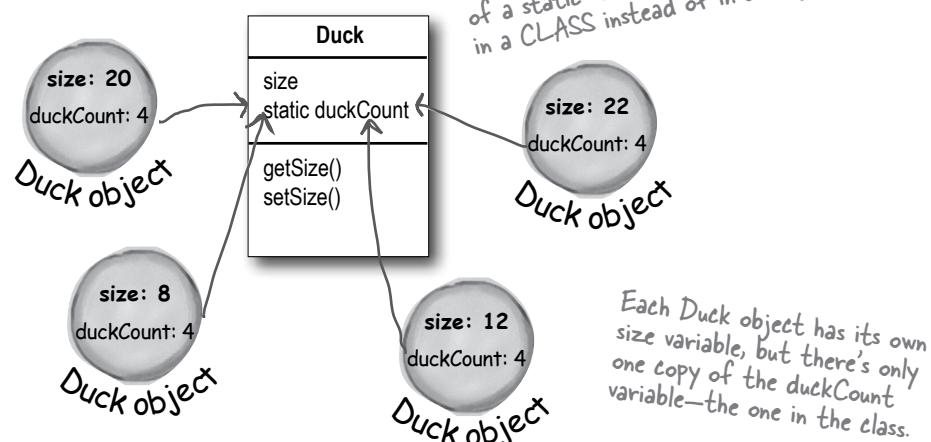
    public Duck() {
        duckCount++; ← Now it will keep
    }

    public void setSize(int s) {
        size = s;
    }

    public int getSize() {
        return size;
    }
}
```

The static `duckCount` variable is initialized ONLY when the class is first loaded, NOT each time a new instance is made.

Now it will keep incrementing each time the Duck constructor runs, because `duckCount` is static and won't be reset to 0.



static variables



Static variables are shared.

All instances of the same class share a single copy of the static variables.

instance variables: 1 per **instance**

static variables: 1 per **class**



Earlier in this chapter, we saw that a private constructor means that the class can't be instantiated from code running outside the class. In other words, only code from within the class can make a new instance of a class with a private constructor. (There's a kind of chicken-and-egg problem here.)

What if you want to write a class in such a way that only ONE instance of it can be created, and anyone who wants to use an instance of the class will always use that one, single instance?

Initializing a static variable

Static variables are initialized when a *class is loaded*. A class is loaded because the JVM decides it's time to load it. Typically, the JVM loads a class because somebody's trying to make a new instance of the class, for the first time, or use a static method or variable of the class. As a programmer, you also have the option of telling the JVM to load a class, but you're not likely to need to do that. In nearly all cases, you're better off letting the JVM decide when to *load* the class.

And there are two guarantees about static initialization:

Static variables in a class are initialized before any *object* of that class can be created.

Static variables in a class are initialized before any *static method* of the class runs.

```
class Player {
    static int playerCount = 0;
    private String name;
    public Player(String n) {
        name = n;
        playerCount++;
    }
}

public class PlayerTestDrive {
    public static void main(String[] args) {
        System.out.println(Player.playerCount);
        Player one = new Player("Tiger Woods");
        System.out.println(Player.playerCount);
    }
}
```



The playerCount is initialized when the class is loaded. We explicitly initialized it to 0, but we don't need to since 0 is the default value for ints. Static variables get default values just like instance variables.

Default values for declared but uninitialized static and instance variables are the same:
 primitive integers (long, short, etc.): 0
 primitive floating points (float, double): 0.0
 boolean: false
 object references: null

Access a static variable just like a static method—with the class name.

Static variables are initialized when the class is loaded. If you don't explicitly initialize a static variable (by assigning it a value at the time you declare it), it gets a default value, so int variables are initialized to zero, which means we didn't need to explicitly say "playerCount = 0". Declaring, but not initializing, a static variable means the static variable will get the default value for that variable type, in exactly the same way that instance variables are given default values when declared.

All static variables in a class are initialized before any object of that class can be created.

```
File Edit Window Help What?
% java PlayerTestDrive
0 ← before any instances are made
1 ← after an object is created
```

static final constants

static final variables are constants

A variable marked **final** means that—once initialized—it can never change. In other words, the value of the static final variable will stay the same as long as the class is loaded. Look up Math.PI in the API, and you'll find:

```
public static final double PI = 3.141592653589793;
```

The variable is marked **public** so that any code can access it.

The variable is marked **static** so that you don't need an instance of class Math (which, remember, you're not allowed to create).

The variable is marked **final** because PI doesn't change (as far as Java is concerned).

There is no other way to designate a variable as a constant, but there is a naming convention that helps you to recognize one.

Constant variable names should be in all caps!

A **static initializer** is a block of code that runs when a class is loaded, before any other code can use the class, so it's a great place to initialize a **static final** variable.

```
class Foo {  
    final static int X;  
    static {  
        X = 42;  
    }  
}
```

Initialize a **final static** variable:

① At the time you declare it:

```
public class Foo {  
    public static final int FOO_X = 25;  
}  
  
OR  
  
notice the naming convention -- static  
final variables are constants, so the  
name should be all uppercase, with an  
underscore separating the words
```

② In a static initializer:

```
public class Bar {  
    public static final double BAR_SIGN;  
  
    static {  
        BAR_SIGN = (double) Math.random();  
    }  
}
```

this code runs as soon as the class is loaded, before any static method variable can be used.

If you don't give a value to a final variable in one of those two places:

```
public class Bar {  
    public static final double BAR_SIGN;  
}  
  
no initialization!
```

The compiler will catch it:

```
File Edit Window Help Jack-in  
% javac Bar.java  
Bar.java:1: variable BAR_SIGN  
might not have been initialized  
1 error
```

final isn't just for static variables...

You can use the keyword `final` to modify non-static variables too, including instance variables, local variables, and even method parameters. In each case, it means the same thing: the value can't be changed. But you can also use final to stop someone from overriding a method or making a subclass.

non-static final variables

```
class Foof {
    final int size = 3; ← now you can't change size
    final int whuffle;

    Foof() {
        whuffle = 42; ← now you can't change whuffle
    }

    void doStuff(final int x) {
        // you can't change x
    }

    void doMore() {
        final int z = 7;
        // you can't change z
    }
}
```

final method

```
class Poof {
    final void calcWhuffle() {
        // important things
        // that must never be overridden
    }
}
```

final class

```
final class MyMostPerfectClass {
    // cannot be extended
}
```

A `final variable` means you can't change its value.

A `final method` means you can't override the method.

A `final class` means you can't extend the class (i.e. you can't make a subclass).



static and final

there are no
Dumb Questions



BULLET POINTS

- A **static method** should be called using the class name rather than an object reference variable:
`Math.random()` vs. `myFoo.go()`
- A static method can be invoked without any instances of the method's class on the heap.
- A static method is good for a utility method that does not (and will never) depend on a particular instance variable value.
- A static method is not associated with a particular instance—only the class—so it cannot access any instance variable values of its class. It wouldn't know *which* instance's values to use.
- A static method cannot access a non-static method, since non-static methods are usually associated with instance variable state.
- If you have a class with only static methods, and you do not want the class to be instantiated, you can mark the constructor private.
- A **static variable** is a variable shared by all members of a given class. There is only one copy of a static variable in a class, rather than one copy per each individual instance for instance variables.
- A static method can access a static variable.
- To make a constant in Java, mark a variable as both static and final.

```
static {  
    DOG_CODE = 420;  
}
```
- A final static variable must be assigned a value either at the time it is declared, or in a static initializer.
- The naming convention for constants (final static variables) is to make the name all uppercase.
- A final variable value cannot be changed once it has been assigned.
- Assigning a value to a final *instance* variable must be either at the time it is declared, or in the constructor.
- A final method cannot be overridden.
- A final class cannot be extended (subclassed).

Q: A static method can't access a non-static variable. But can a non-static method access a static variable?

A: Of course. A non-static method in a class can always call a static method in the class or access a static variable of the class.

Q: Why would I want to make a class final? Doesn't that defeat the whole purpose of OO?

A: Yes and no. A typical reason for making a class final is for security. You can't, for example, make a subclass of the String class. Imagine the havoc if someone extended the String class and substituted their own String subclass objects, polymorphically, where String objects are expected. If you need to count on a particular implementation of the methods in a class, make the class final.

Q: Isn't it redundant to have to mark the methods final if the class is final?

A: If the class is final, you don't need to mark the methods final. Think about it—if a class is final it can never be subclassed, so none of the methods can ever be overridden.

On the other hand, if you *do* want to allow others to extend your class, and you want them to be able to override some, but not all, of the methods, then don't mark the class final but go in and selectively mark specific methods as final. A final method means that a subclass can't override that particular method.



What's Legal?

Given everything you've just learned about static and final, which of these would compile?



- ```

① public class Foo {
 static int x;

 public void go() {
 System.out.println(x);
 }
}

② public class Foo2 {
 int x;

 public static void go() {
 System.out.println(x);
 }
}

③ public class Foo3 {
 final int x;

 public void go() {
 System.out.println(x);
 }
}

④ public class Foo4 {
 static final int x = 12;

 public void go() {
 System.out.println(x);
 }
}

⑤ public class Foo5 {
 static final int x = 12;

 public void go(final int x) {
 System.out.println(x);
 }
}

⑥ public class Foo6 {
 int x = 12;

 public static void go(final int x) {
 System.out.println(x);
 }
}

```

# Math methods

Now that we know how static methods work, let's look at some static methods in class Math. This isn't all of them, just the highlights. Check your API for the rest including sqrt(), tan(), ceil(), floor(), and asin().

## Math.random()

Returns a double between 0.0 through (but not including) 1.0.

```
double r1 = Math.random();
int r2 = (int) (Math.random() * 5);
```

## Math.abs()

Returns a double that is the absolute value of the argument. The method is overloaded, so if you pass it an int it returns an int. Pass it a double it returns a double.

```
int x = Math.abs(-240); // returns 240
double d = Math.abs(240.45); // returns 240.45
```

## Math.round()

Returns an int or a long (depending on whether the argument is a float or a double) rounded to the nearest integer value.

```
int x = Math.round(-24.8f); // returns -25
int y = Math.round(24.45f); // returns 24
```

↑ Remember, floating point literals are assumed to be doubles unless you add the 'f'.

## Math.min()

Returns a value that is the minimum of the two arguments. The method is overloaded to take ints, longs, floats, or doubles.

```
int x = Math.min(24,240); // returns 24
double y = Math.min(90876.5, 90876.49); // returns 90876.49
```

## Math.max()

Returns a value that is the maximum of the two arguments. The method is overloaded to take ints, longs, floats, or doubles.

```
int x = Math.max(24,240); // returns 240
double y = Math.max(90876.5, 90876.49); // returns 90876.5
```

## Wrapping a primitive

Sometimes you want to treat a primitive like an object. For example, in all versions of Java prior to 5.0, you cannot put a primitive directly into a collection like ArrayList or HashMap:

```
int x = 32;
ArrayList list = new ArrayList();
list.add(x);
```

This won't work unless you're using Java 5.0 or greater!! There's no add(int) method in ArrayList that takes an int! (ArrayList only has add() methods that take object references, not primitives.)

There's a wrapper class for every primitive type, and since the wrapper classes are in the java.lang package, you don't need to import them. You can recognize wrapper classes because each one is named after the primitive type it wraps, but with the first letter capitalized to follow the class naming convention.

Oh yeah, for reasons absolutely nobody on the planet is certain of, the API designers decided not to map the names *exactly* from primitive type to class type. You'll see what we mean:

**Boolean**

**Character**

**Byte**

**Short**

**Integer**

**Long**

**Float**

**Double**

Give the primitive to the wrapper constructor. That's it.

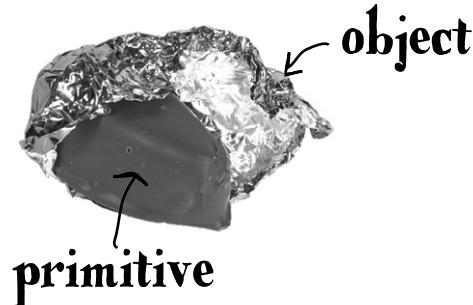
### wrapping a value

```
int i = 288;
Integer iWrap = new Integer(i);
```

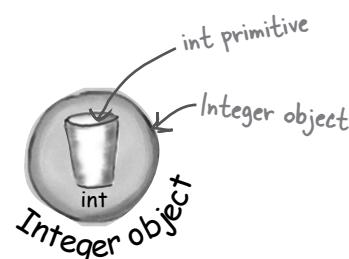
All the wrappers work like this. Boolean has a booleanValue(), Character has a charValue(), etc.

### unwrapping a value

```
int unWrapped = iWrap.intValue();
```



When you need to treat a primitive like an object, wrap it. If you're using any version of Java before 5.0, you'll do this when you need to store a primitive value inside a collection like ArrayList or HashMap.



Note: the picture at the top is a chocolate in a foil wrapper. Get it? Wrapper? Some people think it looks like a baked potato, but that works too.

## static methods



This is stupid. You mean I can't just make an ArrayList of ints??? I have to wrap every single frickin' one in a new Integer object, then unwrap it when I try to access that value in the ArrayList? That's a waste of time and an error waiting to happen...

## Before Java 5.0, YOU had to do the work...

She's right. In all versions of Java prior to 5.0, primitives were primitives and object references were object references, and they were NEVER treated interchangeably. It was always up to you, the programmer, to do the wrapping and unwrapping. There was no way to pass a primitive to a method expecting an object reference, and no way to assign the result of a method returning an object reference directly to a primitive variable—even when the returned reference is to an Integer and the primitive variable is an int. There was simply no relationship between an Integer and an int, other than the fact that Integer has an instance variable of type int (to hold the primitive the Integer wraps). All the work was up to you.

### An ArrayList of primitive ints

#### Without autoboxing (Java versions before 5.0)

```
public void doNumsOldWay() {
 ArrayList listOfNumbers = new ArrayList();

 listOfNumbers.add(new Integer(3)); ← You can't add the primitive '3' to the list,
 so you have to wrap it in an Integer first.

 Integer one = (Integer) listOfNumbers.get(0); ← It comes out as type
 int intOne = one.intValue(); Object, but you can cast
} the Object to an Integer.

Finally you can get the primitive
out of the Integer. ↑
```

Make an ArrayList. (Remember, before 5.0 you could not specify the TYPE, so all ArrayLists were lists of Objects.)

# Autoboxing: blurring the line between primitive and object

The autoboxing feature added to Java 5.0 does the conversion from primitive to wrapper object *automatically!*

Let's see what happens when we want to make an ArrayList to hold ints.

## An ArrayList of primitive ints

### With autoboxing (Java versions 5.0 or greater)

```
public void doNumsNewWay() {
 ArrayList<Integer> listOfNumbers = new ArrayList<Integer>();
```

```
 listOfNumbers.add(3); Just add it!
 int num = listOfNumbers.get(0);
}
```

And the compiler automatically unwraps (unboxes) the Integer object so you can assign the int value directly to a primitive without having to call the intValue() method on the Integer object.

Make an ArrayList of type Integer.  
↓

Although there is NOT a method in ArrayList for add(int), the compiler does all the wrapping (boxing) for you. In other words, there really IS an Integer object stored in the ArrayList, but you get to "pretend" that the ArrayList takes ints. (You can add both ints and Integers to an ArrayList<Integer>.)

**Q:** Why not declare an ArrayList<int> if you want to hold ints?

**A:** Because... *you can't*. Remember, the rule for generic types is that you can specify only class or interface types, *not primitives*. So ArrayList<int> will not compile. But as you can see from the code above, it doesn't really matter, since the compiler lets you put ints into the ArrayList<Integer>. In fact, there's really no way to prevent you from putting primitives into an ArrayList where the type of the list is the type of that primitive's wrapper, if you're using a Java 5.0-compliant compiler, since autoboxing will happen automatically. So, you can put boolean primitives in an ArrayList<Boolean> and chars into an ArrayList<Character>.

**static methods**

## Autoboxing works almost everywhere

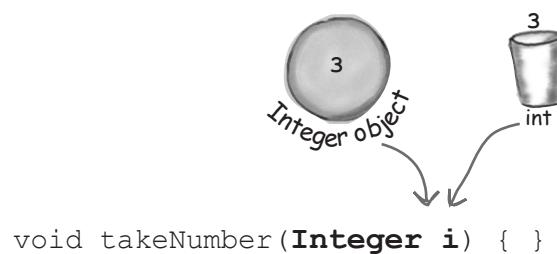
Autoboxing lets you do more than just the obvious wrapping and unwrapping to use primitives in a collection... it also lets you use either a primitive or its wrapper type virtually anywhere one or the other is expected. Think about that!

### Fun with autoboxing

---

#### Method arguments

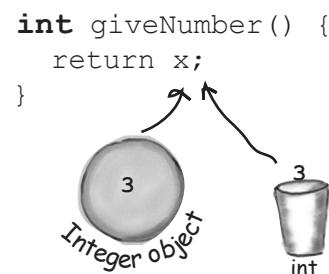
If a method takes a wrapper type, you can pass a reference to a wrapper or a primitive of the matching type. And of course the reverse is true—if a method takes a primitive, you can pass in either a compatible primitive or a reference to a wrapper of that primitive type.



---

#### Return values

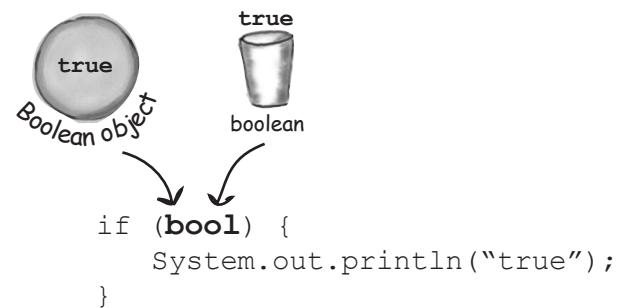
If a method declares a primitive return type, you can return either a compatible primitive or a reference to the wrapper of that primitive type. And if a method declares a wrapper return type, you can return either a reference to the wrapper type or a primitive of the matching type.



---

#### Boolean expressions

Any place a boolean value is expected, you can use either an expression that evaluates to a boolean ( $4 > 2$ ), or a primitive boolean, or a reference to a Boolean wrapper.



## Operations on numbers

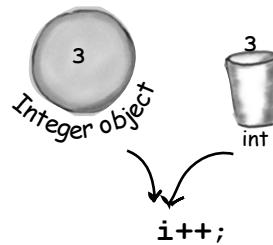
This is probably the strangest one—yes, you can now use a wrapper type as an operand in operations where the primitive type is expected. That means you can apply, say, the increment operator against a reference to an Integer object!

But don't worry—this is just a compiler trick. The language wasn't modified to make the operators work on objects; the compiler simply converts the object to its primitive type before the operation. It sure looks weird, though.

```
Integer i = new Integer(42);
i++;
```

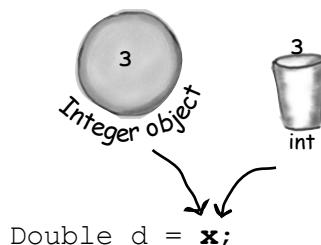
And that means you can also do things like:

```
Integer j = new Integer(5);
Integer k = j + 3;
```



## Assignments

You can assign either a wrapper or primitive to a variable declared as a matching wrapper or primitive. For example, a primitive int variable can be assigned to an Integer reference variable, and vice-versa—a reference to an Integer object can be assigned to a variable declared as an int primitive.



## Sharpen your pencil

```
public class TestBox {
 Integer i;
 int j;

 public static void main (String[] args) {
 TestBox t = new TestBox();
 t.go();
 }

 public void go() {
 j=i;
 System.out.println(j);
 System.out.println(i);
 }
}
```

Will this code compile? Will it run? If it runs, what will it do?

Take your time and think about this one; it brings up an implication of autoboxing that we didn't talk about.

You'll have to go to your compiler to find the answers. (Yes, we're forcing you to experiment, for your own good of course.)

wrapper methods

## But wait! There's more! Wrappers have static utility methods too!

Besides acting like a normal class, the wrappers have a bunch of really useful static methods. We've used one in this book before—Integer.parseInt().

The parse methods take a String and give you back a primitive value.

### Converting a String to a primitive value is easy:

```
String s = "2";
int x = Integer.parseInt(s);
double d = Double.parseDouble("420.24");

boolean b = Boolean.parseBoolean("True");
```

No problem to parse  
"2" into 2.

The (new to 1.5) parseBoolean() method ignores  
the cases of the characters in the String  
argument.

### But if you try to do this:

```
String t = "two";
int y = Integer.parseInt(t);
```

Uh-oh. This compiles just fine, but  
at runtime it blows up. Anything  
that can't be parsed as a number  
will cause a NumberFormatException

### You'll get a runtime exception:

```
File Edit Window Help Clue
% java Wrappers
Exception in thread "main"
java.lang.NumberFormatException: two
at java.lang.Integer.parseInt(Integer.java:409)
at java.lang.Integer.parseInt(Integer.java:458)
at Wrappers.main(Wrappers.java:9)
```

**Every method or constructor that parses a String can throw a NumberFormatException. It's a runtime exception, so you don't have to handle or declare it. But you might want to.**

(We'll talk about Exceptions in the next chapter.)

## And now in reverse... turning a primitive number into a String

There are several ways to turn a number into a String.  
The easiest is to simply concatenate the number to an existing String.

```
double d = 42.5;
String doubleString = "" + d;
```

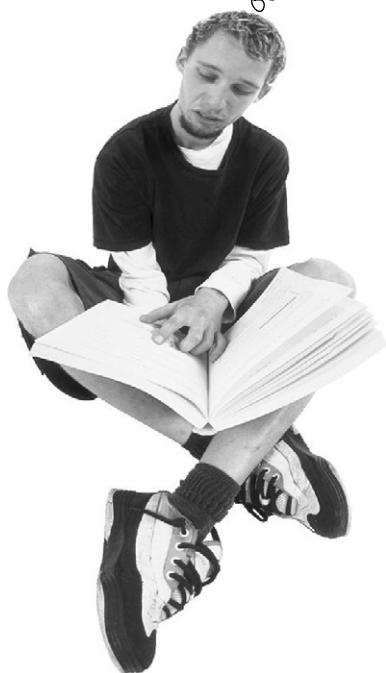
Remember the 't' operator is overloaded  
in Java (the only overloaded operator) as a  
String concatenator. Anything added to a  
String becomes Stringified.

```
double d = 42.5;
String doubleString = Double.toString(d);
```

Another way to do it using a static  
method in class Double.

Yeah,  
but how do I make it  
look like money? With a dollar  
sign and two decimal places  
like \$56.87 or what if I want  
commas like 45,687,890 or  
what if I want it in...

Where's my printf  
like I have in C? Is  
number formatting part of  
the I/O classes?



## number formatting

# Number formatting

In Java, formatting numbers and dates doesn't have to be coupled with I/O. Think about it. One of the most typical ways to display numbers to a user is through a GUI. You put Strings into a scrolling text area, or maybe a table. If formatting was built only into print statements, you'd never be able to format a number into a nice String to display in a GUI. Before Java 5.0, most formatting was handled through classes in the `java.text` package that we won't even look at in this version of the book, now that things have changed.

In Java 5.0, the Java team added more powerful and flexible formatting through a `Formatter` class in `java.util`. But you don't need to create and call methods on the `Formatter` class yourself, because Java 5.0 added convenience methods to some of the I/O classes (including `printf()`) and the `String` class. So it's a simple matter of calling a static `String.format()` method and passing it the thing you want formatted along with formatting instructions.

Of course, you do have to know how to supply the formatting instructions, and that takes a little effort unless you're familiar with the `printf()` function in C/C++. Fortunately, even if you *don't* know `printf()` you can simply follow recipes for the most basic things (that we're showing in this chapter). But you *will* want to learn how to format if you want to mix and match to get *anything* you want.

We'll start here with a basic example, then look at how it works. (Note: we'll revisit formatting again in the I/O chapter.)

## Formatting a number to use commas

```
public class TestFormats {
 public static void main (String[] args) {
 String s = String.format("%, d", 1000000000);
 System.out.println(s);
 }
}
```

1,000,000,000

The number to format (we want it to have commas).

The formatting instructions for how to format the second argument (which in this case is an int value). Remember, there are only two arguments to this method here—the first comma is INSIDE the String literal, so it isn't separating arguments to the `format` method.

Now we get commas inserted into the number.

# Formatting deconstructed...

At the most basic level, formatting consists of two main parts (there is more, but we'll start with this to keep it cleaner):

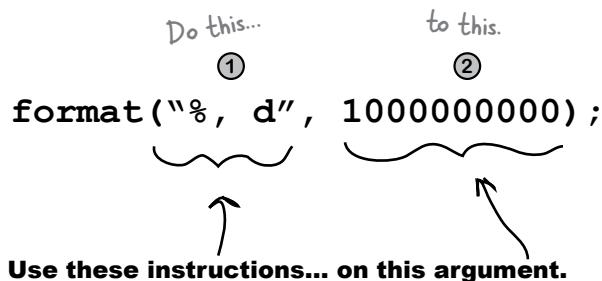
## ① Formatting instructions

You use special format specifiers that describe how the argument should be formatted.

## ② The argument to be formatted.

Although there can be more than one argument, we'll start with just one. The argument type can't be just *anything*... it has to be something that can be formatted using the format specifiers in the formatting instructions. For example, if your formatting instructions specify a *floating point number*, you can't pass in a Dog or even a String that looks like a floating point number.

Note: if you already know printf() from C/C++, you can probably just skim the next few pages. Otherwise, read carefully!



## What do these instructions actually say?

"Take the second argument to this method, and format it as a **decimal integer** and insert **commas**."

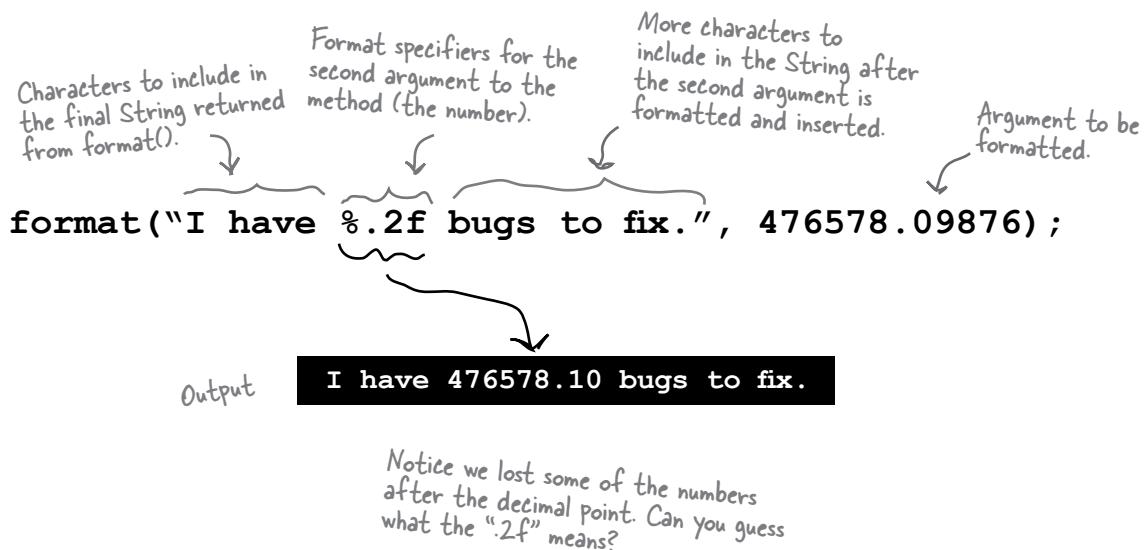
## How do they say that?

On the next page we'll look in more detail at what the syntax "%, d" actually means, but for starters, any time you see the percent sign (%) in a format String (which is always the first argument to a format() method), think of it as representing a variable, and the variable is the other argument to the method. The rest of the characters after the percent sign describe the formatting instructions for the argument.

## the `format()` method

# The percent (%) says, “insert argument here” (and format it using these instructions)

The first argument to a `format()` method is called the format String, and it can actually include characters that you just want printed as-is, without extra formatting. When you see the % sign, though, think of the percent sign as a variable that represents the other argument to the method.



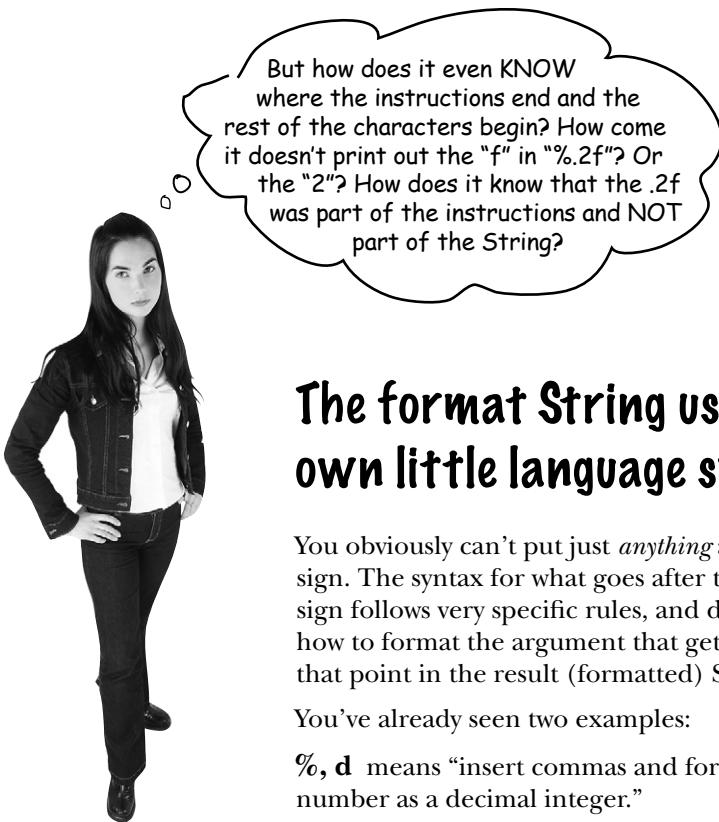
The “%” sign tells the formatter to insert the other method argument (the second argument to `format()`, the number) here, AND format it using the “`.2f`” characters after the percent sign. Then the rest of the format String, “bugs to fix”, is added to the final output.

## Adding a comma

```
format("I have %,.2f bugs to fix.", 476578.09876);
```

I have 476,578.10 bugs to fix.

By changing the format instructions from “`%.2f`” to “`%,.2f`”, we got a comma in the formatted number.



## The format String uses its own little language syntax

You obviously can't put just *anything* after the "%" sign. The syntax for what goes after the percent sign follows very specific rules, and describes how to format the argument that gets inserted at that point in the result (formatted) String.

You've already seen two examples:

**%, d** means "insert commas and format the number as a decimal integer."

and

**%.**2f**** means "format the number as a floating point with a precision of two decimal places."

and

**%.,**2f**** means "insert commas and format the number as a floating point with a precision of two decimal places."

The real question is really, "How do I know what to put after the percent sign to get it to do what I want?" And that includes knowing the symbols (like "d" for decimal and "f" for floating point) as well as the order in which the instructions must be placed following the percent sign. For example, if you put the comma after the "d" like this: "%d," instead of "%,d" it won't work!

Or will it? What do you think this will do:

```
String.format("I have %.2f, bugs to fix.", 476578.09876);
```

(We'll answer that on the next page.)

## format specifier

# The format specifier

Everything after the percent sign up to and including the type indicator (like "d" or "f") are part of the formatting instructions. After the type indicator, the formatter assumes the next set of characters are meant to be part of the output String, until or unless it hits another percent (%) sign. Hmmmm... is that even possible? Can you have more than one formatted argument variable? Put that thought on hold for right now; we'll come back to it in a few minutes. For now, let's look at the syntax for the format specifiers—the things that go after the percent (%) sign and describe how the argument should be formatted.

**A format specifier can have up to five different parts (not including the "%"). Everything in brackets [ ] below is optional, so only the percent (%) and the type are required. But the order is also mandatory, so any parts you DO use must go in this order.**

% [argument number] [flags] [width] [.precision] **type**

We'll get to this later...  
it lets you say WHICH  
argument if there's more  
than one. (Don't worry  
about it just yet.)

These are for  
special formatting  
options like inserting  
commas, or putting  
negative numbers in  
parentheses, or to  
make the numbers  
left justified.

This defines the  
MINIMUM number  
of characters that  
will be used. That's  
**\*minimum\*** not  
**TOTAL**. If the number  
is longer than the  
width, it'll still be used  
in full, but if it's less  
than the width, it'll be  
padded with zeroes.

You already know  
this one...it defines  
the precision. In  
other words, it  
sets the number  
of decimal places.  
Don't forget to  
include the **"."** in  
there.

Type is mandatory  
(see the next page)  
and will usually be  
"d" for a decimal  
integer or "f" for  
a floating point  
number.

% [argument number] [flags] [width] [.precision] **type**

format ("% , 6.1f", 42.000);

There's no "argument number"  
specified in this format String,  
but all the other pieces are there.

## The only required specifier is for TYPE

Although type is the only required specifier, remember that if you *do* put in anything else, type must always come last! There are more than a dozen different type modifiers (not including dates and times; they have their own set), but most of the time you'll probably use %d (decimal) or %f (floating point). And typically you'll combine %f with a precision indicator to set the number of decimal places you want in your output.

**The TYPE is mandatory, everything else is optional.**

**%d** **decimal**  
`format("%d", 42);`  
**42**

A 42.25 would not work! It would be the same as trying to directly assign a double to an int variable.

The argument must be compatible with an int, so that means only byte, short, int, and char (or their wrapper types).

**%f** **floating point**  
`format("%.3f", 42.000000);`  
**42.000**

Here we combined the "f" with a precision indicator ".3" so we ended up with three zeroes.

The argument must be of a floating point type, so that means only a float or double (primitive or wrapper) as well as something called BigDecimal (which we don't look at in this book).

**%x** **hexadecimal**  
`format("%x", 42);`  
**2a**

The argument must be a byte, short, int, long (including both primitive and wrapper types), and BigInteger.

**%c** **character**  
`format("%c", 42);`  
**\***

The number 42 represents the char "\*".

The argument must be a byte, short, char, or int (including both primitive and wrapper types).

You must include a type in your format instructions, and if you specify things besides type, the type must always come last. Most of the time, you'll probably format numbers using either "d" for decimal or "f" for floating point.

## format arguments

# What happens if I have more than one argument?

Imagine you want a String that looks like this:

“The rank is **20,456,654** out of **100,567,890.24**.”

But the numbers are coming from variables. What do you do? You simply add *two* arguments after the format String (first argument), so that means your call to format() will have three arguments instead of two. And inside that first argument (the format String), you’ll have two different format specifiers (two things that start with “%”). The first format specifier will insert the second argument to the method, and the second format specifier will insert the third argument to the method. In other words, the variable insertions in the format String use the order in which the other arguments are passed into the format() method.

```
int one = 20456654;
double two = 100567890.248907;
String s = String.format("The rank is %,d out of %,.2f", one, two);
```

The rank is 20,456,654 out of 100,567,890.25

We added commas to both variables, and restricted the floating point number (the second variable) to two decimal places.

When you have more than one argument, they’re inserted using the order in which you pass them to the format() method.

As you’ll see when we get to date formatting, you might actually want to apply different formatting specifiers to the same argument. That’s probably hard to imagine until you see how *date* formatting (as opposed to the *number* formating we’ve been doing) works. Just know that in a minute, you’ll see how to be more specific about which format specifiers are applied to which arguments.

**Q:** Um, there’s something REALLY strange going on here. Just how many arguments *can* I pass? I mean, how many overloaded format() methods are IN the String class? So, what happens if I want to pass, say, ten different arguments to be formatted for a single output String?

**A:** Good catch. Yes, there *is* something strange (or at least new and different) going on, and no there are *not* a bunch of overloaded format() methods to take a different number of possible arguments. In order to support this new formatting (printf-like) API in Java, the language needed another new feature—*variable argument lists* (called *varargs* for short). We’ll talk about varargs only in the appendix because outside of formatting, you probably won’t use them much in a well-designed system.

## So much for numbers, what about dates?

Imagine you want a String that looks like this: “Sunday, Nov 28 2004”

Nothing special there, you say? Well, imagine that all you have to start with is a variable of type Date—A Java class that can represent a timestamp, and now you want to take that object (as opposed to a number) and send it through the formatter.

The main difference between number and date formatting is that date formats use a two-character type that starts with “t” (as opposed to the single character “f” or “d”, for example). The examples below should give you a good idea of how it works:

### The complete date and time      %tc

```
String.format("%tc", new Date());
```

**Sun Nov 28 14:52:41 MST 2004**

### Just the time      %tr

```
String.format("%tr", new Date());
```

**03:01:47 PM**

### Day of the week, month and day      %tA %tB %td

There isn’t a single format specifier that will do exactly what we want, so we have to combine three of them for day of the week (%tA), month (%tB), and day of the month (%td).

```
Date today = new Date();
String.format("%tA, %tB %td", today, today, today)
```

The comma is not part of the formatting... it's just the character we want printed after the first inserted formatted argument.

But that means we have to pass the Date object in three times, one for each part of the format that we want. In other words, the %tA will give us just the day of the week, but then we have to do it again to get just the month and again for the day of the month.

**Sunday, November 28**

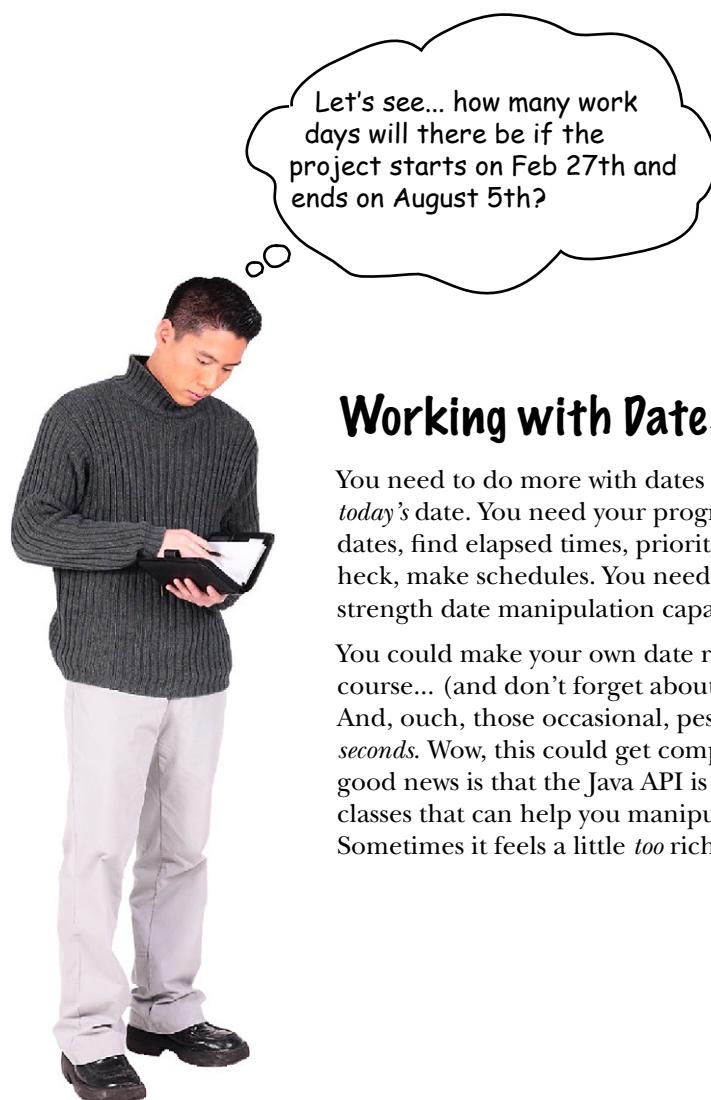
### Same as above, but without duplicating the arguments      %tA %tB %td

```
Date today = new Date();
String.format("%tA, %<tb %<td", today);
```

You can think of this as kind of like calling three different getter methods on the Date object, to get three different pieces of data from it.

The angle-bracket “<” is just another flag in the specifier that tells the formatter to “use the previous argument again.” So it saves you from repeating the arguments, and instead you format the same argument three different ways.

## manipulating dates



## Working with Dates

You need to do more with dates than just get *today*'s date. You need your programs to adjust dates, find elapsed times, prioritize schedules, heck, make schedules. You need industrial strength date manipulation capabilities.

You could make your own date routines of course... (and don't forget about leap years!) And, ouch, those occasional, pesky leap-seconds. Wow, this could get complicated. The good news is that the Java API is rich with classes that can help you manipulate dates. Sometimes it feels a little *too* rich...

## Moving backward and forward in time

Let's say your company's work schedule is Monday through Friday. You've been assigned the task of figuring out the last work day in each calendar month this year...

### **It seems that `java.util.Date` is actually... out of date**

Earlier we used `java.util.Date` to find today's date, so it seems logical that this class would be a good place to start looking for some handy date manipulation capabilities, but when you check out the API you'll find that most of `Date`'s methods have been deprecated!

The `Date` class is still great for getting a "time stamp"—an object that represents the current date and time, so use it when you want to say, "give me NOW".

The good news is that the API recommends `java.util.Calendar` instead, so let's take a look:

### **Use `java.util.Calendar` for your date manipulation**

The designers of the `Calendar` API wanted to think globally, literally. The basic idea is that when you want to work with dates, you ask for a `Calendar` (through a static method of the `Calendar` class that you'll see on the next page), and the JVM hands you back an instance of a concrete subclass of `Calendar`. (`Calendar` is actually an abstract class, so you're always working with a concrete subclass.)

More interesting, though, is that the *kind* of calendar you get back will be *appropriate for your locale*. Much of the world uses the Gregorian calendar, but if you're in an area that doesn't use a Gregorian calendar you can get Java libraries to handle other calendars such as Buddhist, or Islamic or Japanese.

The standard Java API ships with `java.util.GregorianCalendar`, so that's what we'll be using here. For the most part, though, you don't even have to think about the kind of `Calendar` subclass you're using, and instead focus only on the methods of the `Calendar` class.

**For a time-stamp of "now", use `Date`. But for everything else, use `Calendar`.**

## getting a Calendar

# Getting an object that extends Calendar

How in the world do you get an “instance” of an abstract class? Well you don’t of course, this won’t work:

**This WON’T work:**

```
Calendar cal = new Calendar();
```

The compiler won’t allow this!

**Instead, use the static “getInstance()” method:**

```
Calendar cal = Calendar.getInstance();
```

This syntax should look familiar at this point – we’re invoking a static method.

Wait a minute.  
If you can’t make an instance of the Calendar class, what exactly are you assigning to that Calendar reference?



**You can’t get an instance of Calendar, but you can get an instance of a concrete Calendar subclass.**

Obviously you can’t get an instance of Calendar, because Calendar is abstract. But you’re still free to call static methods on Calendar, since *static* methods are called on the *class*, rather than on a particular instance. So you call the static getInstance() on Calendar and it gives you back... an instance of a concrete subclass. Something that extends Calendar (which means it can be polymorphically assigned to Calendar) and which—by contract—can respond to the methods of class Calendar.

In most of the world, and by default for most versions of Java, you’ll be getting back a **java.util.GregorianCalendar** instance.

## Working with Calendar objects

There are several key concepts you'll need to understand in order to work with Calendar objects:

- **Fields hold state** - A Calendar object has many fields that are used to represent aspects of its ultimate state, its date and time. For instance, you can get and set a Calendar's *year* or *month*.
- **Dates and Times can be incremented** - The Calendar class has methods that allow you to add and subtract values from various fields, for example "add one to the month", or "subtract three years".
- **Dates and Times can be represented in milliseconds** - The Calendar class lets you convert your dates into and out of a millisecond representation. (Specifically, the number of milliseconds that have occurred since January 1st, 1970.) This allows you to perform precise calculations such as "elapsed time between two times" or "add 63 hours and 23 minutes and 12 seconds to this time".

### An example of working with a Calendar object:

```
Calendar c = Calendar.getInstance(); Set time to Jan. 7, 2004 at 15:40.

c.set(2004, 0, 7, 15, 40); (Notice the month is zero-based.)

long day1 = c.getTimeInMillis(); Convert this to a big ol'

 amount of milliseconds.

day1 += 1000 * 60 * 60; Add an hour's worth of millis, then update the time.

c.setTimeInMillis(day1); (Notice the "+=", it's like day1 = day1 + ...).

System.out.println("new hour " + c.get(c.HOUR_OF_DAY));

c.add(c.DATE, 35); Add 35 days to the date, which
 should move us into February.

System.out.println("add 35 days " + c.getTime());

c.roll(c.DATE, 35); "Roll" 35 days onto this date. This
 "rolls" the date ahead 35 days, but
 DOES NOT change the month!

c.set(c.DATE, 1); We're not incrementing here, just
 doing a "set" of the date.

System.out.println("set to 1 " + c.getTime());
```

```
File Edit Window Help Time-Flies

new hour 16

add 35 days Wed Feb 11 16:40:41 MST 2004

roll 35 days Tue Feb 17 16:40:41 MST 2004

set to 1 Sun Feb 01 16:40:41 MST 2004
```

This output confirms how millis, add, roll, and set work.

## Calendar API

# Highlights of the Calendar API

We just worked through using a few of the fields and methods in the `Calendar` class. This is a big API, so we're showing only a few of the most common fields and methods that you'll use. Once you get a few of these it should be pretty easy to bend the rest of this API to your will.

### Key Calendar Methods

#### `add(int field, int amount)`

Adds or subtracts time from the calendar's field.

#### `get(int field)`

Returns the value of the given calendar field.

#### `getInstance()`

Returns a `Calendar`, you can specify a locale.

#### `getTimeInMillis()`

Returns this `Calendar`'s time in millis, as a long.

#### `roll(int field, boolean up)`

Adds or subtracts time without changing larger fields.

#### `set(int field, int value)`

Sets the value of a given `Calendar` field.

#### `set(year, month, day, hour, minute) (all ints)`

A common variety of `set` to set a complete time.

#### `setTimeInMillis(long millis)`

Sets a `Calendar`'s time based on a long milli-time.

// more...

### Key Calendar Fields

#### `DATE / DAY_OF_MONTH`

Get / set the day of month

#### `HOUR / HOUR_OF_DAY`

Get / set the 12 hour or 24 hour value.

#### `MILLISECOND`

Get / set the milliseconds.

#### `MINUTE`

Get / set the minute.

#### `MONTH`

Get / set the month.

#### `YEAR`

Get / set the year.

#### `ZONE_OFFSET`

Get / set raw offset of GMT in millis.

// more...

## Even more Statics!... static imports

New to Java 5.0... a real mixed blessing. Some people love this idea, some people hate it. Static imports exist only to save you some typing. If you hate to type, you might just like this feature. The downside to static imports is that - if you're not careful - using them can make your code a lot harder to read.

The basic idea is that whenever you're using a static class, a static variable, or an enum (more on those later), you can import them, and save yourself some typing.

### Some old-fashioned code:

```
import java.lang.Math;

class NoStatic {

 public static void main(String [] args) {

 System.out.println("sqrt " + Math.sqrt(2.0));
 System.out.println("tan " + Math.tan(60));
 }
}
```

The syntax to use when  
declaring static imports.

### Same code, with static imports:

```
import static java.lang.System.out;
import static java.lang.Math.*;

class WithStatic {

 public static void main(String [] args) {

 out.println("sqrt " + sqrt(2.0));
 out.println("tan " + tan(60));
 }
}
```

Static imports in action.

### Use Carefully:

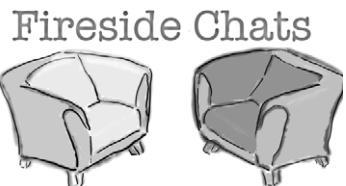
**static imports can  
make your code  
confusing to read**



### - Caveats & Gotchas

- If you're only going to use a static member a few times, we think you should avoid static imports, to help keep the code more readable.
- If you're going to use a static member a lot, (like doing lots of Math calculations), then it's probably OK to use the static import.
- Notice that you can use wildcards (\*), in your static import declaration.
- A big issue with static imports is that it's not too hard to create naming conflicts. For example, if you have two different classes with an "add()" method, how will you and the compiler know which one to use?

## static vs. instance



Tonight's Talk: **An instance variable takes cheap shots at a static variable**

### Instance Variable

I don't even know why we're doing this. Everyone knows static variables are just used for constants. And how many of those are there? I think the whole API must have, what, four? And it's not like anybody ever uses them.

Full of it. Yeah, you can say that again. OK, so there are a few in the Swing library, but everybody knows Swing is just a special case.

Ok, but besides a few GUI things, give me an example of just one static variable that anyone would actually use. In the real world.

Well, that's another special case. And nobody uses that except for debugging anyway.

### Static Variable

You really should check your facts. When was the last time you looked at the API? It's frickin' loaded with statics! It even has entire classes dedicated to holding constant values. There's a class called `SwingConstants`, for example, that's just full of them.

It might be a special case, but it's a really important one! And what about the `Color` class? What a pain if you had to remember the RGB values to make the standard colors? But the color class already has constants defined for blue, purple, white, red, etc. Very handy.

How's `System.out` for starters? The `out` in `System.out` is a static variable of the `System` class. You personally don't make a new instance of the `System`, you just ask the `System` class for its `out` variable.

Oh, like debugging isn't important?

And here's something that probably never crossed your narrow mind—let's face it, static variables are more efficient. One per class instead of one per instance. The memory savings might be huge!

**Instance Variable**

Um, aren't you forgetting something?

Static variables are about as un-OO as it gets!!  
Gee why not just go take a giant backwards  
step and do some procedural programming  
while we're at it.

You're like a global variable, and any  
programmer worth his PDA knows that's  
usually a Bad Thing.

Yeah you live in a class, but they don't call  
it *Class-Oriented* programming. That's just  
stupid. You're a relic. Something to help the  
old-timers make the leap to java.

Well, OK, every once in a while sure, it makes  
sense to use a static, but let me tell you, abuse  
of static variables (and methods) is the mark  
of an immature OO programmer. A designer  
should be thinking about *object* state, not *class*  
state.

Static methods are the worst things of all,  
because it usually means the programmer is  
thinking procedurally instead of about objects  
doing things based on their unique object  
state.

Riiiiiight. Whatever you need to tell yourself...

**Static Variable**

What?

What do you mean *un-OO*?

I am NOT a global variable. There's no such  
thing. I live in a class! That's pretty OO you  
know, a CLASS. I'm not just sitting out there  
in space somewhere; I'm a natural part of the  
state of an object; the only difference is that  
I'm shared by all instances of a class. Very  
efficient.

Alright just stop right there. THAT is  
definitely not true. Some static variables are  
absolutely crucial to a system. And even the  
ones that aren't crucial sure are handy.

Why do you say that? And what's wrong with  
static methods?

Sure, I know that objects should be the focus  
of an OO design, but just because there are  
some clueless programmers out there... don't  
throw the baby out with the bytecode. There's  
a time and place for statics, and when you  
need one, nothing else beats it.

## be the compiler



```
class StaticSuper{

 static {
 System.out.println("super static block");
 }

 StaticSuper{
 System.out.println(
 "super constructor");
 }

 public class StaticTests extends StaticSuper {
 static int rand;

 static {
 rand = (int) (Math.random() * 6);
 System.out.println("static block " + rand);
 }

 StaticTests() {
 System.out.println("constructor");
 }

 public static void main(String [] args) {
 System.out.println("in main");
 StaticTests st = new StaticTests();
 }
 }
```

## BE the compiler

The Java file on this page represents a complete program. Your job is to play compiler and determine whether this file will compile. If it won't compile, how would you fix it, and if it does compile, what would be its output?



If it compiles, which of these is the output?

### Possible Output

```
File Edit Window Help Cling
%java StaticTests
static block 4
in main
super static block
super constructor
constructor
```

```
StaticTests() {
 System.out.println("constructor");
}

public static void main(String [] args) {
 System.out.println("in main");
 StaticTests st = new StaticTests();
}
```

### Possible Output

```
File Edit Window Help Electricity
%java StaticTests
super static block
static block 3
in main
super constructor
constructor
```



This chapter explored the wonderful, static, world of Java. Your job is to decide whether each of the following statements is true or false.

## 👍 TRUE OR FALSE 👎

1. To use the Math class, the first step is to make an instance of it.
2. You can mark a constructor with the **static** keyword.
3. Static methods don't have access to instance variable state of the 'this' object.
4. It is good practice to call a static method using a reference variable.
5. Static variables could be used to count the instances of a class.
6. Constructors are called before static variables are initialized.
7. MAX\_SIZE would be a good name for a static final variable.
8. A static initializer block runs before a class's constructor runs.
9. If a class is marked final, all of its methods must be marked final.
10. A final method can only be overridden if its class is extended.
11. There is no wrapper class for boolean primitives.
12. A wrapper is used when you want to treat a primitive like an object.
13. The parseXxx methods always return a String.
14. Formatting classes (which are decoupled from I/O), are in the java.format package.

code magnets



## Lunar Code Magnets

This one might actually be useful! In addition to what you've learned in the last few pages about manipulating dates, you'll need a little more information... First, full moons happen every 29.52 days or so. Second, there was a full moon on Jan. 7th, 2004. Your job is to reconstruct the code snippets to make a working Java program that produces the output listed below (plus more full moon dates). (You might not need all of the magnets, and add all the curly braces you need.) Oh, by the way, your output will be different if you don't live in the mountain time zone.

```
long day1 = c.getTimeInMillis();
c.set(2004,1,7,15,40);
```

```
import static java.lang.System.out;
static int DAY_IM = 60 * 60 * 24;
```

```
("full moon on %tc", c));
```

```
Calendar c = new Calendar();
```

```
(c.format
```

```
class FullMoons {
```

```
public static void main(String [] args) {
```

```
day1 += (DAY_IM * 29.52);
```

```
for (int x = 0; x < 60; x++) {
```

```
static int DAY_IM = 1000 * 60 * 60 * 24;
```

```
println
```

```
("full moon on %t", c));
```

```
import java.io.*;
```

```
import java.util.*;
```

```
static import java.lang.System.out;
```

```
c.set(2004,0,7,15,40);
```

```
out.println
```

```
c.setTimeInMillis(day1);
```

```
(String.format
```

```
Calendar c = Calendar.getInstance();
```

File Edit Window Help How

```
% java FullMoons
full moon on Fri Feb 06 04:09:35 MST 2004
full moon on Sat Mar 06 16:38:23 MST 2004
full moon on Mon Apr 05 06:07:11 MDT 2004
```

## Exercise Solutions

# BE the compiler

```
StaticSuper() {
 System.out.println(
 "super constructor");
}
```

StaticSuper is a constructor, and must have () in its signature. Notice that as the output below demonstrates, the static blocks for both classes run before either of the constructors run.

### Possible Output

```
File Edit Window Help Cling
%java StaticTests
super static block
static block 3
in main
super constructor
constructor
```

### True or False

1. To use the Math class, the first step is to make an instance of it. **False**
2. You can mark a constructor with the keyword ‘static’. **False**
3. Static methods don’t have access to an object’s instance variables. **True**
4. It is good practice to call a static method using a reference variable. **False**
5. Static variables could be used to count the instances of a class. **True**
6. Constructors are called before static variables are initialized. **False**
7. MAX\_SIZE would be a good name for a static final variable. **True**
8. A static initializer block runs before a class’s constructor runs. **True**
9. If a class is marked final, all of its methods must be marked final. **False**
10. A final method can only be overridden if its class is extended. **False**
11. There is no wrapper class for boolean primitives. **False**
12. A wrapper is used when you want to treat a primitive like an object. **True**
13. The parseXxx methods always return a String. **False**
14. Formatting classes (which are decoupled from I/O), are in the java.format package. **False**

## code magnets solution



## Exercise Solutions

```
import java.util.*;
import static java.lang.System.out;
class FullMoons {
 static int DAY_IM = 1000 * 60 * 60 * 24;
 public static void main(String [] args) {
 Calendar c = Calendar.getInstance();
 c.set(2004,0,7,15,40);
 long day1 = c.getTimeInMillis();
 for (int x = 0; x < 60; x++) {
 day1 += (DAY_IM * 29.52)
 c.setTimeInMillis(day1);
 out.println(String.format("full moon on %tc", c));
 }
 }
}
```

### Notes on the Lunar Code Magnet:

You might discover that a few of the dates produced by this program are off by a day. This astronomical stuff is a little tricky, and if we made it perfect, it would be too complex to make an exercise here.

Hint: one problem you might try to solve is based on differences in time zones. Can you spot the issue?

```
File Edit Window Help Howl
% java FullMoons
full moon on Fri Feb 06 04:09:35 MST 2004
full moon on Sat Mar 06 16:38:23 MST 2004
full moon on Mon Apr 05 06:07:11 MDT 2004
```

## 11 exception handling

# Risky Behavior



**Stuff happens. The file isn't there. The server is down.** No matter how good a programmer you are, you can't control everything. Things can go wrong. *Very wrong.* When you write a risky method, you need code to handle the bad things that might happen. But how do you *know* when a method is risky? And where do you put the code to *handle* the **exceptional** situation? So far in this book, we haven't *really* taken any risks. We've certainly had things go wrong at runtime, but the problems were mostly flaws in our own code. Bugs. And those we should fix at development time. No, the problem-handling code we're talking about here is for code that you *can't* guarantee will work at runtime. Code that expects the file to be in the right directory, the server to be running, or the Thread to stay asleep. And we have to do this *now*. Because in *this* chapter, we're going to build something that uses the risky JavaSound API. We're going to build a MIDI Music Player.

## building the MIDI Music Player

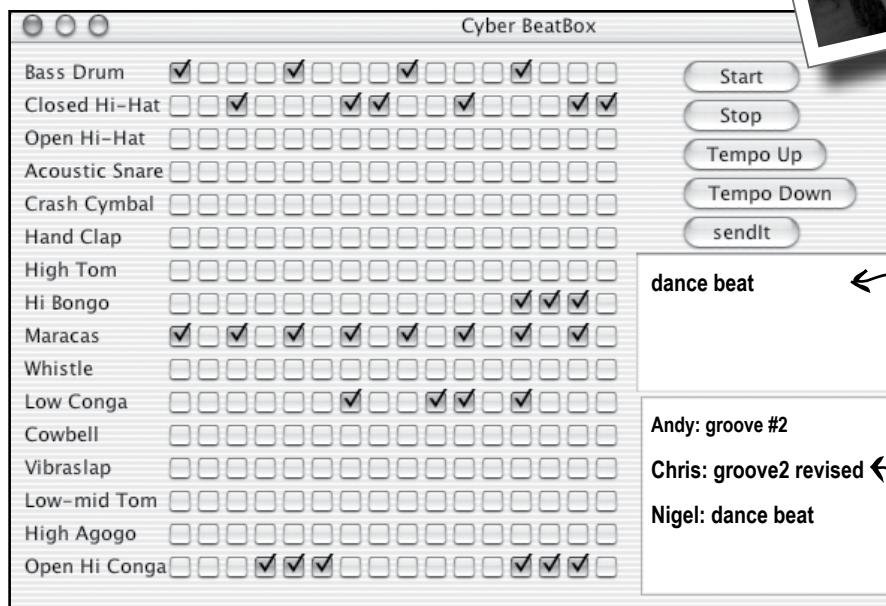
# Let's make a Music Machine

Over the next three chapters, we'll build a few different sound applications, including a BeatBox Drum Machine. In fact, before the book is done, we'll have a multi-player version so you can send your drum loops to another player, kind of like a chat room. You're going to write the whole thing, although you can choose to use Ready-bake code for the GUI parts.

OK, so not every IT department is looking for a new BeatBox server, but we're doing this to *learn* more about Java. Building a BeatBox is just a way to have fun *while* we're learning Java.

**The finished BeatBox looks something like this:**

You make a beatbox loop (a 16-beat drum pattern) by putting checkmarks in the boxes.



your message, that gets sent to the other players, along with your current beat pattern, when you hit "SendIt"

incoming messages from other players. Click one to load the pattern that goes with it, and then click 'Start' to play it.

Put checkmarks in the boxes for each of the 16 'beats'. For example, on beat 1 (of 16) the Bass drum and the Maracas will play, on beat 2 nothing, and on beat 3 the Maracas and Closed Hi-Hat... you get the idea. When you hit 'Start', it plays your pattern in a loop until you hit 'Stop'. At any time, you can "capture" one of your own patterns by sending it to the BeatBox server (which means any other players can listen to it). You can also load any of the incoming patterns by clicking on the message that goes with it.

## We'll start with the basics

Obviously we've got a few things to learn before the whole program is finished, including how to build a Swing GUI, how to *connect* to another machine via networking, and a little I/O so we can *send* something to the other machine.

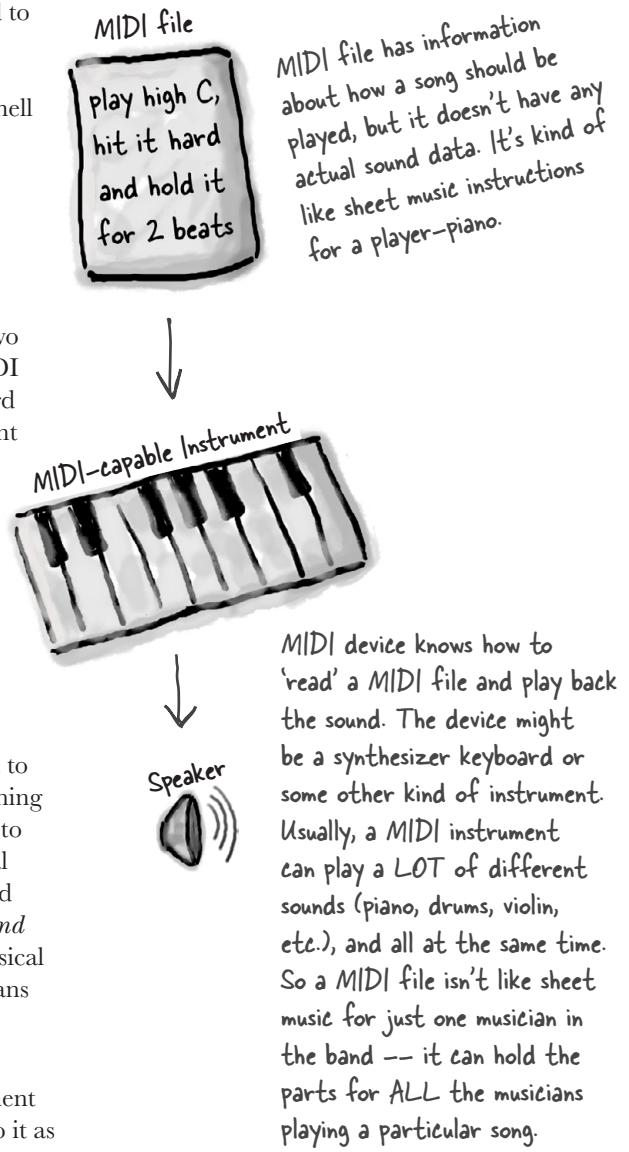
Oh yeah, and the JavaSound API. *That's* where we'll start in this chapter. For now, you can forget the GUI, forget the networking and the I/O, and focus only on getting some MIDI-generated sound to come out of your computer. And don't worry if you don't know a thing about MIDI, or a thing about reading or making music. Everything you need to learn is covered here. You can almost smell the record deal.

### The JavaSound API

JavaSound is a collection of classes and interfaces added to Java starting with version 1.3. These aren't special add-ons; they're part of the standard J2SE class library. JavaSound is split into two parts: MIDI and Sampled. We use only MIDI in this book. MIDI stands for Musical Instrument Digital Interface, and is a standard protocol for getting different kinds of electronic sound equipment to communicate. But for our BeatBox app, you can think of MIDI as *a kind of sheet music* that you feed into some device you can think of like a high-tech 'player piano'. In other words, MIDI data doesn't actually include any *sound*, but it does include the *instructions* that a MIDI-reading instrument can play back. Or for another analogy, you can think of a MIDI file like an HTML document, and the instrument that renders the MIDI file (i.e. *plays* it) is like the Web browser.

MIDI data says *what* to do (play middle C, and here's how hard to hit it, and here's how long to hold it, etc.) but it doesn't say anything at all about the actual *sound* you hear. MIDI doesn't know how to make a flute, piano, or Jimi Hendrix guitar sound. For the actual sound, we need an instrument (a MIDI device) that can read and play a MIDI file. But the device is usually more like an *entire band or orchestra* of instruments. And that instrument might be a physical device, like the electronic keyboard synthesizers the rock musicians play, or it could even be an instrument built entirely in software, living in your computer.

For our BeatBox, we use only the built-in, software-only instrument that you get with Java. It's called a *synthesizer* (some folks refer to it as a *software synth*) because it *creates* sound. Sound that you *hear*.



but it looked so simple

## First we need a Sequencer

Before we can get any sound to play, we need a Sequencer object. The sequencer is the object that takes all the MIDI data and sends it to the right instruments. It's the thing that *plays* the music. A sequencer can do a lot of different things, but in this book, we're using it strictly as a playback device. Like a CD-player on your stereo, but with a few added features. The Sequencer class is in the javax.sound.midi package (part of the standard Java library as of version 1.3). So let's start by making sure we can make (or get) a Sequencer object.

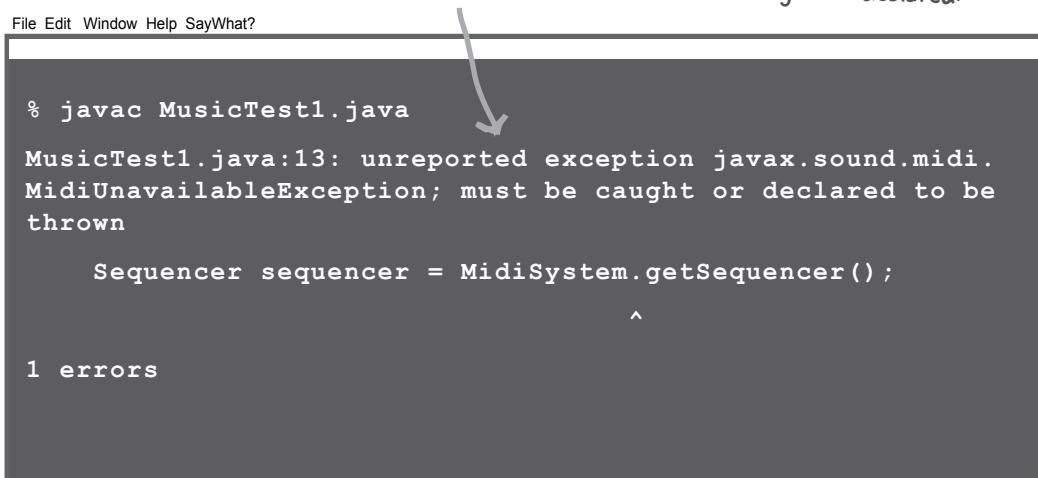
```
import javax.sound.midi.*; ← import the javax.sound.midi package
public class MusicTest1 {
 public void play() {
 Sequencer sequencer = MidiSystem.getSequencer();
 System.out.println("We got a sequencer");
 } // close play

 public static void main(String[] args) {
 MusicTest1 mt = new MusicTest1();
 mt.play();
 } // close main
} // close class
```

We need a Sequencer object. It's the main part of the MIDI device/instrument we're using. It's the thing that, well, sequences all the MIDI information into a 'song'. But we don't make a brand new one ourselves -- we have to ask the MidiSystem to give us one.

### Something's wrong!

This code won't compile! The compiler says there's an 'unreported exception' that must be caught or declared.



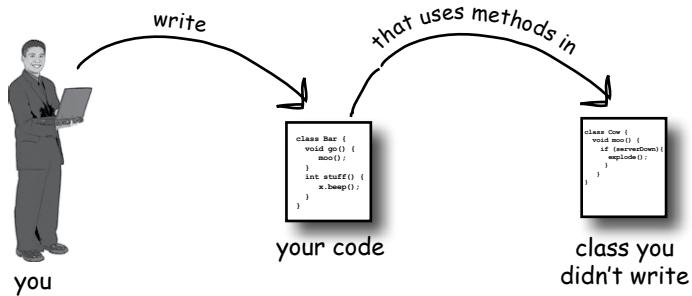
The screenshot shows a terminal window with the following text:

```
File Edit Window Help SayWhat?
% javac MusicTest1.java
MusicTest1.java:13: unreported exception javax.sound.midi.MidiUnavailableException; must be caught or declared to be thrown
 Sequencer sequencer = MidiSystem.getSequencer();
 ^
1 errors
```

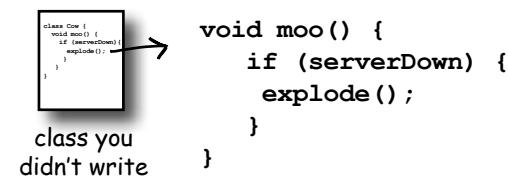
A red arrow points from the explanatory text above to the word "thrown" in the error message. Another red arrow points from the explanatory text to the word "must be caught or declared to be thrown".

## What happens when a method you want to call (probably in a class you didn't write) is risky?

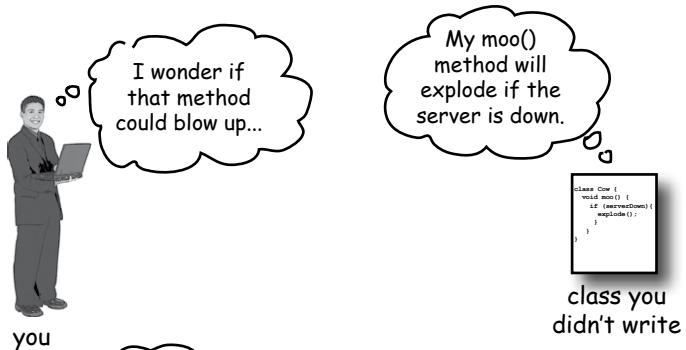
- ① Let's say you want to call a method in a class that you didn't write.



- ② That method does something risky, something that might not work at runtime.



- ③ You need to know that the method you're calling is risky.



- ④ You then write code that can handle the failure if it does happen. You need to be prepared, just in case.



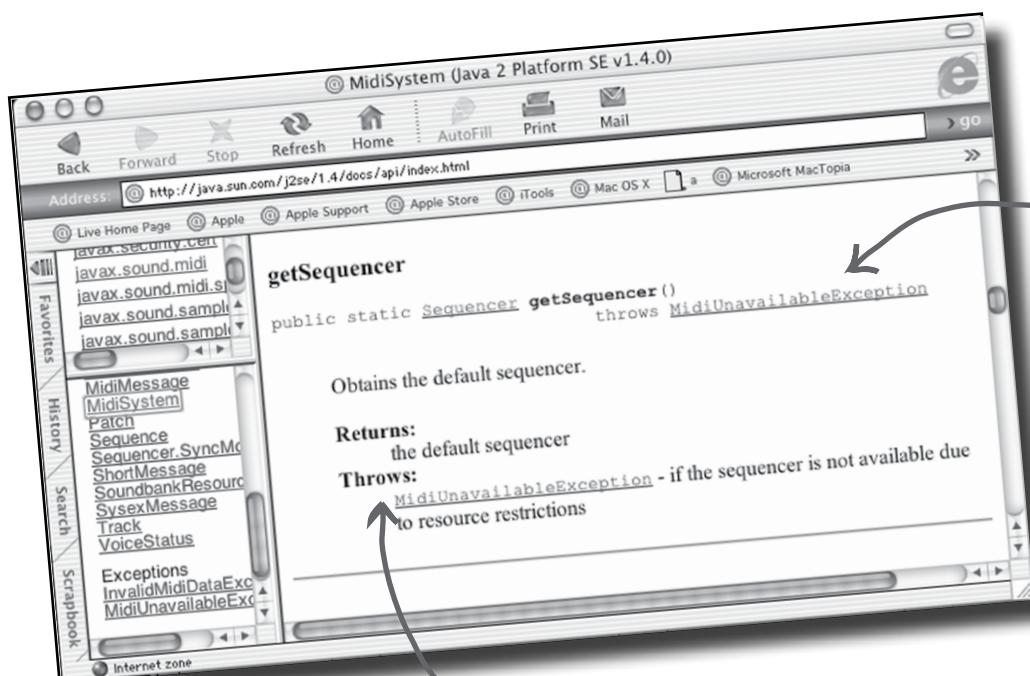
**when things might go wrong**

## Methods in Java use *exceptions* to tell the calling code, “Something Bad Happened. I failed.”

Java’s exception-handling mechanism is a clean, well-lighted way to handle “exceptional situations” that pop up at runtime; it lets you put all your error-handling code in one easy-to-read place. It’s based on you *knowing* that the method you’re calling is risky (i.e. that the method *might* generate an exception), so that you can write code to deal with that possibility. If you *know* you might get an exception when you call a particular method, you can be *prepared* for—possibly even *recover* from—the problem that caused the exception.

So, how *do* you know if a method throws an exception? You find a **throws** clause in the risky method’s declaration.

**The `getSequencer()` method takes a risk. It can fail at runtime.  
So it must ‘declare’ the risk you take when you call it.**



The API does tell you that `getSequencer()` can throw an exception: `MidiUnavailableException`. A method has to declare the exceptions it might throw.

This part tells you WHEN you might get that exception -- in this case, because of resource restrictions (which could just mean the sequencer is already being used).

## The compiler needs to know that YOU know you're calling a risky method.

If you wrap the risky code in something called a **try/catch**, the compiler will relax.

A try/catch block tells the compiler that you *know* an exceptional thing could happen in the method you're calling, and that you're prepared to handle it. That compiler doesn't care *how* you handle it; it cares only that you say you're taking care of it.

```
import javax.sound.midi.*;

public class MusicTest1 {
 public void play() {

 try {
 Sequencer sequencer = MidiSystem.getSequencer(); ← put the risky thing
 System.out.println("Successfully got a sequencer");
 } catch(MidiUnavailableException ex) { ← in a 'try' block.

 System.out.println("Bummer");
 }
 } // close play

 public static void main(String[] args) {
 MusicTest1 mt = new MusicTest1();
 mt.play();
 } // close main
} // close class
```

Dear Compiler,

*I know I'm taking a risk here, but don't you think it's worth it? What should I do?*

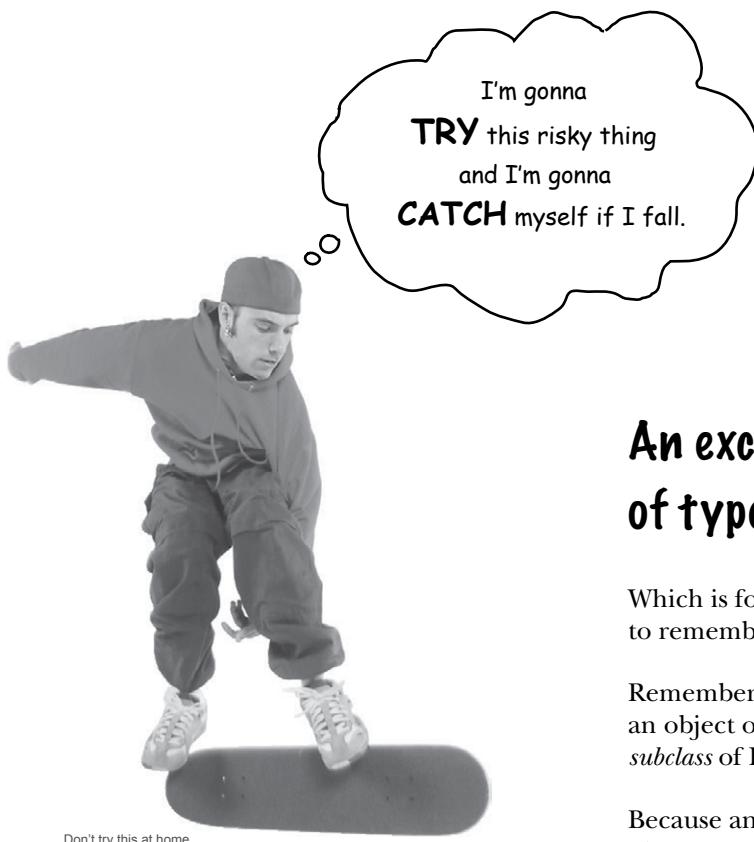
*signed, geeky in Waikiki*

Dear geeky,

*Life is short (especially on the heap). Take the risk. try it. But just in case things don't work out, be sure to catch any problems before all hell breaks loose.*

make a 'catch' block for what to do if the exceptional situation happens -- in other words, a *MidiUnavailableException* is thrown by the call to *getSequencer()*.

exceptions are objects



## An exception is an object... of type Exception.

Which is fortunate, because it would be much harder to remember if exceptions were of type Broccoli.

Remember from your polymorphism chapters that an object of type Exception *can* be an instance of any subclass of Exception.

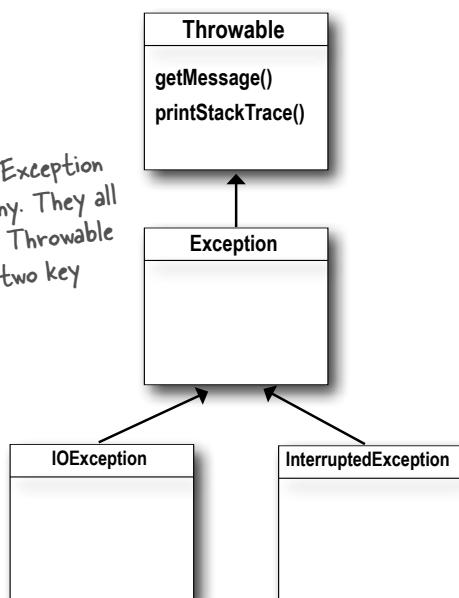
Because an *Exception* is an object, what you *catch* is an object. In the following code, the **catch** argument is declared as type Exception, and the parameter reference variable is *ex*.

```
try {
 // do risky thing
}
} catch (Exception ex) {
 // try to recover
}
```

Part of the Exception class hierarchy. They all extend class Throwable and inherit two key methods.

it's just like declaring a method argument.

This code only runs if an Exception is thrown.



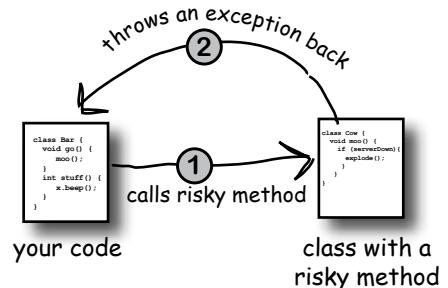
What you write in a catch block depends on the exception that was thrown. For example, if a server is down you might use the catch block to try another server. If the file isn't there, you might ask the user for help finding it.

## If it's your code that catches the exception, then whose code throws it?

You'll spend much more of your Java coding time *handling* exceptions than you'll spend *creating* and *throwing* them yourself. For now, just know that when your code *calls* a risky method—a method that declares an exception—it's the risky method that *throws* the exception back to *you*, the caller.

In reality, it might be you who wrote both classes. It really doesn't matter who writes the code... what matters is knowing which method *throws* the exception and which method *catches* it.

When somebody writes code that could throw an exception, they must *declare* the exception.



### ① Risky, exception-throwing code:

```
public void takeRisk() throws BadException {
 if (abandonAllHope) {
 throw new BadException();
 }
}
```

*create a new object and throw it.*

this method MUST tell the world (by declaring) that it throws a BadException

One method will catch what another method throws. An exception is always thrown back to the caller.

The method that throws has to declare that it might throw the exception.

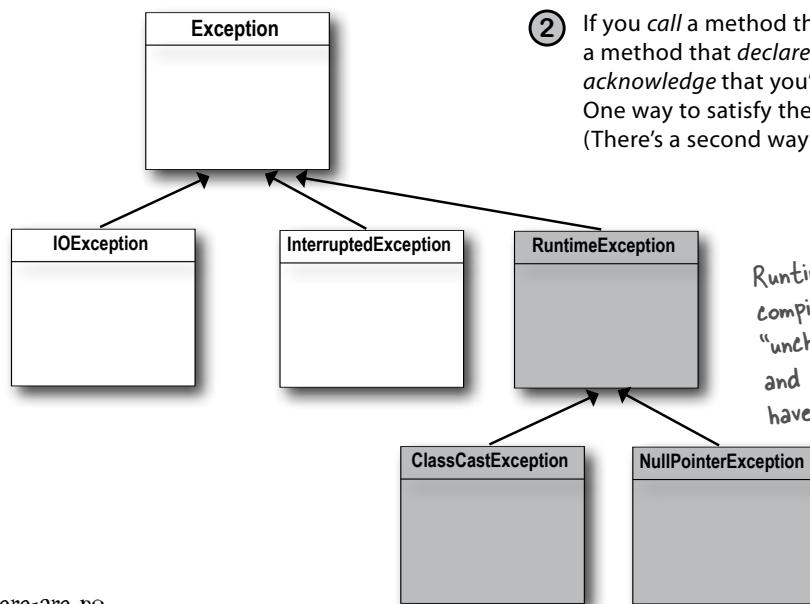
### ② Your code that *calls* the risky method:

```
public void crossFingers() {
 try {
 anObject.takeRisk();
 } catch (BadException ex) {
 System.out.println("Aaargh!");
 ex.printStackTrace();
 }
}
```

If you can't recover from the exception, at LEAST get a stack trace using the printStackTrace() method that all exceptions inherit.

## checked and unchecked exceptions

Exceptions that are NOT subclasses of `RuntimeException` are checked for by the compiler. They're called "checked exceptions"



*there are no  
Dumb Questions*

## The compiler checks for everything except `RuntimeExceptions`.

### The compiler guarantees:

- ① If you *throw* an exception in your code you *must* declare it using the `throws` keyword in your method declaration.
- ② If you *call* a method that throws an exception (in other words, a method that *declares* it throws an exception), you must *acknowledge* that you're aware of the exception possibility. One way to satisfy the compiler is to wrap the call in a try/catch. (There's a second way we'll look at a little later in this chapter.)

`RuntimeExceptions` are NOT checked by the compiler. They're known as (big surprise here) "unchecked exceptions". You can throw, catch, and declare `RuntimeExceptions`, but you don't have to, and the compiler won't check.

**Q:** Wait just a minute! How come this is the FIRST time we've had to try/catch an Exception? What about the exceptions I've already gotten like `NullPointerException` and the exception for `DivideByZero`. I even got a `NumberFormatException` from the `Integer.parseInt()` method. How come we didn't have to catch those?

**A:** The compiler cares about all subclasses of `Exception`, *unless* they are a special type, `RuntimeException`. Any exception class that extends `RuntimeException` gets a free pass. `RuntimeExceptions` can be thrown anywhere, with or without throws declarations or try/catch blocks. The compiler doesn't bother checking whether a method declares that it throws a `RuntimeException`, or whether the caller acknowledges that they might get that exception at runtime.

**Q:** I'll bite. WHY doesn't the compiler care about those runtime exceptions? Aren't they just as likely to bring the whole show to a stop?

**A:** Most `RuntimeExceptions` come from a problem in your code logic, rather than a condition that fails at runtime in ways that you cannot predict or prevent. You *cannot* guarantee the file is there. You *cannot* guarantee the server is up. But you *can* make sure your code doesn't index off the end of an array (that's what the `.length` attribute is for).

You WANT `RuntimeExceptions` to happen at development and testing time. You don't want to code in a try/catch, for example, and have the overhead that goes with it, to catch something that shouldn't happen in the first place.

A try/catch is for handling exceptional situations, not flaws in your code. Use your catch blocks to try to recover from situations you can't guarantee will succeed. Or at the very least, print out a message to the user and a stack trace, so somebody can figure out what happened.

## BULLET POINTS

- A method can throw an exception when something fails at runtime.
  - An exception is always an object of type `Exception`. (Which, as you remember from the polymorphism chapters means the object is from a class that has `Exception` somewhere up its inheritance tree.)
  - The compiler does NOT pay attention to exceptions that are of type `RuntimeException`. A `RuntimeException` does not have to be declared or wrapped in a try/catch (although you're free to do either or both of those things)
  - All Exceptions the compiler cares about are called 'checked exceptions' which really means *compiler-checked* exceptions. Only `RuntimeExceptions` are excluded from compiler checking. All other exceptions must be acknowledged in your code, according to the rules.
  - A method throws an exception with the keyword `throw`, followed by a new exception object:
- ```
throw new NoCaffeineException();
```
- Methods that *might* throw a checked exception **must** announce it with a `throws Exception` declaration.
 - If your code calls a checked-exception-throwing method, it must reassure the compiler that precautions have been taken.
 - If you're prepared to handle the exception, wrap the call in a try/catch, and put your exception handling/recovery code in the catch block.
 - If you're not prepared to handle the exception, you can still make the compiler happy by officially 'ducking' the exception. We'll talk about ducking a little later in this chapter.

Sharpen your pencil

Things you want to do

- connect to a remote server
 access an array beyond its length
 display a window on the screen
 retrieve data from a database
 see if a text file is where you *think* it is
 create a new file
 read a character from the command-line

What might go wrong

the server is down

Which of these do you think might throw an exception that the compiler would care about? We're only looking for the things that you can't control in your code. We did the first one.

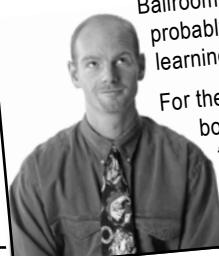
(Because it was the easiest.)

metacognitive tip

If you're trying to learn something new, make that the *last* thing you try to learn before going to sleep. So, once you put this book down (assuming you can tear yourself away from it) don't read anything else more challenging than the back of a Cheerios™ box. Your brain needs time to process what you've read and learned. That could take a few hours. If you try to shove something new in right on top of your Java, some of the Java might not 'stick.'

Of course, this doesn't rule out learning a physical skill. Working on your latest Ballroom KickBoxing routine probably won't affect your Java learning.

For the best results, read this book (or at least look at the pictures) right before going to sleep.



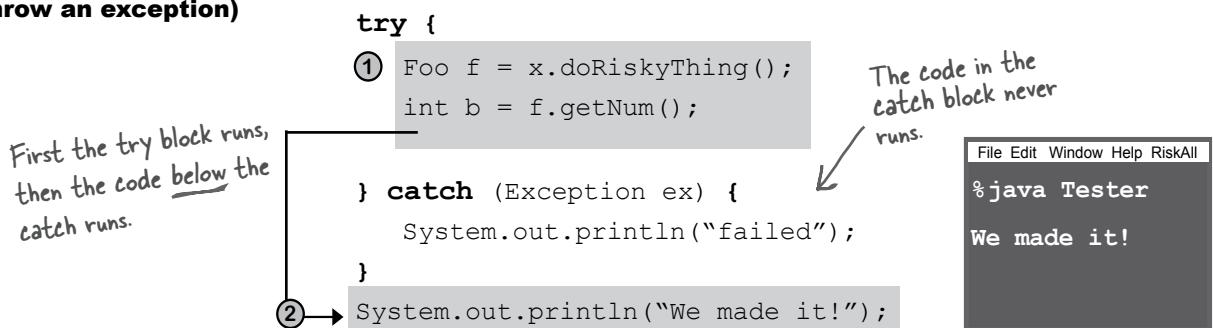
exceptions and flow control

Flow control in try/catch blocks

When you call a risky method, one of two things can happen. The risky method either succeeds, and the try block completes, or the risky method throws an exception back to your calling method.

If the try succeeds

(`doRiskyThing()` does not throw an exception)

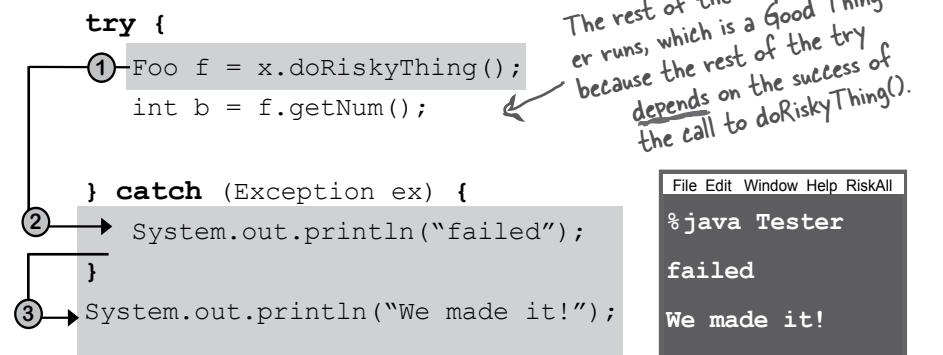


If the try fails

(because `doRiskyThing()` does throw an exception)

The try block runs, but the call to `doRiskyThing()` throws an exception, so the rest of the try block doesn't run.

The catch block runs, then the method continues on.



Finally: for the things you want to do no matter what.

If you try to cook something, you start by turning on the oven.

If the thing you try is a complete **failure**, *you have to turn off the oven*.

If the thing you try **succeeds**, *you have to turn off the oven*.

You have to turn off the oven no matter what!

A finally block is where you put code that must run regardless of an exception.

```
try {
    turnOvenOn();
    x.bake();
} catch (BakingException ex) {
    ex.printStackTrace();
} finally {
    turnOvenOff();
}
```

Without finally, you have to put the turnOvenOff() in *both* the try and the catch because *you have to turn off the oven no matter what*. A finally block lets you put all your important cleanup code in *one* place instead of duplicating it like this:

```
try {
    turnOvenOn();
    x.bake();
    turnOvenOff();
} catch (BakingException ex) {
    ex.printStackTrace();
    turnOvenOff();
}
```



If the try block fails (an exception), flow control immediately moves to the catch block. When the catch block completes, the finally block runs. When the finally block completes, the rest of the method continues on.

If the try block succeeds (no exception), flow control skips over the catch block and moves to the finally block. When the finally block completes, the rest of the method continues on.

If the try or catch block has a return statement, finally will still run! Flow jumps to the finally, then back to the return.